

AD_classification_version_2

September 5, 2025

1 Classifying Alzheimer Disease through Feature Analysis

- Original Paper - <https://ieeexplore.ieee.org/document/11041645>
- Dataset - <https://www.kaggle.com/datasets/brsdincer/alzheimer-features/>
- (Version 1_Experiment 1) Use demented, nondemented, converted data and categorize demented and converted as class 1, nondemented as class 0 - <https://colab.research.google.com/drive/10LysLErBy-Qg2FWZvh2yAr04m2w0cDgq?usp=sharing>
- (This Colab_Version 2_Experiment 2~4) Use only demented and nondemented group, drop converted group which may introduce noise and make the boundary not that clear - https://colab.research.google.com/drive/1N_CkOWYlKjEEvL2YLxS_ZuBVN-AgeYhO?usp=sharing
- (Version 3) Use dataset from original OASIS 2 website, focus on first visit data(150 samples) - <https://colab.research.google.com/drive/1TgLG24-YDWskGqbID0ZdeBd-RB9uHVZw?usp=sharing>

2 Libraries installing

```
[ ]: !pip install lime shap catboost
```

Collecting lime

Downloading lime-0.2.0.1.tar.gz (275 kB)

275.7/275.7

kB 2.3 MB/s eta 0:00:00

Preparing metadata (setup.py) ... done

Requirement already satisfied: shap in /usr/local/lib/python3.11/dist-packages (0.48.0)

Collecting catboost

Downloading catboost-1.2.8-cp311-cp311-manylinux2014_x86_64.whl.metadata (1.2 kB)

Requirement already satisfied: matplotlib in /usr/local/lib/python3.11/dist-packages (from lime) (3.10.0)

Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages (from lime) (2.0.2)

Requirement already satisfied: scipy in /usr/local/lib/python3.11/dist-packages

```

(from lime) (1.16.1)
Requirement already satisfied: tqdm in /usr/local/lib/python3.11/dist-packages
(from lime) (4.67.1)
Requirement already satisfied: scikit-learn>=0.18 in
/usr/local/lib/python3.11/dist-packages (from lime) (1.6.1)
Requirement already satisfied: scikit-image>=0.12 in
/usr/local/lib/python3.11/dist-packages (from lime) (0.25.2)
Requirement already satisfied: pandas in /usr/local/lib/python3.11/dist-packages
(from shap) (2.2.2)
Requirement already satisfied: packaging>20.9 in /usr/local/lib/python3.11/dist-
packages (from shap) (25.0)
Requirement already satisfied: slicer==0.0.8 in /usr/local/lib/python3.11/dist-
packages (from shap) (0.0.8)
Requirement already satisfied: numba>=0.54 in /usr/local/lib/python3.11/dist-
packages (from shap) (0.60.0)
Requirement already satisfied: cloudpickle in /usr/local/lib/python3.11/dist-
packages (from shap) (3.1.1)
Requirement already satisfied: typing-extensions in
/usr/local/lib/python3.11/dist-packages (from shap) (4.14.1)
Requirement already satisfied: graphviz in /usr/local/lib/python3.11/dist-
packages (from catboost) (0.21)
Requirement already satisfied: plotly in /usr/local/lib/python3.11/dist-packages
(from catboost) (5.24.1)
Requirement already satisfied: six in /usr/local/lib/python3.11/dist-packages
(from catboost) (1.17.0)
Requirement already satisfied: llvmlite<0.44,>=0.43.0dev0 in
/usr/local/lib/python3.11/dist-packages (from numba>=0.54->shap) (0.43.0)
Requirement already satisfied: python-dateutil>=2.8.2 in
/usr/local/lib/python3.11/dist-packages (from pandas->shap) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-
packages (from pandas->shap) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-
packages (from pandas->shap) (2025.2)
Requirement already satisfied: networkx>=3.0 in /usr/local/lib/python3.11/dist-
packages (from scikit-image>=0.12->lime) (3.5)
Requirement already satisfied: pillow>=10.1 in /usr/local/lib/python3.11/dist-
packages (from scikit-image>=0.12->lime) (11.3.0)
Requirement already satisfied: imageio!=2.35.0,>=2.33 in
/usr/local/lib/python3.11/dist-packages (from scikit-image>=0.12->lime) (2.37.0)
Requirement already satisfied: tifffile>=2022.8.12 in
/usr/local/lib/python3.11/dist-packages (from scikit-image>=0.12->lime)
(2025.6.11)
Requirement already satisfied: lazy-loader>=0.4 in
/usr/local/lib/python3.11/dist-packages (from scikit-image>=0.12->lime) (0.4)
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.11/dist-
packages (from scikit-learn>=0.18->lime) (1.5.1)
Requirement already satisfied: threadpoolctl>=3.1.0 in
/usr/local/lib/python3.11/dist-packages (from scikit-learn>=0.18->lime) (3.6.0)

```

Requirement already satisfied: contourpy>=1.0.1 in
/usr/local/lib/python3.11/dist-packages (from matplotlib->lime) (1.3.3)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.11/dist-
packages (from matplotlib->lime) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in
/usr/local/lib/python3.11/dist-packages (from matplotlib->lime) (4.59.0)
Requirement already satisfied: kiwisolver>=1.3.1 in
/usr/local/lib/python3.11/dist-packages (from matplotlib->lime) (1.4.9)
Requirement already satisfied: pyparsing>=2.3.1 in
/usr/local/lib/python3.11/dist-packages (from matplotlib->lime) (3.2.3)
Requirement already satisfied: tenacity>=6.2.0 in
/usr/local/lib/python3.11/dist-packages (from plotly->catboost) (9.1.2)
Downloading catboost-1.2.8-cp311-cp311-manylinux2014_x86_64.whl (99.2 MB)

99.2/99.2 MB

25.1 MB/s eta 0:00:00

Building wheels for collected packages: lime

Building wheel for lime (setup.py) ... done

Created wheel for lime: filename=lime-0.2.0.1-py3-none-any.whl size=283834
sha256=7a304cfa1ffa5039c8e358222fd4a3fd0bd421747827ce31a82b38b6dcc9d988

Stored in directory: /root/.cache/pip/wheels/85/fa/a3/9c2d44c9f3cd77cf4e533b58
900b2bf4487f2a17e8ec212a3d

Successfully built lime

Installing collected packages: lime, catboost

Successfully installed catboost-1.2.8 lime-0.2.0.1

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sb
import shap
import lime
```

```
[ ]: from sklearn.metrics import accuracy_score, f1_score, \
    \>roc_auc_score, confusion_matrix, classification_report, precision_score, \
    \>recall_score, roc_curve
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import make_pipeline
from sklearn.model_selection import GridSearchCV, RepeatedStratifiedKFold, \
    \>cross_validate

from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier

from xgboost import XGBClassifier
from lightgbm import LGBMClassifier
```

```
from catboost import CatBoostClassifier
from collections import Counter
```

3 Dataset loading

```
[ ]: df = pd.read_csv('alzheimer.csv')
```

```
[ ]: df.shape
```

```
[ ]: (373, 10)
```

The dataset has 373 samples with 10 features

```
[ ]: df.head()
```

```
[ ]:
      Group M/F  Age  EDUC  SES  MMSE  CDR  eTIV  nWBV  ASF
0  Nondemented  M   87   14  2.0  27.0  0.0  1987  0.696  0.883
1  Nondemented  M   88   14  2.0  30.0  0.0  2004  0.681  0.876
2    Demented  M   75   12  NaN  23.0  0.5  1678  0.736  1.046
3    Demented  M   76   12  NaN  28.0  0.5  1738  0.713  1.010
4    Demented  M   80   12  NaN  22.0  0.5  1698  0.701  1.034
```

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 373 entries, 0 to 372
Data columns (total 10 columns):
 #   Column  Non-Null Count  Dtype
---  -
0   Group   373 non-null      object
1   M/F     373 non-null      object
2   Age     373 non-null      int64
3   EDUC    373 non-null      int64
4   SES     354 non-null      float64
5   MMSE    371 non-null      float64
6   CDR     373 non-null      float64
7   eTIV    373 non-null      int64
8   nWBV    373 non-null      float64
9   ASF     373 non-null      float64
dtypes: float64(5), int64(3), object(2)
memory usage: 29.3+ KB
```

```
[ ]: df.describe()
```

```
[ ]:
      Age      EDUC      SES      MMSE      CDR  \
count  373.000000  373.000000  354.000000  371.000000  373.000000
```

mean	77.013405	14.597855	2.460452	27.342318	0.290885
std	7.640957	2.876339	1.134005	3.683244	0.374557
min	60.000000	6.000000	1.000000	4.000000	0.000000
25%	71.000000	12.000000	2.000000	27.000000	0.000000
50%	77.000000	15.000000	2.000000	29.000000	0.000000
75%	82.000000	16.000000	3.000000	30.000000	0.500000
max	98.000000	23.000000	5.000000	30.000000	2.000000

	eTIV	nWBV	ASF
count	373.000000	373.000000	373.000000
mean	1488.128686	0.729568	1.195461
std	176.139286	0.037135	0.138092
min	1106.000000	0.644000	0.876000
25%	1357.000000	0.700000	1.099000
50%	1470.000000	0.729000	1.194000
75%	1597.000000	0.756000	1.293000
max	2004.000000	0.837000	1.587000

[]:

4 Initial data exploration

```
[ ]: # Distribution chart of each feature
numerical_cols = ['Age', 'EDUC', 'SES', 'MMSE', 'CDR', 'eTIV', 'nWBV', 'ASF']
categorical_cols = ['Group', 'M/F']
all_cols = categorical_cols + numerical_cols


# Set up the matplotlib figure
plt.figure(figsize=(18, 5))

# Loop through each column and plot distribution
for i, col in enumerate(all_cols):
    plt.subplot(2, 5, i + 1)

    if col in numerical_cols:
        sb.histplot(df[col].dropna(), bins=20)
        plt.ylabel("Count")
    else:
        sb.countplot(x=col, data=df)
        plt.ylabel("Count")

    plt.title(f'Distribution of {col}')
    plt.xlabel(col)

plt.tight_layout()
plt.show()
```



#Preprocessing

##Handling Missing Values

From below, we can say that there are 19 missing values in SES and 2 missing values in MMSE.

```
[ ]: df.isnull().sum()
```

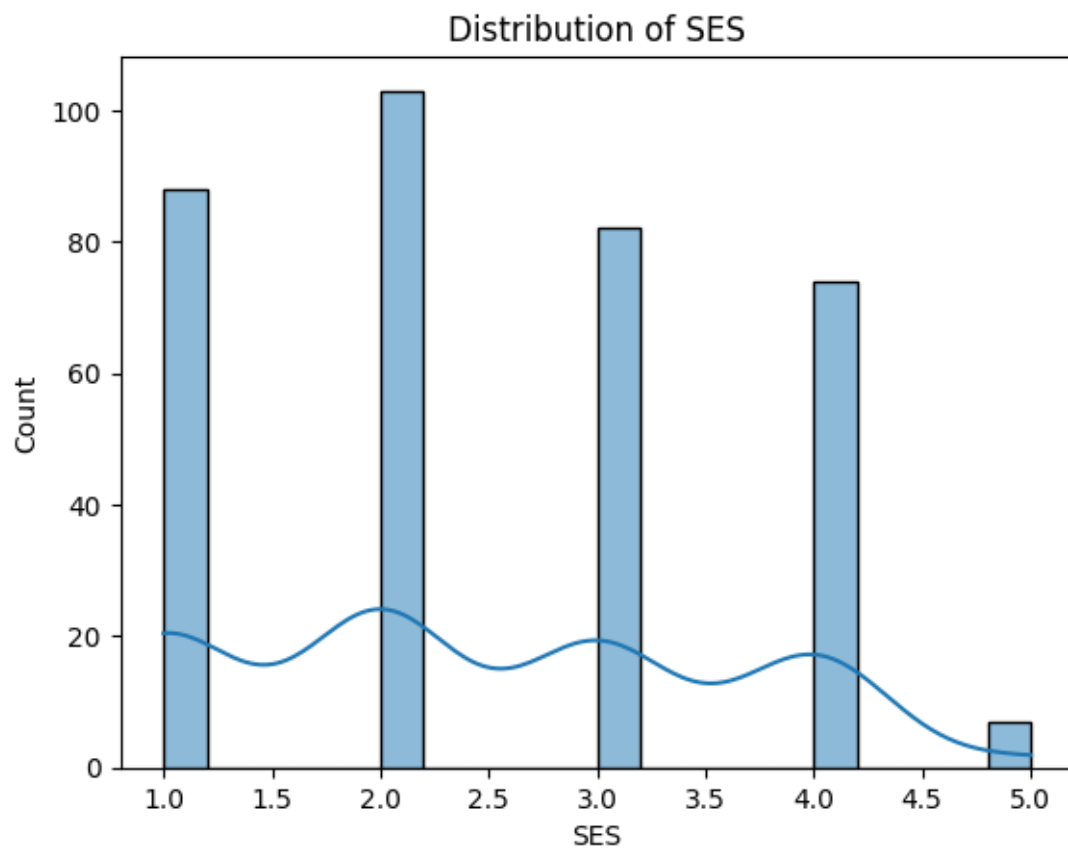
```
[ ]: Group      0
      M/F       0
      Age       0
      EDUC      0
      SES      19
      MMSE      2
      CDR       0
      eTIV      0
      nWBV      0
      ASF       0
      dtype: int64
```


Observing the distribution of 'SES' and 'MMSE' to understand the strategy for imputing missing values.

From below graphs we can say that 'SES' is evenly distributed and we can use 'Mean' as the strategy for imputing, where as 'MMSE' is highly negatively skewed and has outliers, where median would be a better strategy for imputing.

```
[ ]: def plot_distribution(column_name):
      sb.histplot(df[column_name], kde=True, bins=20)
      plt.title(f"Distribution of {column_name}")
      plt.xlabel(column_name)
      plt.ylabel("Count")
      plt.show()
```

```
[ ]: plot_distribution('SES')
      plot_distribution('MMSE')
```





```
[ ]: skew_value = df['SES'].skew()
print(f"Skewness of SES: {skew_value:.2f}")
```

Skewness of SES: 0.22

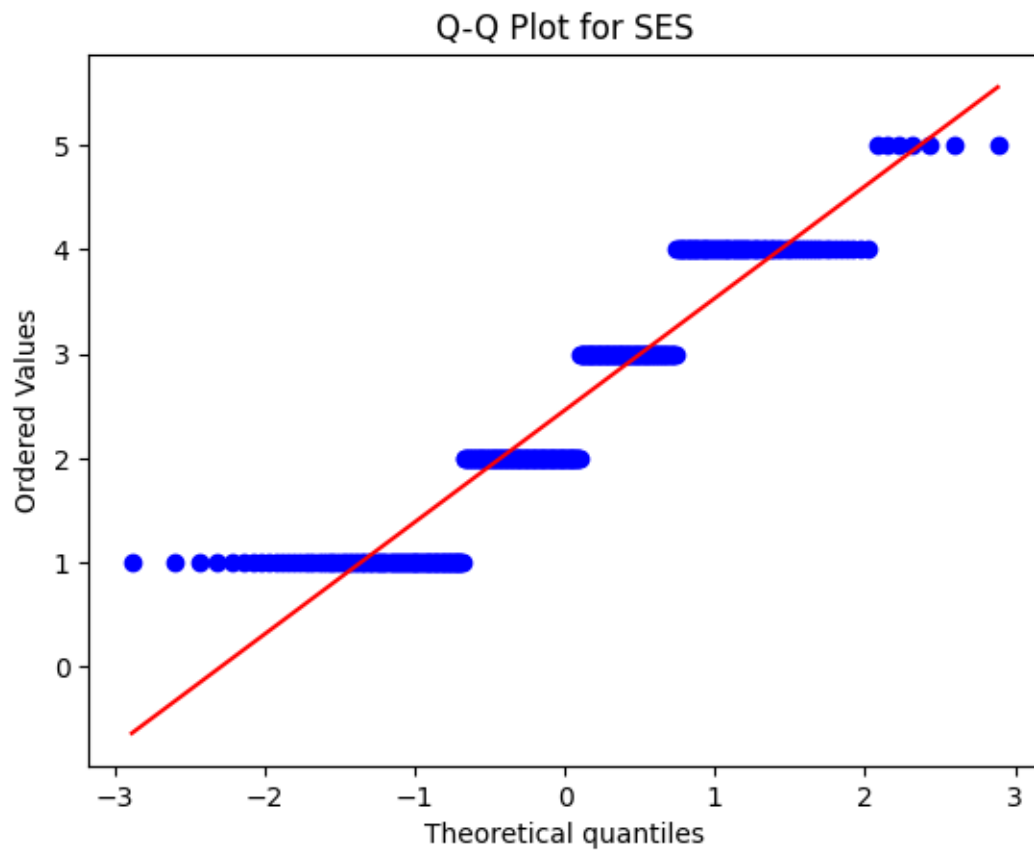
```
[ ]: skew_value = df['MMSE'].skew()
print(f"Skewness of MMSE: {skew_value:.2f}")
```

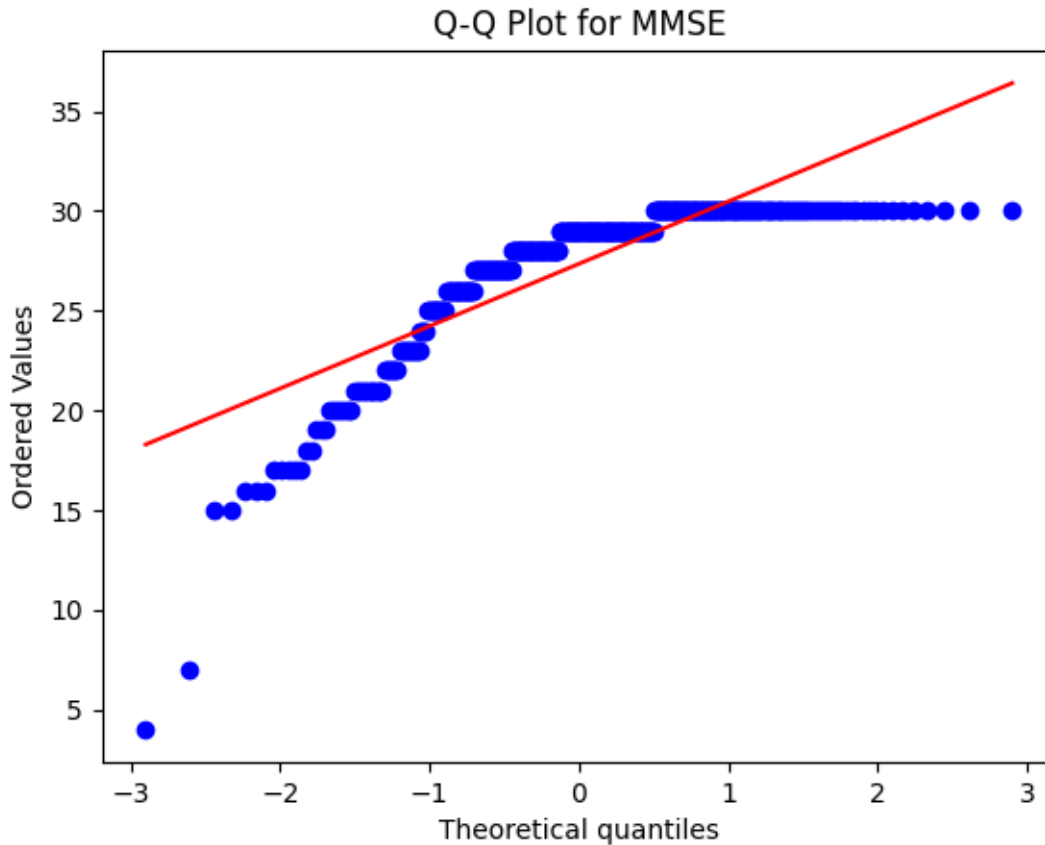
Skewness of MMSE: -2.37

```
[ ]: import scipy.stats as stats

stats.probplot(df['SES'].dropna(), dist="norm", plot=plt)
plt.title("Q-Q Plot for SES")
plt.show()

stats.probplot(df['MMSE'].dropna(), dist="norm", plot=plt)
plt.title("Q-Q Plot for MMSE")
plt.show()
```



```
[ ]: # fill missing values in SES and MMSE columns with average
df['SES'] = df['SES'].fillna(df['SES'].mean())
df['MMSE'] = df['MMSE'].fillna(df['MMSE'].median())
```

```
[ ]: df.isna().sum().sum()
```

```
[ ]: np.int64(0)
```

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 373 entries, 0 to 372
Data columns (total 10 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Group    373 non-null     object
1   M/F      373 non-null     object
2   Age      373 non-null     int64
3   EDUC     373 non-null     int64
4   SES      373 non-null     float64
```

```

5   MMSE      373 non-null    float64
6   CDR       373 non-null    float64
7   eTIV      373 non-null    int64
8   nWBV      373 non-null    float64
9   ASF       373 non-null    float64
dtypes: float64(5), int64(3), object(2)
memory usage: 29.3+ KB

```

##Label Encoding

(Modified)The dataset has three variations of the target value, being Demented(146), Nondemented(190), and Converted(37). For binary classification, it's better to have clear boundary and not to introduce too much noise or unclear information. We will use only demented and nondemented data for this classification task.

```
[ ]: df['M/F'] = df['M/F'].map({'M': 0, 'F': 1})
```

```
[ ]: df['Group'].value_counts()
```

```
[ ]: Group
Nondemented      190
Demented          146
Converted         37
Name: count, dtype: int64
```

```
[ ]: # make a new dataframe to keep only converted data as an additional test dataset
df_converted = df[df['Group'] == 'Converted']
```

```
[ ]: df_converted.head()
df_converted.shape
```

```
[ ]: (37, 10)
```

```
[ ]: # keep only df['Group'] nondemented and demented, remove converted
df = df[df['Group'].isin(['Demented', 'Nondemented'])]
df['Group'].value_counts()
```

```
[ ]: Group
Nondemented      190
Demented          146
Name: count, dtype: int64
```

```
[ ]: df['Group'] = df['Group'].map({'Nondemented': 0, 'Demented': 1})
```

```
[ ]: df['Group'].value_counts()
```

```
[ ]: Group
0      190
1      146
```

Name: count, dtype: int64

```
[ ]: print(df.head(10))
```

	Group	M/F	Age	EDUC	SES	MMSE	CDR	eTIV	nWBV	ASF
0	0	0	87	14	2.000000	27.0	0.0	1987	0.696	0.883
1	0	0	88	14	2.000000	30.0	0.0	2004	0.681	0.876
2	1	0	75	12	2.460452	23.0	0.5	1678	0.736	1.046
3	1	0	76	12	2.460452	28.0	0.5	1738	0.713	1.010
4	1	0	80	12	2.460452	22.0	0.5	1698	0.701	1.034
5	0	1	88	18	3.000000	28.0	0.0	1215	0.710	1.444
6	0	1	90	18	3.000000	27.0	0.0	1200	0.718	1.462
7	0	0	80	12	4.000000	28.0	0.0	1689	0.712	1.039
8	0	0	83	12	4.000000	29.0	0.5	1701	0.711	1.032
9	0	0	85	12	4.000000	30.0	0.0	1699	0.705	1.033

```
[ ]: # export current df to csv file
df.to_csv('alzheimer_cleaned.csv', index=False)
```

5 Quick check on subjective features: CDR and MMSE

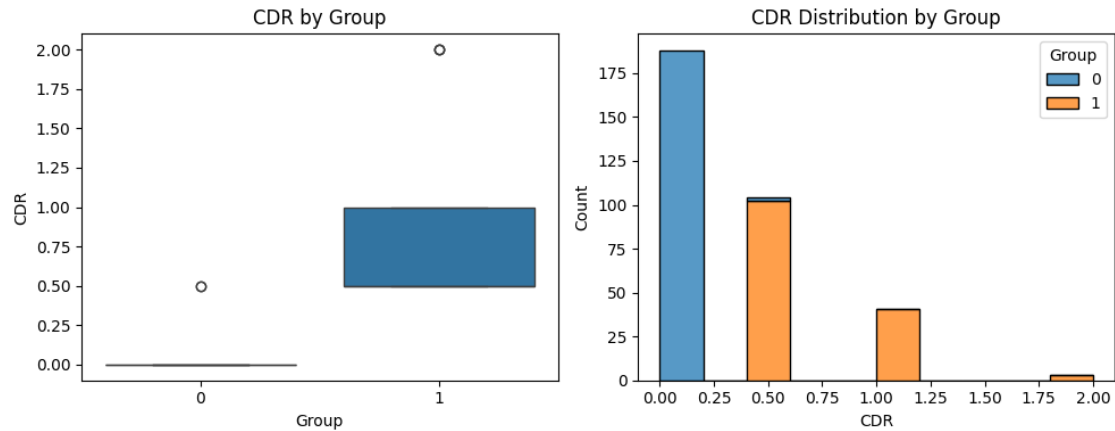
5.1 Check CDR

```
[ ]: # Create a figure with 1 row and 2 columns of subplots
fig, axes = plt.subplots(1, 2, figsize=(10, 4))

# First subplot: Boxplot of CDR by Group
sb.boxplot(x='Group', y='CDR', data=df, ax=axes[0])
axes[0].set_title('CDR by Group')
axes[0].set_xlabel('Group')
axes[0].set_ylabel('CDR')

# Second subplot: Distribution of CDR colored by Group
sb.histplot(data=df, x='CDR', hue='Group', bins=10, multiple='stack',
            ax=axes[1]) # or multiple='dodge'
axes[1].set_title('CDR Distribution by Group')
axes[1].set_xlabel('CDR')
axes[1].set_ylabel('Count')

# Adjust layout to prevent overlap
plt.tight_layout()
plt.show()
```



```
[ ]: # Binarize CDR: 0 stays 0, anything >0 becomes 1
df2 = df[['Group', 'CDR']].dropna().copy()
df2['CDR_bin'] = (df2['CDR'] > 0).astype(int)

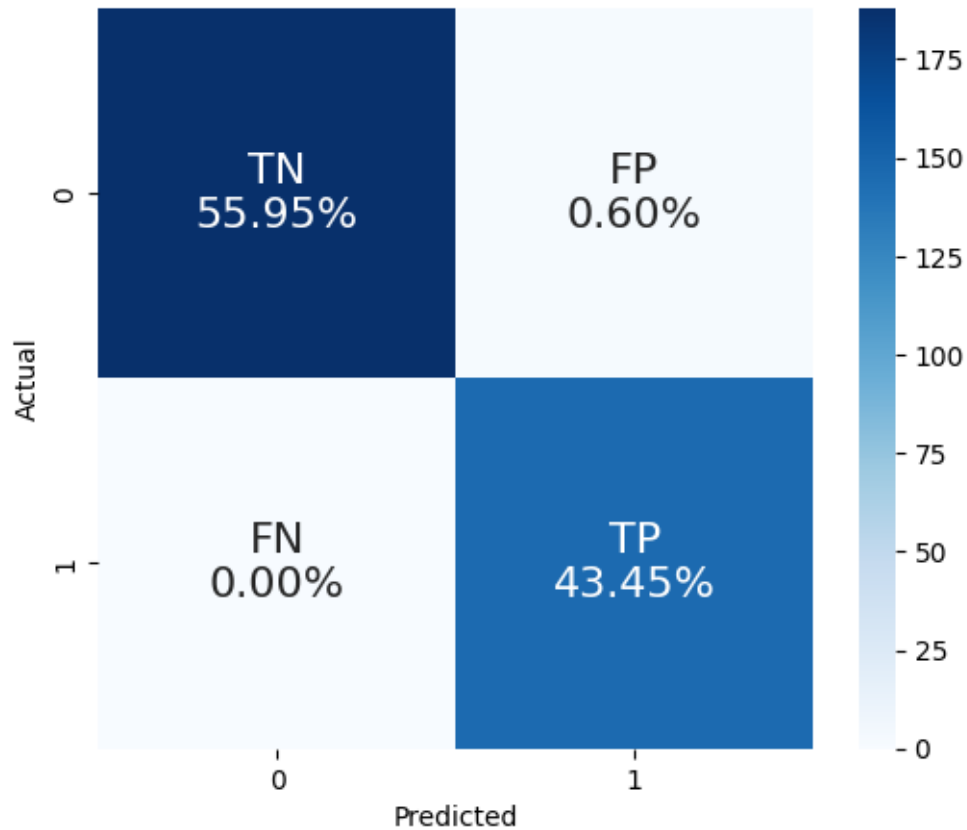
[ ]: cf_matrix = confusion_matrix(df2['Group'], df2['CDR_bin'])

tn, fp, fn, tp = cf_matrix.ravel()/np.sum(cf_matrix)

labels = np.array([
    [f"TN\n{tn:.2%}", f"FP\n{fp:.2%}"],
    [f"FN\n{fn:.2%}", f"TP\n{tp:.2%}"]
])

plt.figure(figsize=(6, 5))

sb.heatmap(cf_matrix, annot=labels,\
            fmt='', cmap='Blues', annot_kws={'size':16})
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```



Group is the ground truth, use CDR_bin to predict ground truth could get 99.4% accuracy. which means that they are almost identical, use CDR_bin is like cheating or data leakage.

```
[ ]: # A quick model running on only one feature_CDR to check if it's cheating to use it

# Use CDR as the only feature
X = df[['CDR']]
y = df['Group']

# train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# train model
model = RandomForestClassifier(random_state=42)
model.fit(X_train, y_train)

# evaluate
y_pred = model.predict(X_test)
```

```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	38
1	1.00	1.00	1.00	30
accuracy			1.00	68
macro avg	1.00	1.00	1.00	68
weighted avg	1.00	1.00	1.00	68

Using only one feature CDR: The value of all the metrics(Accuracy precision recall f1-score) are all 1, means that CDR is highly predictive-strong evidence of target leakage, like another representation of target group. Using CDR as a feature to predict group target is cheating which must be removed from our dataset.

In addition, just because CDR is a cheating feature to predict group target, we can exploit it as the key feature to test our converted dataset(a combination of demented and non-demented) to determine their real group(non-demented or demented)

```
[ ]: X_test_converted = df_converted[['CDR']]
      y_pred_converted = model.predict(X_test_converted)
```

```
[ ]: df_converted['Group'] = y_pred_converted
```

```
[ ]: df_converted.head(5)
```

```
[ ]:
      Group  M/F  Age  EDUC  SES  MMSE  CDR  eTIV  nWBV  ASF
33      0     1   87    14   1.0  30.0  0.0  1406  0.715  1.248
34      0     1   88    14   1.0  29.0  0.0  1398  0.713  1.255
35      1     1   92    14   1.0  27.0  0.5  1423  0.696  1.234
36      0     0   80    20   1.0  29.0  0.0  1587  0.693  1.106
37      1     0   82    20   1.0  28.0  0.5  1606  0.677  1.093
```

```
[ ]: # combine df and df_converted dataset
      # df_combined = pd.concat([df, df_converted])
```

```
[ ]: # df_combined['Group'].value_counts()
```

```
[ ]: # df_combined.head(5)
```

5.2 Check MMSE

```
[ ]: # Create a figure with 1 row and 2 columns of subplots
      fig, axes = plt.subplots(1, 2, figsize=(10, 4))

      # First subplot: Boxplot of MMSE by Group
```


```

sb.boxplot(x='Group', y='MMSE', data=df, ax=axes[0])
axes[0].set_title('MMSE by Group')
axes[0].set_xlabel('Group')
axes[0].set_ylabel('MMSE')

# Second subplot: Distribution of MMSE colored by Group
sb.histplot(data=df, x='MMSE', hue='Group', bins=10, multiple='stack',
            ↪ax=axes[1]) # or multiple='dodge'
axes[1].set_title('MMSE Distribution by Group')
axes[1].set_xlabel('MMSE')
axes[1].set_ylabel('Count')

# Adjust layout to prevent overlap
plt.tight_layout()
plt.show()

```



```

[ ]: # A quick model running on only one feature_MMSE to check if it's cheating to
      ↪use it

# use only MMSE as feature
X = df[['MMSE']]
y = df['Group']

# train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
            ↪random_state=42)

# train model
model = RandomForestClassifier(random_state=42)
model.fit(X_train, y_train)

```



```
# evaluate
y_pred = model.predict(X_test)
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.88	0.97	0.93	38
1	0.96	0.83	0.89	30
accuracy			0.91	68
macro avg	0.92	0.90	0.91	68
weighted avg	0.92	0.91	0.91	68

Using only one feature MMSE: Accuracy is pretty high 0.91, means that MMSE is also a highly predictive-strong evidence of target leakage. Using MMSE as a feature to predict group target is like cheating which must be also removed from our dataset for fairness.


5.3 (additional) Check nWBV

```
[ ]: # Create a figure with 1 row and 2 columns of subplots
fig, axes = plt.subplots(1, 2, figsize=(10, 4))

# First subplot: Boxplot of nWBV by Group
sb.boxplot(x='Group', y='nWBV', data=df, ax=axes[0])
axes[0].set_title('nWBV by Group')
axes[0].set_xlabel('Group')
axes[0].set_ylabel('nWBV')

# Second subplot: Distribution of nWBV colored by Group
sb.histplot(data=df, x='nWBV', hue='Group', bins=10, multiple='stack',
            ax=axes[1]) # or multiple='dodge'
axes[1].set_title('nWBV Distribution by Group')
axes[1].set_xlabel('nWBV')
axes[1].set_ylabel('Count')

# Adjust layout to prevent overlap
plt.tight_layout()
plt.show()
```



```
[ ]: # A quick model running on only one feature_MMSE to check if it's cheating to use it

# use only MMSE as feature
X = df[['nWBV']]
y = df['Group']

# train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# train model
model = RandomForestClassifier(random_state=42)
model.fit(X_train, y_train)

# evaluate
y_pred = model.predict(X_test)
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.56	0.58	0.57	38
1	0.45	0.43	0.44	30
accuracy			0.51	68
macro avg	0.51	0.51	0.51	68
weighted avg	0.51	0.51	0.51	68

Using only one feature nWBV: Accuracy is only 0.50 which means that only nWBV cannot predict the result very well.

6 Experiment 2 __Original process(key paper replication with all features)

```
[ ]: y = df['Group']  
X = df.drop(['Group'], axis = 1)
```

```
[ ]: print(X.head(10))
```

	M/F	Age	EDUC	SES	MMSE	CDR	eTIV	nWBV	ASF
0	0	87	14	2.000000	27.0	0.0	1987	0.696	0.883
1	0	88	14	2.000000	30.0	0.0	2004	0.681	0.876
2	0	75	12	2.460452	23.0	0.5	1678	0.736	1.046
3	0	76	12	2.460452	28.0	0.5	1738	0.713	1.010
4	0	80	12	2.460452	22.0	0.5	1698	0.701	1.034
5	1	88	18	3.000000	28.0	0.0	1215	0.710	1.444
6	1	90	18	3.000000	27.0	0.0	1200	0.718	1.462
7	0	80	12	4.000000	28.0	0.0	1689	0.712	1.039
8	0	83	12	4.000000	29.0	0.5	1701	0.711	1.032
9	0	85	12	4.000000	30.0	0.0	1699	0.705	1.033

```
[ ]: print(y.head(10))
```

```
0    0  
1    0  
2    1  
3    1  
4    1  
5    0  
6    0  
7    0  
8    0  
9    0  
Name: Group, dtype: int64
```

Here we can observe that the data contains nearly double the counts of '0' to '1'.

```
[ ]: y.value_counts()
```

```
[ ]: Group  
0    190  
1    146  
Name: count, dtype: int64
```

##Train test split

```
[ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2,  
↳ random_state = 42)
```

6.1 Model Training

Classifier Models

```
[ ]: models = [  
    "XGBoost",  
    "LightGBM",  
    "CatBoost"  
]  
  
classifiers = [  
    XGBClassifier(),  
    LGBMClassifier(verbosity=-1),  
    CatBoostClassifier(verbose = False)  
]
```

##Hyper-Parameter tuning

Parameter-grid for Hyper-parameter tuning

```
[ ]: parameters_grid = [  
    {  
        'xgbclassifier__n_estimators': [100, 200],  
        'xgbclassifier__max_depth': [3, 5, 7],  
        'xgbclassifier__learning_rate': [0.001, 0.01, 0.1],  
        'xgbclassifier__subsample': [0.7, 0.8, 0.9],  
        'xgbclassifier__colsample_bytree': [0.7, 0.8, 1.0],  
        'xgbclassifier__min_child_weight': [1, 3, 5]  
    },  
  
    {  
        'lgbmclassifier__n_estimators': [100, 200],  
        'lgbmclassifier__max_depth': [3, 5, 7],  
        'lgbmclassifier__learning_rate': [0.001, 0.01, 0.1],  
        'lgbmclassifier__num_leaves': [31, 50, 70]  
    },  
  
    {  
        'catboostclassifier__min_data_in_leaf': [20, 40, 60],  
        'catboostclassifier__rsm': [0.7, 0.8, 1.0],  
        'catboostclassifier__iterations': [100, 200],  
        'catboostclassifier__depth': [3, 5, 7],  
        'catboostclassifier__learning_rate': [0.001, 0.01, 0.1],  
        'catboostclassifier__l2_leaf_reg': [1, 3, 5],  
        'catboostclassifier__bagging_temperature': [0, 0.5, 1]  
    }  
]
```

```
[ ]: from sklearn.base import TransformerMixin, BaseEstimator

class DataFrameStandardScaler(TransformerMixin, BaseEstimator):
    def __init__(self):
        self.scaler = StandardScaler()
        self.columns = None

    def fit(self, X, y=None):
        self.columns = X.columns
        self.scaler.fit(X)
        return self

    def transform(self, X):
        scaled = self.scaler.transform(X)
        return pd.DataFrame(scaled, columns=self.columns, index=X.index)

[ ]: # A dictionary to store trained models
trained_models = {}

for name, clf, param_grid in zip(models, classifiers, parameters_grid):
    pipeline = make_pipeline(DataFrameStandardScaler(), clf)

    grid = GridSearchCV(pipeline, param_grid, cv = 5, scoring = 'accuracy')
    grid.fit(X_train, y_train)

    best_model = grid.best_estimator_

    # Storing the models with their best parameters for SHAP analysis
    trained_models[name] = {
        'model': best_model,
        'best_params': grid.best_params_
    }

    y_pred = best_model.predict(X_test)

    print("-----")

    print(name)
    print("Accuracy:", accuracy_score(y_test, y_pred))
    print("Precision:", precision_score(y_test, y_pred))
    print("Recall:", recall_score(y_test, y_pred))
    print("F1 Score:", f1_score(y_test, y_pred))
    print("AUC:", roc_auc_score(y_test, y_pred))
    print("Best Parameters:", grid.best_params_)
```

XGBoost

```

Accuracy: 1.0
Precision: 1.0
Recall: 1.0
F1 Score: 1.0
AUC: 1.0
Best Parameters: {'xgbclassifier__colsample_bytree': 0.7,
'xgbclassifier__learning_rate': 0.01, 'xgbclassifier__max_depth': 3,
'xgbclassifier__min_child_weight': 1, 'xgbclassifier__n_estimators': 100,
'xgbclassifier__subsample': 0.7}
-----

```

```

LightGBM
Accuracy: 1.0
Precision: 1.0
Recall: 1.0
F1 Score: 1.0
AUC: 1.0
Best Parameters: {'lgbmclassifier__learning_rate': 0.001,
'lgbmclassifier__max_depth': 3, 'lgbmclassifier__n_estimators': 200,
'lgbmclassifier__num_leaves': 31}
-----

```

```

CatBoost
Accuracy: 1.0
Precision: 1.0
Recall: 1.0
F1 Score: 1.0
AUC: 1.0
Best Parameters: {'catboostclassifier__bagging_temperature': 0,
'catboostclassifier__depth': 3, 'catboostclassifier__iterations': 100,
'catboostclassifier__l2_leaf_reg': 1, 'catboostclassifier__learning_rate':
0.001, 'catboostclassifier__min_data_in_leaf': 20, 'catboostclassifier__rsm':
0.7}

```

##Cross-Validation

```

[ ]: from sklearn.model_selection import RepeatedStratifiedKFold, cross_validate, LeaveOneOut

print("Cross Validation: ")

scoring = {
    'accuracy': 'accuracy',
    'precision': 'precision',
    'recall': 'recall',
    'f1': 'f1',
    'roc_auc': 'roc_auc'
}

for i, (name, data) in enumerate(trained_models.items()):

```

```

model = data['model']

# Different random state for each model
cv_different = RepeatedStratifiedKFold(n_splits=5, n_repeats=2,
↪random_state=42+i*10)

scores = cross_validate(model, X_train, y_train, cv=cv_different,
↪scoring=scoring)

print(f"\n{name} (Random State {42+i*10}):")
for metric in scoring.keys():
    mean_score = scores[f'test_{metric}'].mean()
    std_score = scores[f'test_{metric}'].std()
    print(f"{metric.capitalize()}: {mean_score:.3f} ± {std_score:.3f}")

```

Cross Validation:

XGBoost (Random State 42):

Accuracy: 0.993 ± 0.012

Precision: 0.984 ± 0.027

Recall: 1.000 ± 0.000

F1: 0.992 ± 0.014

Roc_auc: 0.997 ± 0.004

LightGBM (Random State 52):

Accuracy: 0.993 ± 0.009

Precision: 0.984 ± 0.020

Recall: 1.000 ± 0.000

F1: 0.992 ± 0.010

Roc_auc: 0.998 ± 0.003

CatBoost (Random State 62):

Accuracy: 0.993 ± 0.009

Precision: 0.984 ± 0.020

Recall: 1.000 ± 0.000

F1: 0.992 ± 0.010

Roc_auc: 0.995 ± 0.006

```

[ ]: from sklearn.model_selection import cross_val_score, LeaveOneOut

loo = LeaveOneOut()

print("Leave-One-Out Cross Validation:")

for i, (name, data) in enumerate(trained_models.items()):
    model = data['model']

```

```

# cross_val_score automatically fits the model on each train/test split
accuracy_scores = cross_val_score(model, X_train, y_train, cv=loo,
↳scoring='accuracy')

print(f"\n{name}:")
print(f"Accuracy: {accuracy_scores.mean():.3f} ± {accuracy_scores.std():.
↳3f}")
print(f"Individual Scores (first 10): {[f'{score:.3f}' for score in
↳accuracy_scores[:10]]}")

```

Leave-One-Out Cross Validation:

XGBoost:

Accuracy: 0.993 ± 0.086

Individual Scores (first 10): ['1.000', '1.000', '1.000', '1.000', '1.000',
'1.000', '1.000', '1.000', '1.000', '1.000']

LightGBM:

Accuracy: 0.993 ± 0.086

Individual Scores (first 10): ['1.000', '1.000', '1.000', '1.000', '1.000',
'1.000', '1.000', '1.000', '1.000', '1.000']

CatBoost:

Accuracy: 0.993 ± 0.086

Individual Scores (first 10): ['1.000', '1.000', '1.000', '1.000', '1.000',
'1.000', '1.000', '1.000', '1.000', '1.000']

##ROC Curve

```

[ ]: from sklearn.metrics import roc_curve, auc

fig, axes = plt.subplots(1, 3, figsize=(15, 5))

for i, (name, data) in enumerate(trained_models.items()):
    # Train ROC
    y_train_proba = data['model'].predict_proba(X_train)[: , 1]
    fpr_train, tpr_train, _ = roc_curve(y_train, y_train_proba)
    axes[i].plot(fpr_train, tpr_train, 'b--', label=f'Train (AUC={auc(fpr_train,
↳tpr_train):.3f})')

    # Test ROC
    y_test_proba = data['model'].predict_proba(X_test)[: , 1]
    fpr_test, tpr_test, _ = roc_curve(y_test, y_test_proba)
    axes[i].plot(fpr_test, tpr_test, 'r-', label=f'Test (AUC={auc(fpr_test,
↳tpr_test):.3f})')

    axes[i].plot([0,1], [0,1], 'k--', alpha=0.5)

```




```

axes[i].set_title(name)
axes[i].legend()

plt.tight_layout()
plt.show()

```



##Confusion Matrix

```
[ ]: print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
```

Confusion Matrix:

```
[[38  0]
 [ 0 30]]
```

##Classification Report

```
[ ]: print("Classification Report:\n", classification_report(y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	38
1	1.00	1.00	1.00	30
accuracy			1.00	68
macro avg	1.00	1.00	1.00	68
weighted avg	1.00	1.00	1.00	68

##SHAP Analysis

###SHAP Bar plot

```
[ ]: for name, data in trained_models.items():
    model = data["model"].named_steps[list(data["model"].named_steps.
    ↪keys())[-1]]
```

```

explainer = shap.Explainer(model)
shap_values = explainer(X_test)


print(f"\n{name} - Best Parameters: {data['best_params']}")
print(f"SHAP values shape: {shap_values.values.shape}")
print(f>Data shape: {shap_values.data.shape}")

fig = plt.figure(figsize=(10,6))
plt.text(0.5, 1.05, f"SHAP Feature Importance - {name}", ha='center',
↪fontsize=16, transform=plt.gca().transAxes)


shap.plots.bar(shap_values)

```

XGBoost - Best Parameters: {'xgbclassifier__colsample_bytree': 0.7,
'xgbclassifier__learning_rate': 0.01, 'xgbclassifier__max_depth': 3,
'xgbclassifier__min_child_weight': 1, 'xgbclassifier__n_estimators': 100,
'xgbclassifier__subsample': 0.7}
SHAP values shape: (68, 9)
Data shape: (68, 9)




LightGBM - Best Parameters: {'lgbmclassifier__learning_rate': 0.001,
'lgbmclassifier__max_depth': 3, 'lgbmclassifier__n_estimators': 200,
'lgbmclassifier__num_leaves': 31}
SHAP values shape: (68, 9)
Data shape: (68, 9)



CatBoost - Best Parameters: {'catboostclassifier__bagging_temperature': 0,
'catboostclassifier__depth': 3, 'catboostclassifier__iterations': 100,
'catboostclassifier__l2_leaf_reg': 1, 'catboostclassifier__learning_rate':
0.001, 'catboostclassifier__min_data_in_leaf': 20, 'catboostclassifier__rsm':
0.7}
SHAP values shape: (68, 9)
Data shape: (68, 9)

SHAP Feature Importance - CatBoost

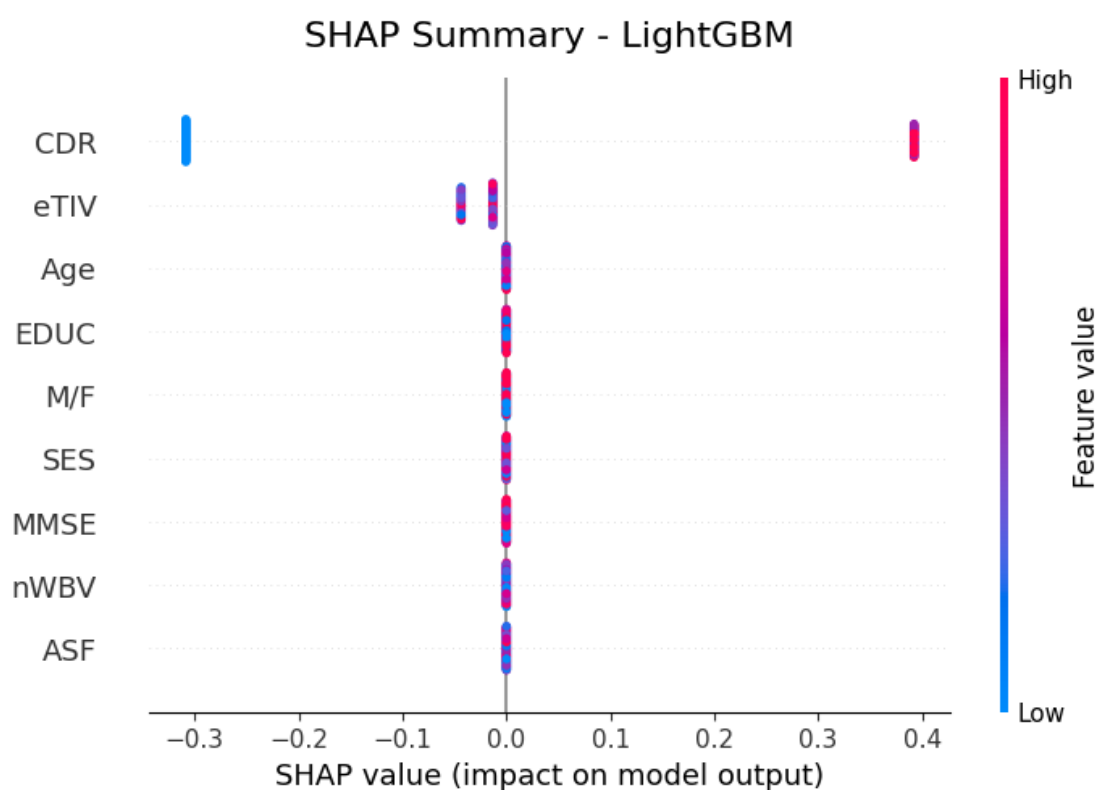




###SHAP Beeswarm plot

```
[ ]: for name, data in trained_models.items():
    model = data["model"].named_steps[list(data["model"].named_steps.
    ↪keys())[-1]]
    explainer = shap.Explainer(model)
    shap_values = explainer(X_test)

    fig = plt.figure(figsize=(10,6))
    plt.text(0.5, 1.05, f"SHAP Summary - {name}", ha='center', fontsize=16,
    ↪transform=plt.gca().transAxes)

    shap.plots.beeswarm(shap_values)
```





7 Experiment 3 & 4 _Parallel Process(remove CDR and MMSE)

7.1 Preprocessing

```
[ ]: df.shape
```

```
[ ]: (336, 10)
```

```
[ ]:
```

```
[ ]: df['Group'].value_counts()
```

```
[ ]: Group
0    190
1    146
Name: count, dtype: int64
```

```
[ ]: df.info()
```

```

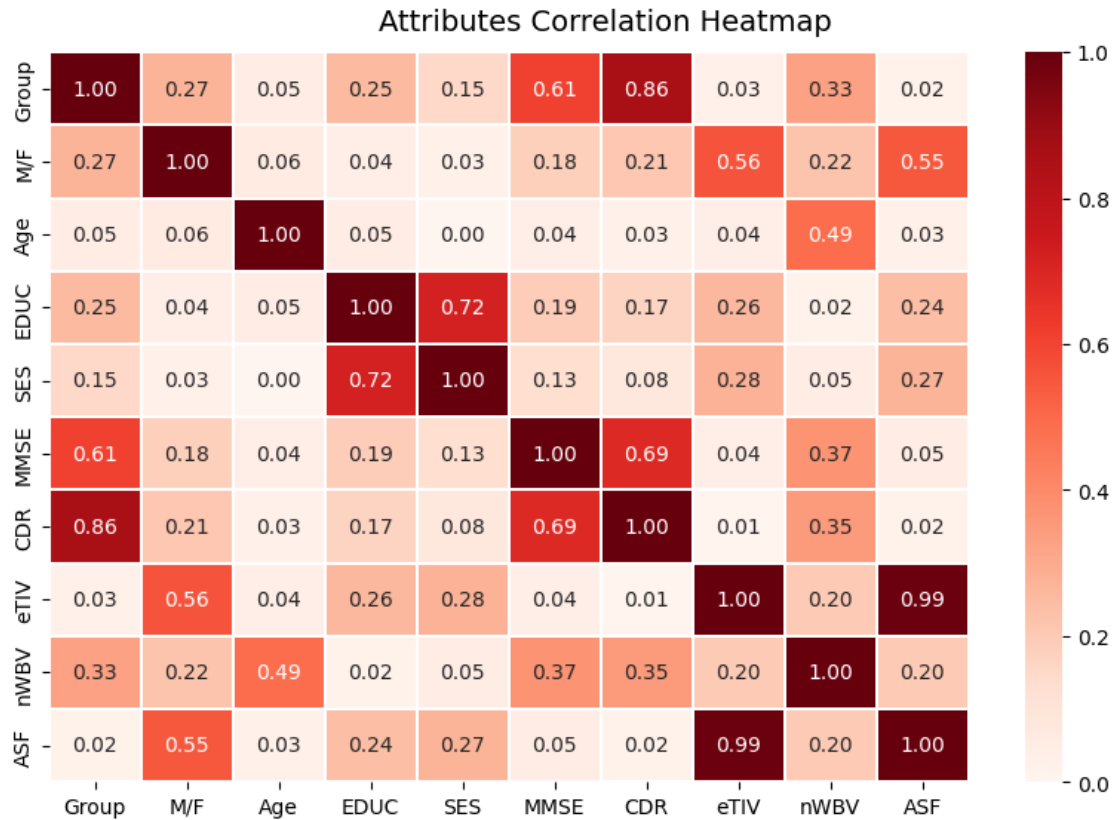
<class 'pandas.core.frame.DataFrame'>
Index: 336 entries, 0 to 372
Data columns (total 10 columns):
#   Column   Non-Null Count  Dtype
---  -
0   Group    336 non-null    int64
1   M/F      336 non-null    int64
2   Age      336 non-null    int64
3   EDUC     336 non-null    int64
4   SES      336 non-null    float64
5   MMSE     336 non-null    float64
6   CDR      336 non-null    float64
7   eTIV     336 non-null    int64
8   nWBV     336 non-null    float64
9   ASF      336 non-null    float64
dtypes: float64(5), int64(5)
memory usage: 28.9 KB

```

```

[ ]: # for visualizing correlations
f, ax = plt.subplots(figsize=(10, 6))
corr = df.corr().abs()
hm = sb.heatmap(round(corr,2), annot=True, ax=ax, cmap="Reds",fmt='.2f',
                 linewidths=.05)
f.subplots_adjust(top=0.93)
t= f.suptitle('Attributes Correlation Heatmap', fontsize=14)

```



ASF and ETIV are strongly related with correlation coefficient $0.99 > 0.85$ (threshold), which make sense because:

eTIV Estimated Total Intracranial Volume Template Intracranial Volume * ASF (Atlas Scaling Factor),

we can consider keep only one of them.

In previous process, we found that use MMSE or CDR to predict is a possible cheating method, so we will remove both of them in our models.

```
[ ]: X = df.drop(['Group', 'CDR', 'MMSE', 'ASF'], axis = 1)
     y = df['Group']
```

```
[ ]: # Split dataset into training and test sets (80% training, 20% testing)
     X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
     random_state=42)

     # Standardize the features skip this because we prefer to see the original data
     for better explanation
     #scaler = StandardScaler()
     #X_train = scaler.fit_transform(X_train)
     #X_test = scaler.transform(X_test)
```



```
[ ]: X_test.columns
```

```
[ ]: Index(['M/F', 'Age', 'EDUC', 'SES', 'eTIV', 'nWBV'], dtype='object')
```

7.2 Model Training

7.2.1 XGBoost

```
[ ]: # Initialize a XGBoost classifier
XGBoost= XGBClassifier(random_state=42)

# Define the parameter grid for Grid
XGB_param_dist = {
    'n_estimators': [100, 200],
    'max_depth': [3, 5, 7],
    'learning_rate': [0.001, 0.01, 0.1],
    'subsample': [0.7, 0.8, 0.9],
    'colsample_bytree': [0.7, 0.8, 1.0],
    'min_child_weight': [1, 3, 5],
}

grid_search = GridSearchCV(XGBoost, XGB_param_dist, cv=5, scoring='f1',
                             n_jobs=-1)
grid_search.fit(X_train, y_train)
print("Best hyperparameters found by GridSearchCV:")
print(grid_search.best_params_)

XGBoost_md1 = grid_search.best_estimator_
y_pred = XGBoost_md1.predict(X_test)
```

Best hyperparameters found by GridSearchCV:

```
{'colsample_bytree': 1.0, 'learning_rate': 0.1, 'max_depth': 5,
 'min_child_weight': 1, 'n_estimators': 200, 'subsample': 0.9}
```

XGBoost Best hyperparameters found by GridSearchCV: {'colsample_bytree': 1.0, 'learning_rate': 0.1, 'max_depth': 5, 'min_child_weight': 1, 'n_estimators': 200, 'subsample': 0.9}

```
[ ]: print("XGBoost: ")
print("\nClassification Report on Test Set:")
print(classification_report(y_test, y_pred))
print("-----")

print("Accuracy:", accuracy_score(y_test, y_pred))
print("Precision:", precision_score(y_test, y_pred))
print("Recall:", recall_score(y_test, y_pred))
print("F1 Score:", f1_score(y_test, y_pred))
print("ROC_AUC:", roc_auc_score(y_test, y_pred))
print("-----")
```

```
# Get feature importances
XGBoost_feature_importances = pd.DataFrame({
    "Feature": X_test.columns,
    "Importance": XGBoost_mdl.feature_importances_
})

XGBoost_feature_importances = XGBoost_feature_importances.
    ↪sort_values(by="Importance", ascending=False)
XGBoost_feature_importances
```

XGBoost:

Classification Report on Test Set:

	precision	recall	f1-score	support
0	0.85	0.87	0.86	38
1	0.83	0.80	0.81	30
accuracy			0.84	68
macro avg	0.84	0.83	0.84	68
weighted avg	0.84	0.84	0.84	68

```
-----
Accuracy: 0.8382352941176471
Precision: 0.8275862068965517
Recall: 0.8
F1 Score: 0.8135593220338984
ROC_AUC: 0.8342105263157895
-----
```

```
[ ]: Feature Importance
2    EDUC    0.236373
0     M/F    0.216841
5    nWBV    0.152848
3     SES    0.137124
1     Age    0.131612
4    eTIV    0.125202
```

```
[ ]: xgb = XGBClassifier(random_state=42,n_jobs=-1)

param_grid = {
    'n_estimators': [100, 200],
    'max_depth': [3, 5, 7],
    'learning_rate': [0.001, 0.01, 0.1],
    'subsample': [0.7, 0.8, 0.9],
    'colsample_bytree': [0.7, 0.8, 1.0],
```

```

        'min_child_weight': [1, 3, 5],
    }

    inner_cv = RepeatedStratifiedKFold(
        n_splits=5, n_repeats=2, random_state=42
    )

    grid = GridSearchCV(
        estimator=xgb,
        param_grid=param_grid,
        scoring='f1',
        cv=inner_cv,
        n_jobs=-1,
        verbose=0
    )

    grid.fit(X_train, y_train)
    best_model = grid.best_estimator_
    print("XGBoost classifier:")
    print("Best params:", grid.best_params_)

    scoring = {
        'accuracy': 'accuracy',
        'precision': 'precision',
        'recall': 'recall',
        'f1': 'f1',
        'roc_auc': 'roc_auc'
    }

    cv_res = cross_validate(
        best_model, X_train, y_train,
        cv=inner_cv, scoring=scoring, n_jobs=-1, return_train_score=False
    )

    def show(name):
        vals = cv_res[f'test_{name}']
        print(f"CV {name:<9}: {vals.mean():.4f} ± {vals.std():.4f}")

    print("\n=== Repeated Stratified 5×2 CV on TRAIN ===")
    for m in ['accuracy', 'precision', 'recall', 'f1', 'roc_auc']:
        show(m)

    y_pred = best_model.predict(X_test)
    y_prob = best_model.predict_proba(X_test)[:, 1]

    acc = accuracy_score(y_test, y_pred)
    prec = precision_score(y_test, y_pred)
    rec = recall_score(y_test, y_pred)

```

```

f1 = f1_score(y_test, y_pred)
auc = roc_auc_score(y_test, y_prob)

print("\n=== Final Performance on HOLD-OUT TEST ===")
print(f"Accuracy : {acc:.4f}")
print(f"Precision: {prec:.4f}")
print(f"Recall    : {rec:.4f}")
print(f"F1 Score  : {f1:.4f}")
print(f"ROC_AUC   : {auc:.4f}")

print("\nClassification Report on Test Set:\n", classification_report(y_test,
↪y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))

feat_imp = pd.DataFrame({
    "Feature": X_train.columns,
    "Importance": best_model.feature_importances_
}).sort_values(by="Importance", ascending=False)
print("\nTop feature importances:\n", feat_imp)

XGBoost_md1 = best_model

```

XGBoost classifier:

Best params: {'colsample_bytree': 0.7, 'learning_rate': 0.1, 'max_depth': 7, 'min_child_weight': 1, 'n_estimators': 200, 'subsample': 0.9}

=== Repeated Stratified 5×2 CV on TRAIN ===

CV accuracy : 0.8117 ± 0.0621
 CV precision: 0.7983 ± 0.0856
 CV recall : 0.7681 ± 0.0716
 CV f1 : 0.7804 ± 0.0660
 CV roc_auc : 0.8674 ± 0.0541

=== Final Performance on HOLD-OUT TEST ===

Accuracy : 0.8676
 Precision: 0.8621
 Recall : 0.8333
 F1 Score : 0.8475
 ROC_AUC : 0.9123

Classification Report on Test Set:

	precision	recall	f1-score	support
0	0.87	0.89	0.88	38
1	0.86	0.83	0.85	30
accuracy			0.87	68

macro avg	0.87	0.86	0.87	68
weighted avg	0.87	0.87	0.87	68

Confusion Matrix:

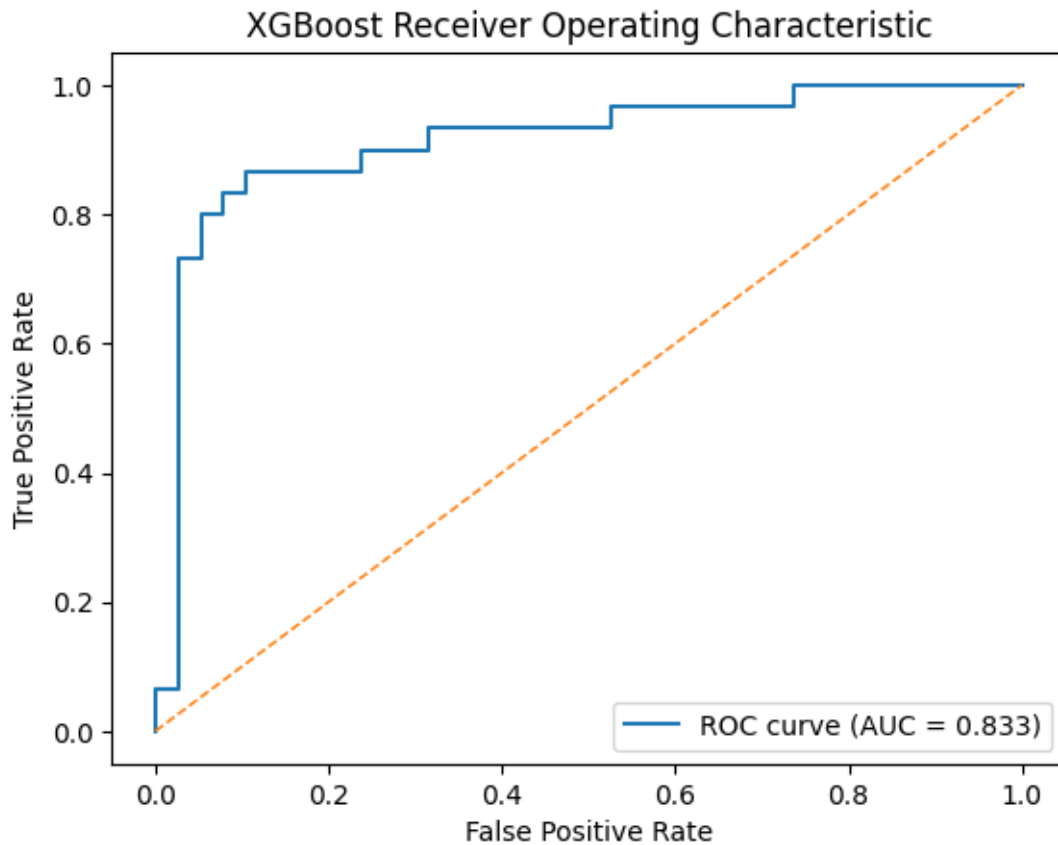
```
[[34  4]
 [ 5 25]]
```

Top feature importances:

	Feature	Importance
0	M/F	0.263357
2	EDUC	0.227721
4	eTIV	0.147784
5	nWBV	0.137234
3	SES	0.121665
1	Age	0.102240

```
[ ]: fpr, tpr, _ = roc_curve(y_test, y_prob)
      roc_auc = recall_score(y_test, y_pred)

      plt.figure()
      plt.plot(fpr, tpr, label=f'ROC curve (AUC = {roc_auc:.3f})')
      plt.plot([0, 1], [0, 1], linestyle='--', linewidth=1)
      plt.xlabel('False Positive Rate')
      plt.ylabel('True Positive Rate')
      plt.title('XGBoost Receiver Operating Characteristic')
      plt.legend(loc='lower right')
      plt.show()
```




```
[ ]: cm = confusion_matrix(y_test, y_pred)
tn, fp, fn, tp = cm.ravel()

labels = np.array([
    [f"TN = {tn}", f"FP = {fp}"],
    [f"FN = {fn}", f"TP = {tp}"]
])

fig, ax = plt.subplots()
ax.matshow(cm, cmap=plt.cm.Blues, alpha=0.3)
for i in range(2):
    for j in range(2):
        ax.text(j, i, labels[i, j],
                va='center', ha='center', fontsize=12, fontweight='bold')

plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix with Labels")
plt.show()
```



```
[ ]: # Find all FN indices in the full test set
FN_all = (~y_pred) & (y_test == 1)
FN_indices = y_test[FN_all].index
print("False Negative indices:", FN_indices)

# Find all FP indices in the full test set
FP_all = (y_pred) & (y_test == 0)
FP_indices = y_test[FP_all].index
print("False Positive indices:", FP_indices)

# Find all TP indices in the full test set
TP_all = (y_pred) & (y_test == 1)
TP_indices = y_test[TP_all].index
print("True Positive indices:", TP_indices)

# Find all TN indices in the full test set
TN_all = (~y_pred) & (y_test == 0)
TN_indices = y_test[TN_all].index
print("True Negative indices:", TN_indices)
```

```
False Negative indices: Index([172, 52, 300, 299, 94], dtype='int64')
False Positive indices: Index([146, 198, 130, 64], dtype='int64')
True Positive indices: Index([124, 332, 250, 317, 154, 25, 90, 106, 285, 87,
215, 127, 3, 239,
162, 345, 72, 39, 89, 51, 88, 16, 329, 365, 275],
dtype='int64')
True Negative indices: Index([ 84, 122, 311, 48, 336, 213, 9, 210, 167, 113,
85, 363, 66, 5,
153, 291, 370, 158, 102, 243, 132, 306, 337, 196, 351, 179, 199, 309,
197, 362, 7, 209, 333, 96],
dtype='int64')
```

```
[ ]: FN_sample_test_idx = X_test.index.get_indexer_for([172, 52, 300, 299, 94])
FP_sample_test_idx = X_test.index.get_indexer_for([146, 198, 130, 64])
TP_sample_test_idx = X_test.index.get_indexer_for([124, 332, 250, 317, 154,
↪25, 90, 106, 285, 87, 215, 127, 3, 239,
162, 345, 72, 39, 89, 51, 88, 16, 329, 365, 275])
TN_sample_test_idx = X_test.index.get_indexer_for([ 84, 122, 311, 48, 336,
↪213, 9, 210, 167, 113, 85, 363, 66, 5,
153, 291, 370, 158, 102, 243, 132, 306, 337, 196, 351, 179, 199, 309,
197, 362, 7, 209, 333, 96])
```

```
[ ]: print("FN_sample_test_idx: ", FN_sample_test_idx)
print("FP_sample_test_idx: ", FP_sample_test_idx)
print("TP_sample_test_idx: ", TP_sample_test_idx)
print("TN_sample_test_idx: ", TN_sample_test_idx)
```

```
FN_sample_test_idx: [16 30 31 37 61]
FP_sample_test_idx: [21 42 50 60]
TP_sample_test_idx: [ 1  2  6  7  9 11 13 14 20 22 28 29 33 35 36 39 40 41 47
49 55 56 64 65
67]
TN_sample_test_idx: [ 0  3  4  5  8 10 12 15 17 18 19 23 24 25 26 27 32 34 38
43 44 45 46 48
51 52 53 54 57 58 59 62 63 66]
```

```
[ ]: XGB_explainer = shap.Explainer(XGBoost_md1)
XGB_shap = XGB_explainer(X_test)
print(type(XGB_explainer))
```


```
<class 'shap.explainers._tree.TreeExplainer'>
```

```
[ ]: print("Values dimensions: %s" % (XGB_shap.values.shape,))
print("Data dimensions: %s" % (XGB_shap.data.shape,))
```


```
Values dimensions: (68, 6)
Data dimensions: (68, 6)
```



```
[ ]: sb.reset_orig()
shap.plots.bar(XGB_shap)
```





```
[ ]: shap.plots.beeswarm(XGB_shap)
```



```
[ ]: mask_f = (X_test['M/F'].values == 1)
mask_m = (X_test['M/F'].values == 0)


shap.plots.beeswarm(XGB_shap[mask_f], show=True)
shap.plots.beeswarm(XGB_shap[mask_m], show=True)
```



From the shap plot above, we can see that M/F(M=0, F=1) female and older age tend to push class to nondemented side.

```
[ ]: # Compute SHAP interaction values
XGB_shap_interaction_values = XGB_explainer.shap_interaction_values(X_test)

# Visualize pairwise interactions (summary plot)
shap.summary_plot(XGB_shap_interaction_values, X_test)
```



```
[ ]: interaction_values = XGB_shap_interaction_values
features = X_test.columns

mean_abs_interactions = np.abs(interaction_values).mean(axis=0)

# Absolute mean measures how strong the interaction is, regardless of direction.

interaction_df = pd.DataFrame(mean_abs_interactions, index=features,
                               columns=features)
interaction_df = interaction_df.where(np.triu(np.ones(interaction_df.shape),
                               k=1).astype(bool))

interaction_df = interaction_df.stack().reset_index()
interaction_df.columns = ['Feature 1', 'Feature 2', 'Mean |Interaction Value|']
interaction_df = interaction_df.sort_values(by='Mean |Interaction Value|',
                               ascending=False)
```

```
print(interaction_df.head(10))
```

	Feature 1	Feature 2	Mean Interaction Value
10	EDUC	eTIV	0.220816
14	eTIV	nWBV	0.212721
11	EDUC	nWBV	0.188962
13	SES	nWBV	0.167425
7	Age	eTIV	0.163801
3	M/F	eTIV	0.142362
8	Age	nWBV	0.137670
1	M/F	EDUC	0.107418
6	Age	SES	0.099767
9	EDUC	SES	0.098318

```
[ ]: # Age and Sex interaction
# Signed mean measures the average directional effect of the interaction
# - positive means on average Age × M/F pushes toward demented,
#- negative means pushes toward nondemented.


age_idx = X_test.columns.get_loc('Age')
gender_idx = X_test.columns.get_loc('M/F')

interaction_values = XGB_shap_interaction_values[age_idx][:, gender_idx]
print("Mean interaction Age × Gender =", interaction_values.mean())
```

Mean interaction Age × Gender = 0.206249

Positive mean SHAP interaction age * gender value(0.20161696) represents gender from male to female(0->1) will improve age's contribution to predict the probability of dementia.

```
[ ]: shap.dependence_plot(
    ('Age', 'M/F'),
    XGB_shap_interaction_values,
    X_test
)
```



SHAP interaction value for Age and M/F: Represents the contribution of the interaction between age and gender to predicting the probability of dementia.

Observation:

We observed a cross-over pattern (points of different colors intersect across age ranges) in the SHAP interaction value for Age and M/F (red: female; blue: male)


- Younger ages (60-70): More red points are positive, more blue points are negative → women tend to push toward demented group; men tend to push toward nondemented group
- Mid-range ages (around 75 years): Interaction values are close to 0 → Gender has little to no influence on the effect of age.
- Older ages (80+ years): Blue points (male) are mostly positive, red points (female) are mostly negative → For older men, increasing age tends to push toward dementia; for older women, it tends to reduce the risk.
- Direction reversal: Across different age ranges, the direction of gender's influence on the age effect flips. This is one of the main reasons why, in our earlier global SHAP analysis, the effects of Age and Sex appeared contrary to expectations — the model has learned locally reversed patterns.

Possible reasons: * Gender proportions differ significantly across age groups in the dataset.


- In the oldest age group, more surviving women are healthy, so “older women” are more likely to be nondemented in the data, leading the model to learn a reversed association.

- The number of male samples is relatively small (M 147: F 189), and for certain age ranges, limited data leads the model to fit local patterns.

```
[ ]: # FN -31
expected_value = XGB_explainer.expected_value
shap.decision_plot(expected_value, XGB_shap.values[31], X_test.iloc[31],
highlight=0)
```



```
[ ]: # FN -16
expected_value = XGB_explainer.expected_value
shap.decision_plot(expected_value, XGB_shap.values[16], X_test.iloc[16],
highlight=0)
```



```
[ ]: # FN -31
shap.initjs()
expected_value = XGB_explainer.expected_value
shap.force_plot(expected_value, XGB_shap.values[31], X_test.iloc[31])
```

<IPython.core.display.HTML object>

```
[ ]: <shap.plots._force.AdditiveForceVisualizer at 0x7f38fc56a690>
```

```
[ ]: # FN -16
shap.initjs()
expected_value = XGB_explainer.expected_value
shap.force_plot(expected_value, XGB_shap.values[16], X_test.iloc[16])
```

<IPython.core.display.HTML object>

```
[ ]: <shap.plots._force.AdditiveForceVisualizer at 0x7f39207ec610>
```

```
[ ]: lime_XGB_explainer = lime.lime_tabular.LimeTabularExplainer(X_test.values,
                                                                feature_names=X_test.columns,
                                                                class_names=['Nondemented', 'Demented'])
```

```
[ ]: # FN case -31
lime_XGB_explainer.explain_instance(X_test.iloc[31].values, \
                                   XGBoost_md1.predict_proba, \
                                   num_features=6). \
```

```
show_in_notebook(predict_proba=True)
```

<IPython.core.display.HTML object>

```
[ ]: # FN case-16
lime_XGB_explainer.explain_instance(X_test.iloc[16].values,\
                                   XGBoost_mdl.predict_proba,\
                                   num_features=6).\
                                   show_in_notebook(predict_proba=True)
```

<IPython.core.display.HTML object>

```
[ ]: # store all fn_results
fn_results = []
feature_counter = Counter()

FN_indices = [16, 30, 31, 37, 61]

for fn_idx in FN_indices:
    instance_values = X_test.iloc[fn_idx].values

    exp = lime_XGB_explainer.explain_instance(
        instance_values,
        XGBoost_mdl.predict_proba,
        num_features=6
    )

    exp_list = exp.as_list()

    pushed_non = [f for f, w in exp_list if w < 0]
    pushed_dem = [f for f, w in exp_list if w > 0]

    fn_results.append({
        'Index': fn_idx,
        'Pushed_Nondemented': pushed_non,
        'Pushed_Demented': pushed_dem
    })

    feature_counter.update(pushed_non)

fn_df = pd.DataFrame(fn_results)

top_causes = pd.DataFrame(feature_counter.most_common(), columns=['Feature', 'Count'])

print(fn_df.head())
```



```
print("\n=== False Negative Feature Frequency ===")
print(top_causes)
```

	Index	Pushed_Nondemented \
0	16	[0.00 < M/F <= 1.00, SES > 3.25, 1491.50 < eTI...
1	30	[EDUC > 16.25, SES <= 2.00]
2	31	[eTIV > 1669.00, 0.00 < M/F <= 1.00, Age > 80...
3	37	[0.00 < M/F <= 1.00, eTIV > 1669.00, 75.00 < A...
4	61	[0.00 < M/F <= 1.00, Age > 80.25, 0.73 < nWBV ...

	Pushed_Demented
0	[EDUC <= 12.00, Age <= 71.00, 0.70 < nWBV <= 0...
1	[M/F <= 0.00, Age <= 71.00, 0.70 < nWBV <= 0.7...
2	[nWBV <= 0.70, 12.00 < EDUC <= 15.00]
3	[nWBV <= 0.70, 12.00 < EDUC <= 15.00]
4	[1391.00 < eTIV <= 1491.50, 12.00 < EDUC <= 15...

=== False Negative Feature Frequency ===

	Feature	Count
0	0.00 < M/F <= 1.00	4
1	SES <= 2.00	4
2	eTIV > 1669.00	2
3	Age > 80.25	2
4	SES > 3.25	1
5	1491.50 < eTIV <= 1669.00	1
6	EDUC > 16.25	1
7	75.00 < Age <= 80.25	1
8	0.73 < nWBV <= 0.76	1

7.2.2 LightGBM

```
[ ]: # Initialize a LightGBM classifier
LightGBM= LGBMClassifier(random_state=42, verbosity=-1)

# Define the parameter grid for Grid
LightGBM_param_dist = {
    'n_estimators': [100, 200],
    'max_depth': [3, 5, 7],
    'learning_rate': [0.001, 0.01, 0.1],
    'num_leaves': [31, 50, 70],
}

grid_search = GridSearchCV(LightGBM, LightGBM_param_dist, cv=5,
    scoring='recall', n_jobs=-1)
grid_search.fit(X_train, y_train)
print("Best hyperparameters found by GridSearchCV:")
print(grid_search.best_params_)
```

```
LightGBM_md1 = grid_search.best_estimator_
y_pred = LightGBM_md1.predict(X_test)
```

Best hyperparameters found by GridSearchCV:

```
{'learning_rate': 0.1, 'max_depth': 7, 'n_estimators': 200, 'num_leaves': 31}
```

LightGBM Best hyperparameters found by GridSearchCV: {'learning_rate': 0.1, 'max_depth': 7, 'n_estimators': 200, 'num_leaves': 31}

```
[ ]: print("LightGBM: ")
print("\nClassification Report on Test Set:")
print(classification_report(y_test, y_pred))
print("-----")

print("Accuracy:", accuracy_score(y_test, y_pred))
print("Precision:", precision_score(y_test, y_pred))
print("Recall:", recall_score(y_test, y_pred))
print("F1 Score:", f1_score(y_test, y_pred))
print("ROC_AUC:", roc_auc_score(y_test, y_pred))
print("-----")

# Get feature importances
feature_importances = pd.DataFrame({
    "Feature": X_test.columns,
    "Importance": LightGBM_md1.feature_importances_
})

feature_importances = feature_importances.sort_values(by="Importance",
↪ascending=False)
feature_importances
```

LightGBM:

Classification Report on Test Set:

	precision	recall	f1-score	support
0	0.89	0.87	0.88	38
1	0.84	0.87	0.85	30
accuracy			0.87	68
macro avg	0.87	0.87	0.87	68
weighted avg	0.87	0.87	0.87	68

Accuracy: 0.8676470588235294

Precision: 0.8387096774193549

Recall: 0.8666666666666667

F1 Score: 0.8524590163934426
ROC_AUC: 0.8675438596491228

```
[ ]: Feature Importance
4    eTIV          610
5    nWBV          478
1     Age          332
2    EDUC          215
3     SES          177
0    M/F           68
```

```
[ ]: LightGBM = LGBMClassifier(random_state=42, verbosity=-1)
```

```
param_grid = {
    'n_estimators': [100, 200],
    'max_depth': [3, 5, 7],
    'learning_rate': [0.001, 0.01, 0.1],
    'num_leaves': [31, 50, 70],
}

inner_cv = RepeatedStratifiedKFold(
    n_splits=5, n_repeats=2, random_state=42
)

grid = GridSearchCV(
    estimator=LightGBM,
    param_grid=param_grid,
    scoring='f1',
    cv=inner_cv,
    n_jobs=-1,
    verbose=0
)

grid.fit(X_train, y_train)
best_model = grid.best_estimator_
print("LightGBM classifier:")
print("Best params:", grid.best_params_)

scoring = {
    'accuracy': 'accuracy',
    'precision': 'precision',
    'recall': 'recall',
    'f1': 'f1',
    'roc_auc': 'roc_auc'
}

cv_res = cross_validate(
```

```

best_model, X_train, y_train,
cv=inner_cv, scoring=scoring, n_jobs=-1, return_train_score=False
)

def show(name):
    vals = cv_res[f'test_{name}']
    print(f"CV {name:<9}: {vals.mean():.4f} ± {vals.std():.4f}")

print("\n=== Repeated Stratified 5×2 CV on TRAIN ===")
for m in ['accuracy', 'precision', 'recall', 'f1', 'roc_auc']:
    show(m)

y_pred = best_model.predict(X_test)
y_prob = best_model.predict_proba(X_test)[: , 1]

acc = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred)
rec = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
auc = roc_auc_score(y_test, y_prob)

print("\n=== Final Performance on HOLD-OUT TEST ===")
print(f"Accuracy : {acc:.4f}")
print(f"Precision: {prec:.4f}")
print(f"Recall   : {rec:.4f}")
print(f"F1 Score  : {f1:.4f}")
print(f"ROC_AUC   : {auc:.4f}")

print("\nClassification Report on Test Set:\n", classification_report(y_test,
↪y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))

feat_imp = pd.DataFrame({
    "Feature": X_train.columns,
    "Importance": best_model.feature_importances_
}).sort_values(by="Importance", ascending=False)
print("\nTop feature importances:\n", feat_imp)

LightGBM_md1 = best_model

```

LightGBM classifier:

Best params: {'learning_rate': 0.1, 'max_depth': 7, 'n_estimators': 200, 'num_leaves': 31}

=== Repeated Stratified 5×2 CV on TRAIN ===

CV accuracy : 0.7969 ± 0.0714

CV precision: 0.7726 ± 0.0991

```
CV recall    : 0.7721 ± 0.0675
CV f1        : 0.7692 ± 0.0711
CV roc_auc   : 0.8443 ± 0.0613
```

=== Final Performance on HOLD-OUT TEST ===

```
Accuracy : 0.8676
Precision: 0.8387
Recall    : 0.8667
F1 Score  : 0.8525
ROC_AUC   : 0.9123
```

Classification Report on Test Set:

	precision	recall	f1-score	support
0	0.89	0.87	0.88	38
1	0.84	0.87	0.85	30
accuracy			0.87	68
macro avg	0.87	0.87	0.87	68
weighted avg	0.87	0.87	0.87	68

Confusion Matrix:

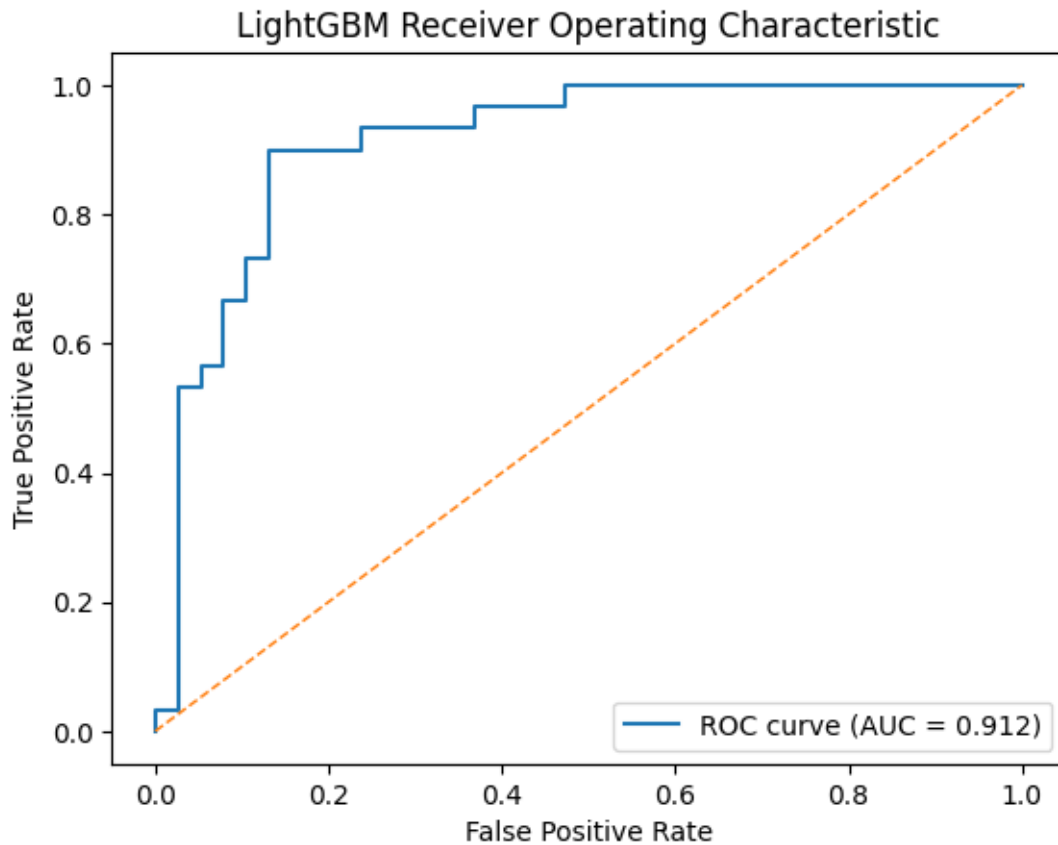
```
[[33  5]
 [ 4 26]]
```

Top feature importances:

	Feature	Importance
4	eTIV	610
5	nWBV	478
1	Age	332
2	EDUC	215
3	SES	177
0	M/F	68

```
[ ]: fpr, tpr, _ = roc_curve(y_test, y_prob)
     roc_auc = roc_auc_score(y_test, y_prob)

     plt.figure()
     plt.plot(fpr, tpr, label=f'ROC curve (AUC = {roc_auc:.3f})')
     plt.plot([0, 1], [0, 1], linestyle='--', linewidth=1)
     plt.xlabel('False Positive Rate')
     plt.ylabel('True Positive Rate')
     plt.title('LightGBM Receiver Operating Characteristic')
     plt.legend(loc='lower right')
     plt.show()
```




```
[ ]: cm = confusion_matrix(y_test, y_pred)
tn, fp, fn, tp = cm.ravel()

labels = np.array([
    [f"TN = {tn}", f"FP = {fp}"],
    [f"FN = {fn}", f"TP = {tp}"]
])

fig, ax = plt.subplots()
ax.matshow(cm, cmap=plt.cm.Blues, alpha=0.3)
for i in range(2):
    for j in range(2):
        ax.text(j, i, labels[i, j],
                va='center', ha='center', fontsize=12, fontweight='bold')

plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix with Labels")
plt.show()
```



```
[ ]: # Find all FN indices in the full test set
FN_all = (~y_pred) & (y_test == 1)
FN_indices = y_test[FN_all].index
print("False Negative indices:", FN_indices)
```

```
False Negative indices: Index([172, 300, 299, 94], dtype='int64')
```

```
[ ]: FN_sample_test_idx = X_test.index.get_indexer_for([172, 300, 299, 94])
FN_sample_test_idx
```

```
[ ]: array([16, 31, 37, 61])
```

```
[ ]: LGB_explainer = shap.Explainer(LightGBM_md1)
LGB_shap = LGB_explainer(X_test)
print(type(LGB_explainer))
```


```
<class 'shap.explainers._tree.TreeExplainer'>
```

```
[ ]: print("Values dimensions: %s" % (LGB_shap.values.shape,))
print("Data dimensions: %s" % (LGB_shap.data.shape,))
```

Values dimensions: (68, 6)

Data dimensions: (68, 6)


```
[ ]: sb.reset_orig()  
shap.plots.bar(LGB_shap)
```




```
[ ]: shap.plots.beeswarm(LGB_shap)
```




```
[ ]: # LGB-FN-16
sb.reset_orig()
expected_value = LGB_explainer.expected_value
shap.decision_plot(expected_value, LGB_shap.values[16], X_test.iloc[16],
highlight=0)
```



```
[ ]: # LGB-FN-31
expected_value = LGB_explainer.expected_value
shap.decision_plot(expected_value, LGB_shap.values[31], X_test.iloc[31],
highlight=0)
```



```
[ ]: # LGB-FN-16
shap.initjs()
expected_value = LGB_explainer.expected_value
shap.force_plot(expected_value, LGB_shap.values[16], X_test.iloc[16])
```

<IPython.core.display.HTML object>

```
[ ]: <shap.plots._force.AdditiveForceVisualizer at 0x7f38fdb96650>
```

```
[ ]: # LGB-FN-31
shap.initjs()
expected_value = LGB_explainer.expected_value
shap.force_plot(expected_value, LGB_shap.values[31], X_test.iloc[31])
```

<IPython.core.display.HTML object>

```
[ ]: <shap.plots._force.AdditiveForceVisualizer at 0x7f391eb30f50>
```

```
[ ]: lime_LGB_explainer = lime.lime_tabular.LimeTabularExplainer(X_test.values,
                                                                feature_names=X_test.columns,
                                                                class_names=['Nondemented', 'Demented'])
```

```
[ ]: # LGB-FN-16
lime_LGB_explainer.explain_instance(X_test.iloc[16].values, \
                                   LightGBM_md1.predict_proba, \
                                   num_features=6). \
```

```
show_in_notebook(predict_proba=True)
```

```
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739:  
UserWarning: X does not have valid feature names, but LGBMClassifier was fitted  
with feature names
```

```
warnings.warn(  

```

```
<IPython.core.display.HTML object>
```

```
[ ]: # LGB-FN-31  
lime_LGB_explainer.explain_instance(X_test.iloc[31].values,\  
                                   LightGBM_mdl.predict_proba,\  
                                   num_features=6).\  
                                   show_in_notebook(predict_proba=True)
```

```
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739:  
UserWarning: X does not have valid feature names, but LGBMClassifier was fitted  
with feature names
```

```
warnings.warn(  

```

```
<IPython.core.display.HTML object>
```

```
[ ]: # store all fn_results  
fn_results = []  
feature_counter = Counter()  
  
LGB_FN_indices = [16, 31, 37, 61]  
  
for fn_idx in LGB_FN_indices:  
    LGB_instance_values = X_test.iloc[fn_idx].values  
  
    exp = lime_LGB_explainer.explain_instance(  
        LGB_instance_values,  
        LightGBM_mdl.predict_proba,  
        num_features=6  
    )  
  
    exp_list = exp.as_list()  
  
    pushed_non = [f for f, w in exp_list if w < 0]  
    pushed_dem = [f for f, w in exp_list if w > 0]  
  
    fn_results.append({  
        'Index': fn_idx,  
        'Pushed_Nondemented': pushed_non,  
        'Pushed_Demented': pushed_dem  
    })
```

```

feature_counter.update(pushed_non)

fn_df = pd.DataFrame(fn_results)

top_causes = pd.DataFrame(feature_counter.most_common(), columns=['Feature', 'Count'])

print(fn_df.head())

print("\n=== LGB False Negative Feature Frequency ===")
print(top_causes)

```

```

/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739:
UserWarning: X does not have valid feature names, but LGBMClassifier was fitted
with feature names
  warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739:
UserWarning: X does not have valid feature names, but LGBMClassifier was fitted
with feature names
  warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739:
UserWarning: X does not have valid feature names, but LGBMClassifier was fitted
with feature names
  warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739:
UserWarning: X does not have valid feature names, but LGBMClassifier was fitted
with feature names
  warnings.warn(

```

	Index	Pushed_Nondemented \
0	16	[0.00 < M/F <= 1.00, SES > 3.25, 1491.50 < eTIV <= 1669.00]
1	31	[eTIV > 1669.00, 0.00 < M/F <= 1.00, Age > 80.25]
2	37	[0.00 < M/F <= 1.00, eTIV > 1669.00, 75.00 < Age <= 80.25]
3	61	[0.00 < M/F <= 1.00, Age > 80.25, 0.73 < nWBV <= 0.70]

	Pushed_Demented
0	[EDUC <= 12.00, Age <= 71.00, 0.70 < nWBV <= 0.70]
1	[nWBV <= 0.70, 12.00 < EDUC <= 15.00]
2	[nWBV <= 0.70, 12.00 < EDUC <= 15.00]
3	[1391.00 < eTIV <= 1491.50, 12.00 < EDUC <= 15.00]

=== LGB False Negative Feature Frequency ===

	Feature	Count
0	0.00 < M/F <= 1.00	4
1	SES <= 2.00	3
2	eTIV > 1669.00	2
3	Age > 80.25	2
4	SES > 3.25	1

```

5  1491.50 < eTIV <= 1669.00      1
6      75.00 < Age <= 80.25      1
7      0.73 < nWBV <= 0.76      1

```

```
###CatBoost
```

```

[ ]: # Initialize a CatBoost classifier
CatBoost = CatBoostClassifier(random_state=42, verbose = False)

# Define the parameter grid for Grid
CatBoost_param_dist = {
    'min_data_in_leaf': [20, 40, 60],
    'rsm': [0.7, 0.8, 1.0],
    'iterations': [100, 200],
    'depth': [3, 5, 7],
    'learning_rate': [0.001, 0.01, 0.1],
    'l2_leaf_reg': [1, 3, 5],
    'bagging_temperature': [0, 0.5, 1.0],
}

grid_search = GridSearchCV(CatBoost, CatBoost_param_dist, cv=5,
    ↪scoring='recall', n_jobs=-1)
grid_search.fit(X_train, y_train)
print("Best hyperparameters found by GridSearchCV:")
print(grid_search.best_params_)

CatBoost_md1 = grid_search.best_estimator_
y_pred = CatBoost_md1.predict(X_test)

```

Best hyperparameters found by GridSearchCV:

```
{'bagging_temperature': 0, 'depth': 7, 'iterations': 200, 'l2_leaf_reg': 5,
'learning_rate': 0.1, 'min_data_in_leaf': 20, 'rsm': 0.8}
```

CatBoost Best hyperparameters found by GridSearchCV:

```
{'bagging_temperature': 0, 'depth': 7, 'iterations': 200, 'l2_leaf_reg': 5, 'learning_rate': 0.1,
'min_data_in_leaf': 20, 'rsm': 0.8}
```

```

[ ]: print("CatBoost: ")
print("\nClassification Report on Test Set:")
print(classification_report(y_test, y_pred))
print("-----")

print("Accuracy:", accuracy_score(y_test, y_pred))
print("Precision:", precision_score(y_test, y_pred))
print("Recall:", recall_score(y_test, y_pred))
print("F1 Score:", f1_score(y_test, y_pred))
print("ROC_AUC:", roc_auc_score(y_test, y_pred))
print("-----")

```

```
# Get feature importances
feature_importances = pd.DataFrame({
    "Feature": X_test.columns,
    "Importance": CatBoost_md1.feature_importances_
})

feature_importances = feature_importances.sort_values(by="Importance",
↪ascending=False)
feature_importances
```

CatBoost:

Classification Report on Test Set:

	precision	recall	f1-score	support
0	0.85	0.92	0.89	38
1	0.89	0.80	0.84	30
accuracy			0.87	68
macro avg	0.87	0.86	0.86	68
weighted avg	0.87	0.87	0.87	68

Accuracy: 0.8676470588235294

Precision: 0.8888888888888888

Recall: 0.8

F1 Score: 0.8421052631578947

ROC_AUC: 0.8605263157894737

```
[ ]: Feature Importance
4    eTIV    22.606796
5    nWBV    21.504276
2    EDUC    17.237605
1    Age     16.420387
3    SES     14.022414
0    M/F     8.208523
```

```
[ ]: CatBoost = CatBoostClassifier(random_state=42)
```

```
param_grid = {
    'min_data_in_leaf': [20, 40, 60],
    'rsm': [0.7, 0.8, 1.0],
    'iterations': [100, 200],
    'depth': [3, 5, 7],
    'learning_rate': [0.001, 0.01, 0.1],
```

```

    'l2_leaf_reg': [1, 3, 5],
    'bagging_temperature': [0, 0.5, 1.0],
}

inner_cv = RepeatedStratifiedKFold(
    n_splits=5, n_repeats=2, random_state=42
)

grid = GridSearchCV(
    estimator=CatBoost,
    param_grid=param_grid,
    scoring='f1',
    cv=inner_cv,
    n_jobs=-1,
    verbose=0
)

grid.fit(X_train, y_train)
best_model = grid.best_estimator_
print("CatBoostClassifier:")
print("Best params:", grid.best_params_)

scoring = {
    'accuracy': 'accuracy',
    'precision': 'precision',
    'recall': 'recall',
    'f1': 'f1',
    'roc_auc': 'roc_auc'
}

cv_res = cross_validate(
    best_model, X_train, y_train,
    cv=inner_cv, scoring=scoring, n_jobs=-1, return_train_score=False
)

def show(name):
    vals = cv_res[f'test_{name}']
    print(f"CV {name:<9}: {vals.mean():.4f} ± {vals.std():.4f}")

print("\n=== Repeated Stratified 5×2 CV on TRAIN ===")
for m in ['accuracy', 'precision', 'recall', 'f1', 'roc_auc']:
    show(m)

y_pred = best_model.predict(X_test)
y_prob = best_model.predict_proba(X_test)[:, 1]

acc = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred)

```

```

rec = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
auc = roc_auc_score(y_test, y_prob)

print("\n=== Final Performance on HOLD-OUT TEST ===")
print(f"Accuracy : {acc:.4f}")
print(f"Precision: {prec:.4f}")
print(f"Recall    : {rec:.4f}")
print(f"F1 Score  : {f1:.4f}")
print(f"ROC_AUC   : {auc:.4f}")

print("\nClassification Report on Test Set:\n", classification_report(y_test,
↪y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))

feat_imp = pd.DataFrame({
    "Feature": X_train.columns,
    "Importance": best_model.feature_importances_
}).sort_values(by="Importance", ascending=False)
print("\nTop feature importances:\n", feat_imp)

CatBoost_md1 = best_model

```

0:	learn: 0.6276889	total: 47.7ms	remaining: 9.5s
1:	learn: 0.5745076	total: 49ms	remaining: 4.85s
2:	learn: 0.5643154	total: 49.6ms	remaining: 3.26s
3:	learn: 0.5272631	total: 50.8ms	remaining: 2.49s
4:	learn: 0.4983803	total: 51.9ms	remaining: 2.02s
5:	learn: 0.4609295	total: 53.1ms	remaining: 1.72s
6:	learn: 0.4404177	total: 54ms	remaining: 1.49s
7:	learn: 0.4183808	total: 55.1ms	remaining: 1.32s
8:	learn: 0.3995695	total: 56.3ms	remaining: 1.2s
9:	learn: 0.3771613	total: 57.8ms	remaining: 1.1s
10:	learn: 0.3653814	total: 59.2ms	remaining: 1.02s
11:	learn: 0.3591362	total: 60.3ms	remaining: 945ms
12:	learn: 0.3424529	total: 61.4ms	remaining: 883ms
13:	learn: 0.3281506	total: 62.3ms	remaining: 828ms
14:	learn: 0.3188232	total: 63.4ms	remaining: 782ms
15:	learn: 0.3085403	total: 64.3ms	remaining: 739ms
16:	learn: 0.2969568	total: 65.3ms	remaining: 703ms
17:	learn: 0.2809187	total: 66.4ms	remaining: 672ms
18:	learn: 0.2730252	total: 67.4ms	remaining: 642ms
19:	learn: 0.2641290	total: 68.5ms	remaining: 617ms
20:	learn: 0.2565324	total: 69.5ms	remaining: 592ms
21:	learn: 0.2436529	total: 70.6ms	remaining: 571ms
22:	learn: 0.2405098	total: 71.4ms	remaining: 549ms
23:	learn: 0.2348255	total: 72.4ms	remaining: 531ms

24:	learn: 0.2293342	total: 73.5ms	remaining: 514ms
25:	learn: 0.2215201	total: 74.4ms	remaining: 498ms
26:	learn: 0.2125021	total: 75.5ms	remaining: 484ms
27:	learn: 0.2037552	total: 76.5ms	remaining: 470ms
28:	learn: 0.1996584	total: 77.4ms	remaining: 456ms
29:	learn: 0.1929344	total: 78.3ms	remaining: 444ms
30:	learn: 0.1845759	total: 79.3ms	remaining: 432ms
31:	learn: 0.1794482	total: 80.4ms	remaining: 422ms
32:	learn: 0.1721649	total: 81.2ms	remaining: 411ms
33:	learn: 0.1670736	total: 82ms	remaining: 401ms
34:	learn: 0.1623246	total: 83.1ms	remaining: 392ms
35:	learn: 0.1572181	total: 84.1ms	remaining: 383ms
36:	learn: 0.1572009	total: 84.5ms	remaining: 372ms
37:	learn: 0.1539729	total: 85.5ms	remaining: 365ms
38:	learn: 0.1535214	total: 86ms	remaining: 355ms
39:	learn: 0.1505825	total: 87ms	remaining: 348ms
40:	learn: 0.1451667	total: 88.1ms	remaining: 342ms
41:	learn: 0.1424512	total: 89.2ms	remaining: 336ms
42:	learn: 0.1386530	total: 90.2ms	remaining: 329ms
43:	learn: 0.1344911	total: 91.2ms	remaining: 324ms
44:	learn: 0.1334826	total: 91.9ms	remaining: 316ms
45:	learn: 0.1312075	total: 92.9ms	remaining: 311ms
46:	learn: 0.1277087	total: 93.9ms	remaining: 306ms
47:	learn: 0.1252132	total: 94.9ms	remaining: 300ms
48:	learn: 0.1224295	total: 95.9ms	remaining: 295ms
49:	learn: 0.1205716	total: 96.9ms	remaining: 291ms
50:	learn: 0.1168845	total: 97.9ms	remaining: 286ms
51:	learn: 0.1135585	total: 99ms	remaining: 282ms
52:	learn: 0.1132123	total: 99.5ms	remaining: 276ms
53:	learn: 0.1100292	total: 100ms	remaining: 272ms
54:	learn: 0.1086760	total: 101ms	remaining: 267ms
55:	learn: 0.1058472	total: 102ms	remaining: 263ms
56:	learn: 0.1002297	total: 103ms	remaining: 259ms
57:	learn: 0.0977531	total: 104ms	remaining: 255ms
58:	learn: 0.0958838	total: 105ms	remaining: 252ms
59:	learn: 0.0935505	total: 106ms	remaining: 248ms
60:	learn: 0.0909632	total: 107ms	remaining: 245ms
61:	learn: 0.0873322	total: 108ms	remaining: 241ms
62:	learn: 0.0847598	total: 109ms	remaining: 238ms
63:	learn: 0.0831981	total: 110ms	remaining: 235ms
64:	learn: 0.0801172	total: 111ms	remaining: 231ms
65:	learn: 0.0778482	total: 113ms	remaining: 229ms
66:	learn: 0.0762141	total: 114ms	remaining: 226ms
67:	learn: 0.0745551	total: 115ms	remaining: 222ms
68:	learn: 0.0733157	total: 116ms	remaining: 220ms
69:	learn: 0.0709646	total: 117ms	remaining: 217ms
70:	learn: 0.0693339	total: 118ms	remaining: 214ms
71:	learn: 0.0676360	total: 119ms	remaining: 211ms

72:	learn: 0.0666511	total: 119ms	remaining: 208ms
73:	learn: 0.0658717	total: 120ms	remaining: 205ms
74:	learn: 0.0642397	total: 122ms	remaining: 203ms
75:	learn: 0.0637420	total: 122ms	remaining: 200ms
76:	learn: 0.0627294	total: 123ms	remaining: 197ms
77:	learn: 0.0614020	total: 124ms	remaining: 195ms
78:	learn: 0.0597003	total: 125ms	remaining: 192ms
79:	learn: 0.0581487	total: 127ms	remaining: 190ms
80:	learn: 0.0576556	total: 128ms	remaining: 187ms
81:	learn: 0.0559929	total: 129ms	remaining: 185ms
82:	learn: 0.0550454	total: 130ms	remaining: 183ms
83:	learn: 0.0540910	total: 131ms	remaining: 180ms
84:	learn: 0.0525559	total: 132ms	remaining: 178ms
85:	learn: 0.0508644	total: 133ms	remaining: 176ms
86:	learn: 0.0494802	total: 134ms	remaining: 174ms
87:	learn: 0.0484818	total: 135ms	remaining: 171ms
88:	learn: 0.0476333	total: 136ms	remaining: 169ms
89:	learn: 0.0466150	total: 137ms	remaining: 167ms
90:	learn: 0.0459159	total: 138ms	remaining: 165ms
91:	learn: 0.0442754	total: 139ms	remaining: 163ms
92:	learn: 0.0439857	total: 139ms	remaining: 160ms
93:	learn: 0.0433327	total: 140ms	remaining: 158ms
94:	learn: 0.0421936	total: 141ms	remaining: 156ms
95:	learn: 0.0412142	total: 142ms	remaining: 154ms
96:	learn: 0.0400072	total: 143ms	remaining: 152ms
97:	learn: 0.0386757	total: 144ms	remaining: 150ms
98:	learn: 0.0377154	total: 145ms	remaining: 148ms
99:	learn: 0.0369504	total: 146ms	remaining: 146ms
100:	learn: 0.0363817	total: 147ms	remaining: 144ms
101:	learn: 0.0357451	total: 148ms	remaining: 142ms
102:	learn: 0.0352235	total: 149ms	remaining: 140ms
103:	learn: 0.0346965	total: 150ms	remaining: 139ms
104:	learn: 0.0344951	total: 151ms	remaining: 137ms
105:	learn: 0.0341423	total: 152ms	remaining: 135ms
106:	learn: 0.0338168	total: 153ms	remaining: 133ms
107:	learn: 0.0331411	total: 154ms	remaining: 131ms
108:	learn: 0.0321822	total: 155ms	remaining: 130ms
109:	learn: 0.0319019	total: 156ms	remaining: 128ms
110:	learn: 0.0314680	total: 157ms	remaining: 126ms
111:	learn: 0.0305845	total: 158ms	remaining: 124ms
112:	learn: 0.0300843	total: 159ms	remaining: 122ms
113:	learn: 0.0296826	total: 160ms	remaining: 121ms
114:	learn: 0.0289181	total: 161ms	remaining: 119ms
115:	learn: 0.0286427	total: 162ms	remaining: 117ms
116:	learn: 0.0284335	total: 163ms	remaining: 116ms
117:	learn: 0.0279876	total: 164ms	remaining: 114ms
118:	learn: 0.0276698	total: 165ms	remaining: 112ms
119:	learn: 0.0272298	total: 166ms	remaining: 110ms

120:	learn: 0.0269485	total: 167ms	remaining: 109ms
121:	learn: 0.0266267	total: 168ms	remaining: 107ms
122:	learn: 0.0262445	total: 169ms	remaining: 106ms
123:	learn: 0.0257281	total: 170ms	remaining: 104ms
124:	learn: 0.0254373	total: 171ms	remaining: 102ms
125:	learn: 0.0250424	total: 172ms	remaining: 101ms
126:	learn: 0.0247979	total: 173ms	remaining: 99.2ms
127:	learn: 0.0241344	total: 174ms	remaining: 97.6ms
128:	learn: 0.0237783	total: 175ms	remaining: 96.1ms
129:	learn: 0.0234832	total: 175ms	remaining: 94.5ms
130:	learn: 0.0233156	total: 176ms	remaining: 93ms
131:	learn: 0.0231681	total: 177ms	remaining: 91.4ms
132:	learn: 0.0228606	total: 179ms	remaining: 90.1ms
133:	learn: 0.0223851	total: 180ms	remaining: 88.5ms
134:	learn: 0.0219639	total: 181ms	remaining: 87ms
135:	learn: 0.0217298	total: 182ms	remaining: 85.6ms
136:	learn: 0.0215157	total: 183ms	remaining: 84.1ms
137:	learn: 0.0212368	total: 184ms	remaining: 82.6ms
138:	learn: 0.0210671	total: 185ms	remaining: 81ms
139:	learn: 0.0206571	total: 185ms	remaining: 79.5ms
140:	learn: 0.0201631	total: 186ms	remaining: 78ms
141:	learn: 0.0197564	total: 187ms	remaining: 76.4ms
142:	learn: 0.0195333	total: 188ms	remaining: 74.9ms
143:	learn: 0.0192705	total: 189ms	remaining: 73.4ms
144:	learn: 0.0190205	total: 190ms	remaining: 72ms
145:	learn: 0.0187692	total: 191ms	remaining: 70.6ms
146:	learn: 0.0184775	total: 192ms	remaining: 69.2ms
147:	learn: 0.0183194	total: 193ms	remaining: 67.8ms
148:	learn: 0.0181327	total: 194ms	remaining: 66.4ms
149:	learn: 0.0180176	total: 195ms	remaining: 65ms
150:	learn: 0.0177636	total: 196ms	remaining: 63.6ms
151:	learn: 0.0175364	total: 197ms	remaining: 62.2ms
152:	learn: 0.0173777	total: 198ms	remaining: 60.9ms
153:	learn: 0.0172088	total: 199ms	remaining: 59.5ms
154:	learn: 0.0169582	total: 200ms	remaining: 58.1ms
155:	learn: 0.0167966	total: 201ms	remaining: 56.7ms
156:	learn: 0.0165438	total: 202ms	remaining: 55.3ms
157:	learn: 0.0163692	total: 203ms	remaining: 54ms
158:	learn: 0.0162543	total: 204ms	remaining: 52.6ms
159:	learn: 0.0160601	total: 205ms	remaining: 51.3ms
160:	learn: 0.0158619	total: 206ms	remaining: 49.9ms
161:	learn: 0.0156578	total: 207ms	remaining: 48.5ms
162:	learn: 0.0155133	total: 208ms	remaining: 47.2ms
163:	learn: 0.0153781	total: 209ms	remaining: 45.9ms
164:	learn: 0.0152523	total: 210ms	remaining: 44.5ms
165:	learn: 0.0150373	total: 211ms	remaining: 43.2ms
166:	learn: 0.0149248	total: 212ms	remaining: 41.9ms
167:	learn: 0.0147364	total: 213ms	remaining: 40.6ms

168:	learn: 0.0146610	total: 214ms	remaining: 39.2ms
169:	learn: 0.0144629	total: 215ms	remaining: 37.9ms
170:	learn: 0.0142976	total: 216ms	remaining: 36.6ms
171:	learn: 0.0141761	total: 217ms	remaining: 35.3ms
172:	learn: 0.0140164	total: 218ms	remaining: 34ms
173:	learn: 0.0138057	total: 219ms	remaining: 32.7ms
174:	learn: 0.0136879	total: 220ms	remaining: 31.5ms
175:	learn: 0.0136131	total: 221ms	remaining: 30.2ms
176:	learn: 0.0134512	total: 222ms	remaining: 28.9ms
177:	learn: 0.0133220	total: 223ms	remaining: 27.6ms
178:	learn: 0.0132624	total: 224ms	remaining: 26.3ms
179:	learn: 0.0131619	total: 225ms	remaining: 25ms
180:	learn: 0.0130354	total: 226ms	remaining: 23.8ms
181:	learn: 0.0129432	total: 227ms	remaining: 22.5ms
182:	learn: 0.0128168	total: 228ms	remaining: 21.2ms
183:	learn: 0.0126758	total: 229ms	remaining: 19.9ms
184:	learn: 0.0125520	total: 230ms	remaining: 18.7ms
185:	learn: 0.0124253	total: 231ms	remaining: 17.4ms
186:	learn: 0.0122516	total: 232ms	remaining: 16.2ms
187:	learn: 0.0121749	total: 233ms	remaining: 14.9ms
188:	learn: 0.0121189	total: 234ms	remaining: 13.6ms
189:	learn: 0.0120047	total: 235ms	remaining: 12.4ms
190:	learn: 0.0119277	total: 236ms	remaining: 11.1ms
191:	learn: 0.0117016	total: 237ms	remaining: 9.89ms
192:	learn: 0.0115816	total: 238ms	remaining: 8.65ms
193:	learn: 0.0114349	total: 240ms	remaining: 7.42ms
194:	learn: 0.0113053	total: 241ms	remaining: 6.17ms
195:	learn: 0.0112673	total: 242ms	remaining: 4.93ms
196:	learn: 0.0111205	total: 243ms	remaining: 3.7ms
197:	learn: 0.0109900	total: 244ms	remaining: 2.47ms
198:	learn: 0.0108626	total: 245ms	remaining: 1.23ms
199:	learn: 0.0107928	total: 246ms	remaining: 0ms

LightGBM classifier:

Best params: {'bagging_temperature': 0, 'depth': 7, 'iterations': 200, 'l2_leaf_reg': 1, 'learning_rate': 0.1, 'min_data_in_leaf': 20, 'rsm': 0.8}

=== Repeated Stratified 5×2 CV on TRAIN ===

CV accuracy : 0.8601 ± 0.0651
 CV precision: 0.8547 ± 0.0891
 CV recall : 0.8236 ± 0.0728
 CV f1 : 0.8370 ± 0.0711
 CV roc_auc : 0.9116 ± 0.0463

=== Final Performance on HOLD-OUT TEST ===

Accuracy : 0.8824
 Precision: 0.9231
 Recall : 0.8000
 F1 Score : 0.8571

ROC_AUC : 0.9605

Classification Report on Test Set:

	precision	recall	f1-score	support
0	0.86	0.95	0.90	38
1	0.92	0.80	0.86	30
accuracy			0.88	68
macro avg	0.89	0.87	0.88	68
weighted avg	0.89	0.88	0.88	68

Confusion Matrix:


```
[[36  2]
 [ 6 24]]
```

Top feature importances:

	Feature	Importance
4	eTIV	22.004060
5	nWBV	21.028797
3	SES	17.430076
2	EDUC	15.873306
1	Age	15.560894
0	M/F	8.102866

```
[ ]: fpr, tpr, _ = roc_curve(y_test, y_prob)
      roc_auc = roc_auc_score(y_test, y_prob)

      plt.figure()
      plt.plot(fpr, tpr, label=f'ROC curve (AUC = {roc_auc:.3f})')
      plt.plot([0, 1], [0, 1], linestyle='--', linewidth=1)
      plt.xlabel('False Positive Rate')
      plt.ylabel('True Positive Rate')
      plt.title('CatBoost Receiver Operating Characteristic')
      plt.legend(loc='lower right')
      plt.show()
```




```
[ ]: cm = confusion_matrix(y_test, y_pred)
tn, fp, fn, tp = cm.ravel()

labels = np.array([
    [f"TN = {tn}", f"FP = {fp}"],
    [f"FN = {fn}", f"TP = {tp}"]
])

fig, ax = plt.subplots()
ax.matshow(cm, cmap=plt.cm.Blues, alpha=0.3)
for i in range(2):
    for j in range(2):
        ax.text(j, i, labels[i, j],
                va='center', ha='center', fontsize=12, fontweight='bold')

plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix with Labels")
plt.show()
```



```
[ ]: # Find all FN indices in the full test set
FN_all = (~y_pred) & (y_test == 1)
FN_indices = y_test[FN_all].index
print("False Negative indices:", FN_indices)
```

```
False Negative indices: Index([332, 52, 300, 299, 51, 94], dtype='int64')
```

```
[ ]: FN_sample_test_idx = X_test.index.get_indexer_for([332, 52, 300, 299, 51, 94])
FN_sample_test_idx
```

```
[ ]: array([ 2, 30, 31, 37, 49, 61])
```

```
[ ]: CB_explainer = shap.Explainer(CatBoost_md1)
CB_shap = CB_explainer(X_test)
print(type(CB_explainer))
```


```
<class 'shap.explainers._tree.TreeExplainer'>
```

```
[ ]: print("Values dimensions: %s" % (CB_shap.values.shape,))
print("Data dimensions: %s" % (CB_shap.data.shape,))
```

Values dimensions: (68, 6)

Data dimensions: (68, 6)

```
[ ]: sb.reset_orig()  
shap.plots.bar(CB_shap)
```



```
[ ]: shap.plots.beeswarm(CB_shap)
```





```
[ ]: print("X_test.iloc[2]: ")
      print(X_test.iloc[2])
      print(y_test.iloc[2], y_pred[2])
      print("-----")
      print("X_test.iloc[31]: ")
      print(X_test.iloc[31])
      print(y_test.iloc[31], y_pred[31])
```


```
X_test.iloc[2]:
M/F      0.000
Age      80.000
EDUC     12.000
SES       3.000
eTIV    1755.000
nWBV      0.696
Name: 332, dtype: float64
1 0
```

```
-----
X_test.iloc[31]:
M/F       1.000
Age      82.000
EDUC     15.000
SES       2.000
eTIV    1729.000
nWBV      0.698
Name: 300, dtype: float64
1 0
```

```
[ ]: # CB-FN-2
      expected_value = CB_explainer.expected_value
      shap.decision_plot(expected_value, CB_shap.values[2], X_test.iloc[2],
      ↪highlight=0)
```



```
[ ]: # CB-FN-31
expected_value = CB_explainer.expected_value
shap.decision_plot(expected_value, CB_shap.values[31], X_test.iloc[31],
highlight=0)
```



```
[ ]: # CB-FN-2
shap.initjs()
expected_value = CB_explainer.expected_value
shap.force_plot(expected_value, CB_shap.values[2], X_test.iloc[2])
```

<IPython.core.display.HTML object>

```
[ ]: <shap.plots._force.AdditiveForceVisualizer at 0x7f38fd7e2690>
```

```
[ ]: # CB-FN-31
shap.initjs()
expected_value = CB_explainer.expected_value
shap.force_plot(expected_value, CB_shap.values[31], X_test.iloc[31])
```

<IPython.core.display.HTML object>

```
[ ]: <shap.plots._force.AdditiveForceVisualizer at 0x7f38fd554710>
```

```
[ ]: lime_CB_explainer = lime.lime_tabular.LimeTabularExplainer(X_test.values,
                                                                feature_names=X_test.columns,
                                                                class_names=['Nondemented', 'Demented'])
```

```
[ ]: # CB-FN-2
lime_CB_explainer.explain_instance(X_test.iloc[2].values,\
                                   CatBoost_md1.predict_proba,\
                                   num_features=6).\
                                   show_in_notebook(predict_proba=True)
```

<IPython.core.display.HTML object>

```
[ ]: # CB-FN-31
lime_CB_explainer.explain_instance(X_test.iloc[31].values,\
                                   CatBoost_md1.predict_proba,\
                                   num_features=6).\
                                   show_in_notebook(predict_proba=True)
```

<IPython.core.display.HTML object>

```
[ ]: # store all fn_results
fn_results = []
feature_counter = Counter()

CB_FN_indices = [ 2, 30, 31, 37, 49, 61]

for fn_idx in CB_FN_indices:
    CB_instance_values = X_test.iloc[fn_idx].values
```

```

exp = lime_CB_explainer.explain_instance(
    CB_instance_values,
    CatBoost_mdl.predict_proba,
    num_features=6
)

exp_list = exp.as_list()

pushed_non = [f for f, w in exp_list if w < 0]
pushed_dem = [f for f, w in exp_list if w > 0]

fn_results.append({
    'Index': fn_idx,
    'Pushed_Nondemented': pushed_non,
    'Pushed_Demented': pushed_dem
})

feature_counter.update(pushed_non)

fn_df = pd.DataFrame(fn_results)

top_causes = pd.DataFrame(feature_counter.most_common(), columns=['Feature', 'Count'])

print(fn_df.head())

print("\n=== CB False Negative Feature Frequency ===")
print(top_causes)

```

	Index	Pushed_Nondemented \
0	2	[eTIV > 1669.00, 75.00 < Age <= 80.25]
1	30	[EDUC > 16.25, SES <= 2.00, 1491.50 < eTIV <= ...
2	31	[Age > 80.25, 0.00 < M/F <= 1.00, eTIV > 1669...
3	37	[0.00 < M/F <= 1.00, eTIV > 1669.00, 75.00 < A...
4	49	[EDUC > 16.25, 0.73 < nWBV <= 0.76, SES <= 2.00]

	Pushed_Demented
0	[nWBV <= 0.70, EDUC <= 12.00, M/F <= 0.00, 2.0...
1	[Age <= 71.00, M/F <= 0.00, 0.70 < nWBV <= 0.73]
2	[nWBV <= 0.70, 12.00 < EDUC <= 15.00]
3	[nWBV <= 0.70, 12.00 < EDUC <= 15.00]
4	[Age <= 71.00, M/F <= 0.00, 1491.50 < eTIV <= ...


```

=== CB False Negative Feature Frequency ===

```

	Feature	Count
0	SES <= 2.00	5
1	eTIV > 1669.00	3

2	0.00 < M/F <= 1.00	3
3	75.00 < Age <= 80.25	2
4	EDUC > 16.25	2
5	Age > 80.25	2
6	0.73 < nWBV <= 0.76	2
7	1491.50 < eTIV <= 1669.00	1

7.3 Additional testing

Among all three gradient boosting models, from the shap global analysis, we found that older age and female tend to push result to the nondemented side, which is contrary to the facts. In reality, age is positively correlated with Alzheimer's disease. Women, because they live longer than men, are more likely to be included in Alzheimer's disease samples than men. (Some studies have also shown that no significant gender differences were found in analyses of the same age group.)


This result maybe due to sample selection bias.

7.3.1 Check Gender and Age distribution by Group

```
[ ]: sb.countplot(x='Group', hue='M/F', data=df)

pct = pd.crosstab(df['Group'], df['M/F'], normalize='index') * 100
plt.title('M/F Distribution by Group')
print(pct)
```

M/F	F	M
Group		
Demented	41.095890	58.904110
Nondemented	67.894737	32.105263



```
[ ]: df['M/F'].value_counts()
```

```
[ ]: M/F
      F    189
      M    147
      Name: count, dtype: int64
```

in Nondemented group 68% is female


in demented group, 41% is female.

```
[ ]: # double check Age by group
      # Create a figure with 1 row and 2 columns of subplots
      fig, axes = plt.subplots(1, 2, figsize=(10, 4))

      # First subplot: Boxplot of Age by group
      sb.boxplot(x='Group', y='Age', data=df, ax=axes[0])
      axes[0].set_title('Age by Group')
      axes[0].set_xlabel('Group')
      axes[0].set_ylabel('Age')
```

```
# Second subplot: Distribution of Age colored by Group
sb.histplot(data=df, x='Age', hue='Group', bins=10, multiple='stack',
            ax=axes[1]) # or multiple='dodge'
axes[1].set_title('Age Distribution by Group')
axes[1].set_xlabel('Age')
axes[1].set_ylabel('Count')


# Adjust layout to prevent overlap
plt.tight_layout()
plt.show()
```



```
[ ]: sb.boxplot(x='Group', y='Age', hue='M/F', data=df)
df.groupby('Group')['Age'].describe()
```

```
[ ]:
```

	count	mean	std	min	25%	50%	75%	max
Group								
Demented	146.0	76.260274	6.940193	61.0	71.0	76.0	81.0	98.0
Nondemented	190.0	77.057895	8.096104	60.0	71.0	77.0	82.0	97.0



```
[ ]: from scipy.stats import ttest_ind, mannwhitneyu
```

```
nd_age = df[df['Group']=='Nondemented']['Age']
```

```
d_age = df[df['Group']=='Demented']['Age']
```

```
# t test
```

```
t_stat, p_val = ttest_ind(nd_age, d_age, equal_var=False)
```

```
print(f"T-test: t={t_stat:.3f}, p={p_val:.3f}")
```

```
# Mann-Whitney U
```

```
u_stat, p_val_u = mannwhitneyu(nd_age, d_age, alternative='two-sided')
```

```
print(f"Mann-Whitney U: U={u_stat:.3f}, p={p_val_u:.3f}")
```

T-test: t=0.971, p=0.332

Mann-Whitney U: U=14745.000, p=0.321

```
[ ]: # 1) age group
```

```
bins = [0, 70, 80, np.inf]
```

```
labels = ['<70', '70-80', '>80']
```

```
df['AgeBand'] = pd.cut(df['Age'].astype(float), bins=bins, labels=labels, right=False)
```



```

counts = df.groupby(['AgeBand', 'Group', 'M/F']).size().
↳reset_index(name='count')

g = sb.catplot(
    data=counts,
    x='AgeBand',
    y='count',
    hue='M/F',
    col='Group',
    kind='bar',
    palette={'F': '#FF9999', 'M': '#9999FF'},
    height=4,
    aspect=1
)
g.set_axis_labels("Age band", "Count")
g.set_titles("{col_name}")
g._legend.set_title("Gender")

plt.show()

# 2) count every group's Demented% and Female%
summary = (
    df.groupby('AgeBand')
    .agg(
        n_total      = ('Group', 'size'),
        n_demented    = ('Group', lambda x: (x.str.lower() == 'demented').sum()),
        n_female      = ('M/F', lambda x: (x.str.upper() == 'F').sum())
    )
)

summary['pct_demented'] = (summary['n_demented'] / summary['n_total'] * 100).
↳round(1)
summary['pct_female']   = (summary['n_female'] / summary['n_total'] * 100).
↳round(1)

print(summary.reset_index())


```

/tmp/ipython-input-1705621936.py:6: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```

counts = df.groupby(['AgeBand', 'Group',
'M/F']).size().reset_index(name='count')

```



	AgeBand	n_total	n_demented	n_female	pct_demented	pct_female
0	<70	60	25	30	41.7	50.0
1	70-80	154	74	86	48.1	55.8
2	>80	122	47	73	38.5	59.8

/tmp/ipython-input-1705621936.py:27: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```
df.groupby('AgeBand')
```

1. Dementia proportion (pct_demented)

- The <70 group has the lowest proportion (41.7%).
- The 70–80 age group has the highest dementia proportion (48.1%).
- In the older group (>80), the dementia proportion actually drops to 38.5%.

This pattern does not fully align with the common expectation that “older age means higher risk,” suggesting the presence of healthy older survivors (especially women) in our dataset. This survivor bias likely contributes to the lower dementia proportion in the >80 group.

2. Female proportion (pct_female)

The proportion of females increases with age: <70 (50%) → 70–80 (55.8%) → >80 (59.8%), which is an expected pattern, as women generally live longer, leading to a higher proportion of females in the oldest age group.

3. Implication for our SHAP results

In the >80 group, the proportion of females is very high, and a substantial portion of these women are nondemented. This can lead the model to learn the pattern “female & older age → nondemented.”

Combined with the drop in dementia proportion for the >80 group, these two factors together may explain why, in the global SHAP analysis, Age and Female show contributions toward the nondemented class — a direction contrary to typical clinical expectations.

7.3.2 Model training-XGBoost_remove SES

```
[ ]: # Try remove SES feature
X = df.drop(['Group', 'CDR', 'MMSE', 'ASF', 'SES'], axis = 1)
y = df['Group']

# Split dataset into training and test sets (80% training, 20% testing)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)
```

```
[ ]: X_test.columns
```

```
[ ]: Index(['M/F', 'Age', 'EDUC', 'eTIV', 'nWBV'], dtype='object')
```

```
[ ]: xgb = XGBClassifier(random_state=42,n_jobs=-1)

param_grid = {
    'n_estimators': [100, 200],
    'max_depth': [3, 5, 7],
    'learning_rate': [0.001, 0.01, 0.1],
    'subsample': [0.7, 0.8, 0.9],
    'colsample_bytree': [0.7, 0.8, 1.0],
    'min_child_weight': [1, 3, 5],
}

inner_cv = RepeatedStratifiedKFold(
    n_splits=5, n_repeats=2, random_state=42
)

grid = GridSearchCV(
    estimator=xgb,
    param_grid=param_grid,
    scoring='f1',
    cv=inner_cv,
    n_jobs=-1,
    verbose=0
)

grid.fit(X_train, y_train)
best_model = grid.best_estimator_
print("XGBoost classifier:")
print("Best params:", grid.best_params_)
```

```

scoring = {
    'accuracy': 'accuracy',
    'precision': 'precision',
    'recall': 'recall',
    'f1': 'f1',
    'roc_auc': 'roc_auc'
}
cv_res = cross_validate(
    best_model, X_train, y_train,
    cv=inner_cv, scoring=scoring, n_jobs=-1, return_train_score=False
)

def show(name):
    vals = cv_res[f'test_{name}']
    print(f"CV {name:<9}: {vals.mean():.4f} ± {vals.std():.4f}")

print("\n=== Repeated Stratified 5×2 CV on TRAIN ===")
for m in ['accuracy', 'precision', 'recall', 'f1', 'roc_auc']:
    show(m)

y_pred = best_model.predict(X_test)
y_prob = best_model.predict_proba(X_test)[:, 1]

acc = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred)
rec = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
auc = roc_auc_score(y_test, y_prob)

print("\n=== Final Performance on HOLD-OUT TEST ===")
print(f"Accuracy : {acc:.4f}")
print(f"Precision: {prec:.4f}")
print(f"Recall    : {rec:.4f}")
print(f"F1 Score  : {f1:.4f}")
print(f"ROC_AUC   : {auc:.4f}")

print("\nClassification Report on Test Set:\n", classification_report(y_test,
↪y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))

feat_imp = pd.DataFrame({
    "Feature": X_train.columns,
    "Importance": best_model.feature_importances_
}).sort_values(by="Importance", ascending=False)
print("\nTop feature importances:\n", feat_imp)

XGBoost_md1 = best_model

```

XGBoost classifier:

Best params: {'colsample_bytree': 1.0, 'learning_rate': 0.1, 'max_depth': 5, 'min_child_weight': 1, 'n_estimators': 200, 'subsample': 0.9}

=== Repeated Stratified 5×2 CV on TRAIN ===

CV accuracy : 0.8005 ± 0.0613

CV precision: 0.7818 ± 0.0930

CV recall : 0.7638 ± 0.0690

CV f1 : 0.7696 ± 0.0647

CV roc_auc : 0.8551 ± 0.0533

=== Final Performance on HOLD-OUT TEST ===

Accuracy : 0.8382

Precision: 0.7879

Recall : 0.8667

F1 Score : 0.8254

ROC_AUC : 0.8860

Classification Report on Test Set:

	precision	recall	f1-score	support
0	0.89	0.82	0.85	38
1	0.79	0.87	0.83	30
accuracy			0.84	68
macro avg	0.84	0.84	0.84	68
weighted avg	0.84	0.84	0.84	68

Confusion Matrix:

[[31 7]

[4 26]]


Top feature importances:

	Feature	Importance
2	EDUC	0.279822
0	M/F	0.253008
4	nWBV	0.174518
3	eTIV	0.149782
1	Age	0.142870

```
[ ]: fpr, tpr, _ = roc_curve(y_test, y_prob)
      roc_auc = roc_auc_score(y_test, y_prob)

      plt.figure()
      plt.plot(fpr, tpr, label=f'ROC curve (AUC = {roc_auc:.3f})')
      plt.plot([0, 1], [0, 1], linestyle='--', linewidth=1)
      plt.xlabel('False Positive Rate')
```

```
plt.ylabel('True Positive Rate')
plt.title('XGBoost Receiver Operating Characteristic')
plt.legend(loc='lower right')
plt.show()
```




```
[ ]: cm = confusion_matrix(y_test, y_pred)
tn, fp, fn, tp = cm.ravel()

labels = np.array([
    [f"TN = {tn}", f"FP = {fp}"],
    [f"FN = {fn}", f"TP = {tp}"]
])

fig, ax = plt.subplots()
ax.matshow(cm, cmap=plt.cm.Blues, alpha=0.3)
for i in range(2):
    for j in range(2):
        ax.text(j, i, labels[i, j],
                va='center', ha='center', fontsize=12, fontweight='bold')
```

```
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix with Labels")
plt.show()
```



```
[ ]: # Find all FN indices in the full test set
FN_all = (~y_pred) & (y_test == 1)
FN_indices = y_test[FN_all].index
print("False Negative indices:", FN_indices)

# Find all FP indices in the full test set
FP_all = (y_pred) & (y_test == 0)
FP_indices = y_test[FP_all].index
print("False Positive indices:", FP_indices)

# Find all TP indices in the full test set
TP_all = (y_pred) & (y_test == 1)
TP_indices = y_test[TP_all].index
print("True Positive indices:", TP_indices)
```

```
# Find all FN indices in the full test set
TN_all = (~y_pred) & (y_test == 0)
TN_indices = y_test[TN_all].index
print("True Negative indices:", TN_indices)
```

```
False Negative indices: Index([300, 299, 51, 94], dtype='int64')
False Positive indices: Index([167, 146, 198, 130, 199, 7, 64], dtype='int64')
True Positive indices: Index([124, 332, 250, 317, 154, 25, 90, 106, 172, 285,
87, 215, 127, 52,
3, 239, 162, 345, 72, 39, 89, 88, 16, 329, 365, 275],
dtype='int64')
True Negative indices: Index([ 84, 122, 311, 48, 336, 213, 9, 210, 113, 85,
363, 66, 5, 153,
291, 370, 158, 102, 243, 132, 306, 337, 196, 351, 179, 309, 197, 362,
209, 333, 96],
dtype='int64')
```

```
[ ]: FN_sample_test_idx = X_test.index.get_indexer_for([300, 299, 51, 94])
FP_sample_test_idx = X_test.index.get_indexer_for([167, 146, 198, 130, 199, 7, 64])
TP_sample_test_idx = X_test.index.get_indexer_for([124, 332, 250, 317, 154, 25, 90, 106, 172, 285, 87, 215, 127, 52, 3, 239, 162, 345, 72, 39, 89, 88, 16, 329, 365, 275])
TN_sample_test_idx = X_test.index.get_indexer_for([84, 122, 311, 48, 336, 213, 9, 210, 113, 85, 363, 66, 5, 153, 291, 370, 158, 102, 243, 132, 306, 337, 196, 351, 179, 309, 197, 362, 209, 333, 96])
```

```
[ ]: print("FN_sample_test_idx: ", FN_sample_test_idx)
print("FP_sample_test_idx: ", FP_sample_test_idx)
print("TP_sample_test_idx: ", TP_sample_test_idx)
print("TN_sample_test_idx: ", TN_sample_test_idx)
```

```
FN_sample_test_idx: [31 37 49 61]
FP_sample_test_idx: [17 21 42 50 53 59 60]
TP_sample_test_idx: [ 1  2  6  7  9 11 13 14 16 20 22 28 29 30 33 35 36 39 40
41 47 55 56 64
65 67]
TN_sample_test_idx: [ 0  3  4  5  8 10 12 15 18 19 23 24 25 26 27 32 34 38 43
44 45 46 48 51
52 54 57 58 62 63 66]
```

```
[ ]: XGB_explainer = shap.Explainer(XGBoost_md1)
XGB_shap = XGB_explainer(X_test)
print(type(XGB_explainer))
```




```
<class 'shap.explainers._tree.TreeExplainer'>
```

```
[ ]: print("Values dimensions: %s" % (XGB_shap.values.shape,))  
print("Data dimensions: %s" % (XGB_shap.data.shape,))
```

```
Values dimensions: (68, 5)
```

```
Data dimensions: (68, 5)
```

```
[ ]: sb.reset_orig()  
shap.plots.bar(XGB_shap)
```



```
[ ]: shap.plots.beeswarm(XGB_shap)
```



From the shap plot above, we can see that M/F(M=0, F=1) female and older age tend to push class to nondemented side.

```
[ ]: interaction_values = XGB_shap_interaction_values
features = X_test.columns

mean_abs_interactions = np.abs(interaction_values).mean(axis=0)

# Absolute mean measures how strong the interaction is, regardless of direction.


interaction_df = pd.DataFrame(mean_abs_interactions, index=features,
                               ↪columns=features)
interaction_df = interaction_df.where(np.triu(np.ones(interaction_df.shape),
                               ↪k=1).astype(bool))

interaction_df = interaction_df.stack().reset_index()
interaction_df.columns = ['Feature 1', 'Feature 2', 'Mean |Interaction Value|']
interaction_df = interaction_df.sort_values(by='Mean |Interaction Value|',
                               ↪ascending=False)


print(interaction_df.head(10))
```

	Feature 1	Feature 2	Mean Interaction Value
5	Age	eTIV	0.304100
9	eTIV	nWBV	0.282089
7	EDUC	eTIV	0.268376
8	EDUC	nWBV	0.227692
6	Age	nWBV	0.170623
1	M/F	EDUC	0.157386
2	M/F	eTIV	0.134400
4	Age	EDUC	0.115562
0	M/F	Age	0.094171
3	M/F	nWBV	0.062230

```
[ ]: # FN -31
expected_value = XGB_explainer.expected_value
shap.decision_plot(expected_value, XGB_shap.values[31], X_test.iloc[31],
                   ↪highlight=0)
```



```
[ ]: # FN -37
expected_value = XGB_explainer.expected_value
shap.decision_plot(expected_value, XGB_shap.values[37], X_test.iloc[16],
highlight=0)
```



```
[ ]: # FN -31
shap.initjs()
expected_value = XGB_explainer.expected_value
shap.force_plot(expected_value, XGB_shap.values[31], X_test.iloc[31])
```

<IPython.core.display.HTML object>

```
[ ]: <shap.plots._force.AdditiveForceVisualizer at 0x786ee8999350>
```

```
[ ]: # FN -37
shap.initjs()
expected_value = XGB_explainer.expected_value
shap.force_plot(expected_value, XGB_shap.values[16], X_test.iloc[16])
```

<IPython.core.display.HTML object>

```
[ ]: <shap.plots._force.AdditiveForceVisualizer at 0x786ee9146010>
```

```
[ ]: lime_XGB_explainer = lime.lime_tabular.LimeTabularExplainer(X_test.values,
                                                                feature_names=X_test.columns,
                                                                class_names=['Nondemented', 'Demented'])
```

```
[ ]: # FN case -31
lime_XGB_explainer.explain_instance(X_test.iloc[31].values, \
                                   XGBoost_md1.predict_proba, \
                                   num_features=5). \
    show_in_notebook(predict_proba=True)
```

<IPython.core.display.HTML object>

```
[ ]: # FN case-37
lime_XGB_explainer.explain_instance(X_test.iloc[37].values, \
                                   XGBoost_md1.predict_proba, \
                                   num_features=5). \
    show_in_notebook(predict_proba=True)
```

<IPython.core.display.HTML object>

```
[ ]: # store all fn_results
fn_results = []
feature_counter = Counter()

FN_indices = [31, 37, 49, 61]

for fn_idx in FN_indices:
    instance_values = X_test.iloc[fn_idx].values

    exp = lime_XGB_explainer.explain_instance(
        instance_values,
        XGBoost_md1.predict_proba,
        num_features=6
    )
```

```

exp_list = exp.as_list()

pushed_non = [f for f, w in exp_list if w < 0]
pushed_dem = [f for f, w in exp_list if w > 0]

fn_results.append({
    'Index': fn_idx,
    'Pushed_Nondemented': pushed_non,
    'Pushed_Demented': pushed_dem
})

feature_counter.update(pushes_non)

fn_df = pd.DataFrame(fn_results)

top_causes = pd.DataFrame(feature_counter.most_common(), columns=['Feature', 'Count'])

print(fn_df.head())

print("\n=== False Negative Feature Frequency ===")
print(top_causes)

```

	Index	Pushed_Nondemented \
0	31	[0.00 < M/F <= 1.00, Age > 80.25, eTIV > 1669.00]
1	37	[0.00 < M/F <= 1.00, eTIV > 1669.00, 75.00 < A...
2	49	[EDUC > 16.25, 0.73 < nWBV <= 0.76, 1491.50 < ...
3	61	[0.00 < M/F <= 1.00, Age > 80.25, 0.73 < nWBV ...

	Pushed_Demented
0	[nWBV <= 0.70, 12.00 < EDUC <= 15.00]
1	[nWBV <= 0.70, 12.00 < EDUC <= 15.00]
2	[M/F <= 0.00, Age <= 71.00]
3	[1391.00 < eTIV <= 1491.50, 12.00 < EDUC <= 15...

=== False Negative Feature Frequency ===

	Feature	Count
0	0.00 < M/F <= 1.00	3
1	Age > 80.25	2
2	eTIV > 1669.00	2
3	0.73 < nWBV <= 0.76	2
4	75.00 < Age <= 80.25	1
5	EDUC > 16.25	1
6	1491.50 < eTIV <= 1669.00	1

removing SES didn't notably decrease XGB model's performance. Fewer FN cases(5→4), recall is improved

7.3.3 Model training-LightGBM_remove SES

```
[ ]: LightGBM = LGBMClassifier(random_state=42, verbosity=-1)

param_grid = {
    'n_estimators': [100, 200],
    'max_depth': [3, 5, 7],
    'learning_rate': [0.001, 0.01, 0.1],
    'num_leaves': [31, 50, 70],
}

inner_cv = RepeatedStratifiedKFold(
    n_splits=5, n_repeats=2, random_state=42
)

grid = GridSearchCV(
    estimator=LightGBM,
    param_grid=param_grid,
    scoring='f1',
    cv=inner_cv,
    n_jobs=-1,
    verbose=0
)

grid.fit(X_train, y_train)
best_model = grid.best_estimator_
print("LightGBM classifier:")
print("Best params:", grid.best_params_)

scoring = {
    'accuracy': 'accuracy',
    'precision': 'precision',
    'recall': 'recall',
    'f1': 'f1',
    'roc_auc': 'roc_auc'
}

cv_res = cross_validate(
    best_model, X_train, y_train,
    cv=inner_cv, scoring=scoring, n_jobs=-1, return_train_score=False
)

def show(name):
    vals = cv_res[f'test_{name}']
    print(f"CV {name:<9}: {vals.mean():.4f} ± {vals.std():.4f}")

print("\n=== Repeated Stratified 5x2 CV on TRAIN ===")
for m in ['accuracy', 'precision', 'recall', 'f1', 'roc_auc']:
```

```

show(m)

y_pred = best_model.predict(X_test)
y_prob = best_model.predict_proba(X_test)[: , 1]

acc = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred)
rec = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
auc = roc_auc_score(y_test, y_prob)

print("\n=== Final Performance on HOLD-OUT TEST ===")
print(f"Accuracy : {acc:.4f}")
print(f"Precision: {prec:.4f}")
print(f"Recall    : {rec:.4f}")
print(f"F1 Score  : {f1:.4f}")
print(f"ROC_AUC   : {auc:.4f}")

print("\nClassification Report on Test Set:\n", classification_report(y_test,
    ↪y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))

feat_imp = pd.DataFrame({
    "Feature": X_train.columns,
    "Importance": best_model.feature_importances_
}).sort_values(by="Importance", ascending=False)
print("\nTop feature importances:\n", feat_imp)

LightGBM_md1 = best_model

```

LightGBM classifier:

Best params: {'learning_rate': 0.1, 'max_depth': 5, 'n_estimators': 200, 'num_leaves': 31}

=== Repeated Stratified 5×2 CV on TRAIN ===

CV accuracy : 0.7783 ± 0.0724
 CV precision: 0.7604 ± 0.1032
 CV recall : 0.7335 ± 0.0954
 CV f1 : 0.7416 ± 0.0783
 CV roc_auc : 0.8376 ± 0.0706

=== Final Performance on HOLD-OUT TEST ===

Accuracy : 0.7647
 Precision: 0.7188
 Recall : 0.7667
 F1 Score : 0.7419
 ROC_AUC : 0.8658

Classification Report on Test Set:

	precision	recall	f1-score	support
0	0.81	0.76	0.78	38
1	0.72	0.77	0.74	30
accuracy			0.76	68
macro avg	0.76	0.76	0.76	68
weighted avg	0.77	0.76	0.77	68

Confusion Matrix:


```
[[29  9]
 [ 7 23]]
```

Top feature importances:

	Feature	Importance
3	eTIV	617
4	nWBV	434
1	Age	337
2	EDUC	194
0	M/F	68

```
[ ]: fpr, tpr, _ = roc_curve(y_test, y_prob)
     roc_auc = roc_auc_score(y_test, y_prob)

     plt.figure()
     plt.plot(fpr, tpr, label=f'ROC curve (AUC = {roc_auc:.3f})')
     plt.plot([0, 1], [0, 1], linestyle='--', linewidth=1)
     plt.xlabel('False Positive Rate')
     plt.ylabel('True Positive Rate')
     plt.title('LightGBM Receiver Operating Characteristic')
     plt.legend(loc='lower right')
     plt.show()
```





```
[ ]: cm = confusion_matrix(y_test, y_pred)
tn, fp, fn, tp = cm.ravel()

labels = np.array([
    [f"TN = {tn}", f"FP = {fp}"],
    [f"FN = {fn}", f"TP = {tp}"]
])

fig, ax = plt.subplots()
ax.matshow(cm, cmap=plt.cm.Blues, alpha=0.3)
for i in range(2):
    for j in range(2):
        ax.text(j, i, labels[i, j],
                va='center', ha='center', fontsize=12, fontweight='bold')

plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix with Labels")
plt.show()
```



```
[ ]: # Find all FN indices in the full test set
FN_all = (~y_pred) & (y_test == 1)
FN_indices = y_test[FN_all].index
print("False Negative indices:", FN_indices)
```

```
False Negative indices: Index([154, 300, 299, 39, 51, 94, 329], dtype='int64')
```

```
[ ]: FN_sample_test_idx = X_test.index.get_indexer_for([154, 300, 299, 39, 51, 94, ↵
↵329])
FN_sample_test_idx
```

```
[ ]: array([ 9, 31, 37, 41, 49, 61, 64])
```

```
[ ]: LGB_explainer = shap.Explainer(LightGBM_md1)
LGB_shap = LGB_explainer(X_test)
print(type(LGB_explainer))
```


```
<class 'shap.explainers._tree.TreeExplainer'>
```

```
[ ]: print("Values dimensions: %s" % (LGB_shap.values.shape,))
print("Data dimensions: %s" % (LGB_shap.data.shape,))
```


Values dimensions: (68, 5)

Data dimensions: (68, 5)


```
[ ]: sb.reset_orig()
shap.plots.bar(LGB_shap)
```




```
[ ]: shap.plots.beeswarm(LGB_shap)
```



```
[ ]: # LGB-FN-31
sb.reset_orig()
expected_value = LGB_explainer.expected_value
shap.decision_plot(expected_value, LGB_shap.values[31], X_test.iloc[31],
highlight=0)
```



```
[ ]: # LGB-FN-37
expected_value = LGB_explainer.expected_value
shap.decision_plot(expected_value, LGB_shap.values[37], X_test.iloc[37],
highlight=0)
```



```
[ ]: # LGB-FN-31
shap.initjs()
expected_value = LGB_explainer.expected_value
shap.force_plot(expected_value, LGB_shap.values[31], X_test.iloc[31])
```

<IPython.core.display.HTML object>

```
[ ]: <shap.plots._force.AdditiveForceVisualizer at 0x786ee95a1110>
```

```
[ ]: # LGB-FN-37
shap.initjs()
expected_value = LGB_explainer.expected_value
shap.force_plot(expected_value, LGB_shap.values[37], X_test.iloc[37])
```

<IPython.core.display.HTML object>

```
[ ]: <shap.plots._force.AdditiveForceVisualizer at 0x786eea188450>
```

```
[ ]: lime_LGB_explainer = lime.lime_tabular.LimeTabularExplainer(X_test.values,
                                                                feature_names=X_test.columns,
                                                                class_names=['Nondemented', 'Demented'],
                                                                class_labels={'Nondemented': 0, 'Demented': 1})
```

```
[ ]: # LGB-FN-31
lime_LGB_explainer.explain_instance(X_test.iloc[31].values, \
                                   LightGBM_md1.predict_proba, \
                                   num_features=6). \
    show_in_notebook(predict_proba=True)
```

/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739:
UserWarning: X does not have valid feature names, but LGBMClassifier was fitted with feature names

warnings.warn(

<IPython.core.display.HTML object>

```
[ ]: # LGB-FN-37
lime_LGB_explainer.explain_instance(X_test.iloc[37].values, \
                                   LightGBM_md1.predict_proba, \
                                   num_features=6). \
    show_in_notebook(predict_proba=True)
```

/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739:
UserWarning: X does not have valid feature names, but LGBMClassifier was fitted with feature names

warnings.warn(

<IPython.core.display.HTML object>

```
[ ]: # LGB-FN-49
lime_LGB_explainer.explain_instance(X_test.iloc[49].values,\
                                   LightGBM_md1.predict_proba,\
                                   num_features=6).\
                                   show_in_notebook(predict_proba=True)
```

/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739:
UserWarning: X does not have valid feature names, but LGBMClassifier was fitted
with feature names
warnings.warn(

<IPython.core.display.HTML object>

```
[ ]: # store all fn_results
fn_results = []
feature_counter = Counter()

LGB_FN_indices = [ 9, 31, 37, 41, 49, 61, 64]

for fn_idx in LGB_FN_indices:
    LGB_instance_values = X_test.iloc[fn_idx].values

    exp = lime_LGB_explainer.explain_instance(
        LGB_instance_values,
        LightGBM_md1.predict_proba,
        num_features=6
    )

    exp_list = exp.as_list()

    pushed_non = [f for f, w in exp_list if w < 0]
    pushed_dem = [f for f, w in exp_list if w > 0]

    fn_results.append({
        'Index': fn_idx,
        'Pushed_Nondemented': pushed_non,
        'Pushed_Demented': pushed_dem
    })

    feature_counter.update(pushed_non)

fn_df = pd.DataFrame(fn_results)

top_causes = pd.DataFrame(feature_counter.most_common(), columns=['Feature',
↪ 'Count'])
```

```
print(fn_df.head())

print("\n=== LGB False Negative Feature Frequency ===")
print(top_causes)
```

```
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739:
UserWarning: X does not have valid feature names, but LGBMClassifier was fitted
with feature names
  warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739:
UserWarning: X does not have valid feature names, but LGBMClassifier was fitted
with feature names
  warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739:
UserWarning: X does not have valid feature names, but LGBMClassifier was fitted
with feature names
  warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739:
UserWarning: X does not have valid feature names, but LGBMClassifier was fitted
with feature names
  warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739:
UserWarning: X does not have valid feature names, but LGBMClassifier was fitted
with feature names
  warnings.warn(
```

	Index	Pushed_Nondemented \
0	9	[0.00 < M/F <= 1.00, Age > 80.25]
1	31	[0.00 < M/F <= 1.00, eTIV > 1669.00, Age > 80.25]
2	37	[0.00 < M/F <= 1.00, eTIV > 1669.00, 75.00 < A...
3	41	[EDUC > 16.25, eTIV > 1669.00, 71.00 < Age <= ...
4	49	[EDUC > 16.25, 0.73 < nWBV <= 0.76, 1491.50 < ...

	Pushed_Demented
0	[nWBV <= 0.70, eTIV <= 1391.00, 12.00 < EDUC <...
1	[nWBV <= 0.70, 12.00 < EDUC <= 15.00]
2	[nWBV <= 0.70, 12.00 < EDUC <= 15.00]
3	[M/F <= 0.00, 0.70 < nWBV <= 0.73]
4	[M/F <= 0.00, Age <= 71.00]

=== LGB False Negative Feature Frequency ===

	Feature	Count
0	0.00 < M/F <= 1.00	4
1	eTIV > 1669.00	4
2	Age > 80.25	3
3	EDUC > 16.25	3
4	0.73 < nWBV <= 0.76	3
5	75.00 < Age <= 80.25	2

```
6      71.00 < Age <= 75.00      1
7  1491.50 < eTIV <= 1669.00    1
```

```
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739:
UserWarning: X does not have valid feature names, but LGBMClassifier was fitted
with feature names
```

```
warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739:
UserWarning: X does not have valid feature names, but LGBMClassifier was fitted
with feature names
```

```
warnings.warn(
```

compared with before removing SES, performance of LightGBM decreases. More FN cases(4->7)

7.3.4 Model training-CatBoost_remove SES

```
[ ]: CatBoost = CatBoostClassifier(random_state=42)
```

```
param_grid = {
    'min_data_in_leaf': [20, 40, 60],
    'rsm': [0.7, 0.8, 1.0],
    'iterations': [100, 200],
    'depth': [3, 5, 7],
    'learning_rate': [0.001, 0.01, 0.1],
    'l2_leaf_reg': [1, 3, 5],
    'bagging_temperature': [0, 0.5, 1.0],
}

inner_cv = RepeatedStratifiedKFold(
    n_splits=5, n_repeats=2, random_state=42
)

grid = GridSearchCV(
    estimator=CatBoost,
    param_grid=param_grid,
    scoring='f1',
    cv=inner_cv,
    n_jobs=-1,
    verbose=0
)

grid.fit(X_train, y_train)
best_model = grid.best_estimator_
print("CatBoostClassifier:")
print("Best params:", grid.best_params_)

scoring = {
    'accuracy': 'accuracy',
```



```

        'precision': 'precision',
        'recall': 'recall',
        'f1': 'f1',
        'roc_auc': 'roc_auc'
    }
    cv_res = cross_validate(
        best_model, X_train, y_train,
        cv=inner_cv, scoring=scoring, n_jobs=-1, return_train_score=False
    )

    def show(name):
        vals = cv_res[f'test_{name}']
        print(f"CV {name:<9}: {vals.mean():.4f} ± {vals.std():.4f}")

    print("\n=== Repeated Stratified 5×2 CV on TRAIN ===")
    for m in ['accuracy', 'precision', 'recall', 'f1', 'roc_auc']:
        show(m)

    y_pred = best_model.predict(X_test)
    y_prob = best_model.predict_proba(X_test)[:, 1]

    acc = accuracy_score(y_test, y_pred)
    prec = precision_score(y_test, y_pred)
    rec = recall_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)
    auc = roc_auc_score(y_test, y_prob)

    print("\n=== Final Performance on HOLD-OUT TEST ===")
    print(f"Accuracy : {acc:.4f}")
    print(f"Precision: {prec:.4f}")
    print(f"Recall   : {rec:.4f}")
    print(f"F1 Score  : {f1:.4f}")
    print(f"ROC_AUC   : {auc:.4f}")

    print("\nClassification Report on Test Set:\n", classification_report(y_test,
        ↪y_pred))
    print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))

    feat_imp = pd.DataFrame({
        "Feature": X_train.columns,
        "Importance": best_model.feature_importances_
    }).sort_values(by="Importance", ascending=False)
    print("\nTop feature importances:\n", feat_imp)

    CatBoost_md1 = best_model

```

0: learn: 0.6593913 total: 47.4ms remaining: 9.44s

1:	learn: 0.6161128	total: 48.4ms	remaining: 4.79s
2:	learn: 0.5747853	total: 49.4ms	remaining: 3.25s
3:	learn: 0.5599539	total: 50ms	remaining: 2.45s
4:	learn: 0.5253404	total: 51.1ms	remaining: 1.99s
5:	learn: 0.5001249	total: 52.1ms	remaining: 1.68s
6:	learn: 0.4951362	total: 52.7ms	remaining: 1.45s
7:	learn: 0.4747394	total: 53.9ms	remaining: 1.29s
8:	learn: 0.4546658	total: 55ms	remaining: 1.17s
9:	learn: 0.4454416	total: 55.6ms	remaining: 1.06s
10:	learn: 0.4294976	total: 56.5ms	remaining: 971ms
11:	learn: 0.4152974	total: 57.5ms	remaining: 900ms
12:	learn: 0.3948456	total: 58.4ms	remaining: 840ms
13:	learn: 0.3842047	total: 59.2ms	remaining: 787ms
14:	learn: 0.3773444	total: 60.1ms	remaining: 741ms
15:	learn: 0.3576999	total: 61.1ms	remaining: 702ms
16:	learn: 0.3499457	total: 62ms	remaining: 668ms
17:	learn: 0.3356367	total: 63.1ms	remaining: 638ms
18:	learn: 0.3232779	total: 64ms	remaining: 610ms
19:	learn: 0.3199074	total: 64.6ms	remaining: 582ms
20:	learn: 0.3055672	total: 65.7ms	remaining: 560ms
21:	learn: 0.2965582	total: 66.7ms	remaining: 540ms
22:	learn: 0.2830015	total: 67.9ms	remaining: 522ms
23:	learn: 0.2797083	total: 69.2ms	remaining: 507ms
24:	learn: 0.2692293	total: 70.2ms	remaining: 492ms
25:	learn: 0.2608171	total: 71.3ms	remaining: 477ms
26:	learn: 0.2554928	total: 72.3ms	remaining: 463ms
27:	learn: 0.2477168	total: 73.3ms	remaining: 450ms
28:	learn: 0.2382749	total: 74.3ms	remaining: 438ms
29:	learn: 0.2352668	total: 75.2ms	remaining: 426ms
30:	learn: 0.2282407	total: 76.2ms	remaining: 415ms
31:	learn: 0.2227499	total: 77.2ms	remaining: 405ms
32:	learn: 0.2161551	total: 78.1ms	remaining: 395ms
33:	learn: 0.2127196	total: 79.2ms	remaining: 387ms
34:	learn: 0.2070432	total: 80.1ms	remaining: 378ms
35:	learn: 0.2017828	total: 81.2ms	remaining: 370ms
36:	learn: 0.1926881	total: 82.2ms	remaining: 362ms
37:	learn: 0.1885964	total: 83.2ms	remaining: 355ms
38:	learn: 0.1843307	total: 84.3ms	remaining: 348ms
39:	learn: 0.1812445	total: 85.3ms	remaining: 341ms
40:	learn: 0.1789055	total: 86.2ms	remaining: 334ms
41:	learn: 0.1770552	total: 87.4ms	remaining: 329ms
42:	learn: 0.1769425	total: 87.9ms	remaining: 321ms
43:	learn: 0.1749386	total: 88.8ms	remaining: 315ms
44:	learn: 0.1749027	total: 89.3ms	remaining: 308ms
45:	learn: 0.1717094	total: 90.2ms	remaining: 302ms
46:	learn: 0.1689454	total: 91.2ms	remaining: 297ms
47:	learn: 0.1684493	total: 91.9ms	remaining: 291ms
48:	learn: 0.1652703	total: 92.8ms	remaining: 286ms

49:	learn: 0.1602030	total: 93.8ms	remaining: 281ms
50:	learn: 0.1569368	total: 94.7ms	remaining: 277ms
51:	learn: 0.1532159	total: 95.7ms	remaining: 272ms
52:	learn: 0.1470671	total: 96.7ms	remaining: 268ms
53:	learn: 0.1443544	total: 97.8ms	remaining: 265ms
54:	learn: 0.1423886	total: 99ms	remaining: 261ms
55:	learn: 0.1357377	total: 100ms	remaining: 257ms
56:	learn: 0.1304753	total: 101ms	remaining: 254ms
57:	learn: 0.1277747	total: 102ms	remaining: 250ms
58:	learn: 0.1250609	total: 103ms	remaining: 247ms
59:	learn: 0.1200222	total: 104ms	remaining: 243ms
60:	learn: 0.1170928	total: 105ms	remaining: 239ms
61:	learn: 0.1144087	total: 106ms	remaining: 236ms
62:	learn: 0.1119146	total: 107ms	remaining: 233ms
63:	learn: 0.1101312	total: 108ms	remaining: 230ms
64:	learn: 0.1078354	total: 109ms	remaining: 227ms
65:	learn: 0.1064887	total: 110ms	remaining: 224ms
66:	learn: 0.1042670	total: 111ms	remaining: 221ms
67:	learn: 0.0997478	total: 112ms	remaining: 218ms
68:	learn: 0.0978650	total: 113ms	remaining: 215ms
69:	learn: 0.0968160	total: 114ms	remaining: 212ms
70:	learn: 0.0952747	total: 115ms	remaining: 209ms
71:	learn: 0.0936632	total: 116ms	remaining: 206ms
72:	learn: 0.0911765	total: 117ms	remaining: 204ms
73:	learn: 0.0877513	total: 118ms	remaining: 201ms
74:	learn: 0.0860806	total: 119ms	remaining: 199ms
75:	learn: 0.0849196	total: 120ms	remaining: 196ms
76:	learn: 0.0836052	total: 121ms	remaining: 194ms
77:	learn: 0.0809637	total: 122ms	remaining: 192ms
78:	learn: 0.0797554	total: 124ms	remaining: 189ms
79:	learn: 0.0769329	total: 125ms	remaining: 187ms
80:	learn: 0.0760161	total: 126ms	remaining: 185ms
81:	learn: 0.0752541	total: 127ms	remaining: 183ms
82:	learn: 0.0740942	total: 128ms	remaining: 180ms
83:	learn: 0.0726738	total: 129ms	remaining: 178ms
84:	learn: 0.0708871	total: 130ms	remaining: 176ms
85:	learn: 0.0698076	total: 131ms	remaining: 174ms
86:	learn: 0.0685995	total: 132ms	remaining: 172ms
87:	learn: 0.0675179	total: 133ms	remaining: 170ms
88:	learn: 0.0662989	total: 134ms	remaining: 167ms
89:	learn: 0.0648178	total: 135ms	remaining: 165ms
90:	learn: 0.0637765	total: 137ms	remaining: 164ms
91:	learn: 0.0626972	total: 139ms	remaining: 163ms
92:	learn: 0.0617175	total: 140ms	remaining: 161ms
93:	learn: 0.0609860	total: 141ms	remaining: 159ms
94:	learn: 0.0596594	total: 142ms	remaining: 157ms
95:	learn: 0.0587644	total: 143ms	remaining: 155ms
96:	learn: 0.0574071	total: 144ms	remaining: 153ms

97:	learn: 0.0566022	total: 145ms	remaining: 151ms
98:	learn: 0.0557687	total: 147ms	remaining: 150ms
99:	learn: 0.0546145	total: 148ms	remaining: 148ms
100:	learn: 0.0538387	total: 149ms	remaining: 146ms
101:	learn: 0.0530472	total: 150ms	remaining: 144ms
102:	learn: 0.0519145	total: 151ms	remaining: 142ms
103:	learn: 0.0511723	total: 152ms	remaining: 140ms
104:	learn: 0.0503818	total: 153ms	remaining: 138ms
105:	learn: 0.0494722	total: 154ms	remaining: 137ms
106:	learn: 0.0485186	total: 155ms	remaining: 135ms
107:	learn: 0.0477060	total: 156ms	remaining: 133ms
108:	learn: 0.0468543	total: 157ms	remaining: 131ms
109:	learn: 0.0456000	total: 158ms	remaining: 129ms
110:	learn: 0.0443978	total: 159ms	remaining: 128ms
111:	learn: 0.0435653	total: 160ms	remaining: 126ms
112:	learn: 0.0432561	total: 161ms	remaining: 124ms
113:	learn: 0.0419331	total: 162ms	remaining: 122ms
114:	learn: 0.0417388	total: 163ms	remaining: 121ms
115:	learn: 0.0409675	total: 164ms	remaining: 119ms
116:	learn: 0.0405151	total: 165ms	remaining: 117ms
117:	learn: 0.0399008	total: 166ms	remaining: 115ms
118:	learn: 0.0384224	total: 167ms	remaining: 114ms
119:	learn: 0.0377821	total: 168ms	remaining: 112ms
120:	learn: 0.0374588	total: 169ms	remaining: 110ms
121:	learn: 0.0367639	total: 170ms	remaining: 109ms
122:	learn: 0.0365966	total: 171ms	remaining: 107ms
123:	learn: 0.0361286	total: 172ms	remaining: 105ms
124:	learn: 0.0356571	total: 173ms	remaining: 104ms
125:	learn: 0.0352078	total: 174ms	remaining: 102ms
126:	learn: 0.0345354	total: 175ms	remaining: 101ms
127:	learn: 0.0339649	total: 176ms	remaining: 98.9ms
128:	learn: 0.0335745	total: 177ms	remaining: 97.4ms
129:	learn: 0.0325093	total: 178ms	remaining: 95.9ms
130:	learn: 0.0319434	total: 179ms	remaining: 94.3ms
131:	learn: 0.0312770	total: 180ms	remaining: 92.7ms
132:	learn: 0.0310424	total: 181ms	remaining: 91.2ms
133:	learn: 0.0301217	total: 182ms	remaining: 89.6ms
134:	learn: 0.0298561	total: 183ms	remaining: 88.1ms
135:	learn: 0.0294238	total: 184ms	remaining: 86.5ms
136:	learn: 0.0289752	total: 185ms	remaining: 85ms
137:	learn: 0.0286738	total: 186ms	remaining: 83.6ms
138:	learn: 0.0284155	total: 187ms	remaining: 82ms
139:	learn: 0.0282291	total: 188ms	remaining: 80.5ms
140:	learn: 0.0278695	total: 189ms	remaining: 79.1ms
141:	learn: 0.0275022	total: 190ms	remaining: 77.6ms
142:	learn: 0.0273527	total: 191ms	remaining: 76.1ms
143:	learn: 0.0270758	total: 192ms	remaining: 74.6ms
144:	learn: 0.0268138	total: 193ms	remaining: 73.3ms

145:	learn: 0.0261164	total: 194ms	remaining: 71.9ms
146:	learn: 0.0254907	total: 195ms	remaining: 70.5ms
147:	learn: 0.0253650	total: 196ms	remaining: 69ms
148:	learn: 0.0251657	total: 197ms	remaining: 67.6ms
149:	learn: 0.0248262	total: 198ms	remaining: 66.1ms
150:	learn: 0.0246040	total: 199ms	remaining: 64.7ms
151:	learn: 0.0243856	total: 200ms	remaining: 63.3ms
152:	learn: 0.0243007	total: 201ms	remaining: 61.8ms
153:	learn: 0.0240222	total: 202ms	remaining: 60.4ms
154:	learn: 0.0239106	total: 203ms	remaining: 59ms
155:	learn: 0.0235797	total: 204ms	remaining: 57.6ms
156:	learn: 0.0233511	total: 205ms	remaining: 56.2ms
157:	learn: 0.0231489	total: 206ms	remaining: 54.8ms
158:	learn: 0.0228279	total: 207ms	remaining: 53.4ms
159:	learn: 0.0225688	total: 208ms	remaining: 52.1ms
160:	learn: 0.0220480	total: 209ms	remaining: 50.7ms
161:	learn: 0.0217246	total: 210ms	remaining: 49.3ms
162:	learn: 0.0215687	total: 211ms	remaining: 48ms
163:	learn: 0.0213136	total: 212ms	remaining: 46.6ms
164:	learn: 0.0211194	total: 213ms	remaining: 45.2ms
165:	learn: 0.0208881	total: 214ms	remaining: 43.8ms
166:	learn: 0.0207319	total: 215ms	remaining: 42.5ms
167:	learn: 0.0205932	total: 216ms	remaining: 41.1ms
168:	learn: 0.0205057	total: 217ms	remaining: 39.7ms
169:	learn: 0.0204389	total: 217ms	remaining: 38.4ms
170:	learn: 0.0200581	total: 219ms	remaining: 37.1ms
171:	learn: 0.0199939	total: 220ms	remaining: 35.7ms
172:	learn: 0.0198272	total: 221ms	remaining: 34.4ms
173:	learn: 0.0195395	total: 222ms	remaining: 33.1ms
174:	learn: 0.0194205	total: 223ms	remaining: 31.8ms
175:	learn: 0.0192243	total: 224ms	remaining: 30.5ms
176:	learn: 0.0190095	total: 225ms	remaining: 29.2ms
177:	learn: 0.0188153	total: 226ms	remaining: 27.9ms
178:	learn: 0.0185343	total: 227ms	remaining: 26.6ms
179:	learn: 0.0182716	total: 228ms	remaining: 25.3ms
180:	learn: 0.0180765	total: 229ms	remaining: 24ms
181:	learn: 0.0179598	total: 231ms	remaining: 22.8ms
182:	learn: 0.0177268	total: 232ms	remaining: 21.5ms
183:	learn: 0.0175665	total: 234ms	remaining: 20.3ms
184:	learn: 0.0173862	total: 235ms	remaining: 19.1ms
185:	learn: 0.0171053	total: 237ms	remaining: 17.8ms
186:	learn: 0.0169377	total: 240ms	remaining: 16.7ms
187:	learn: 0.0168117	total: 241ms	remaining: 15.4ms
188:	learn: 0.0167093	total: 242ms	remaining: 14.1ms
189:	learn: 0.0166480	total: 244ms	remaining: 12.9ms
190:	learn: 0.0164774	total: 246ms	remaining: 11.6ms
191:	learn: 0.0164050	total: 247ms	remaining: 10.3ms
192:	learn: 0.0163232	total: 249ms	remaining: 9.04ms

```

193:   learn: 0.0162641      total: 251ms   remaining: 7.76ms
194:   learn: 0.0161572      total: 253ms   remaining: 6.47ms
195:   learn: 0.0159926      total: 254ms   remaining: 5.17ms
196:   learn: 0.0158394      total: 255ms   remaining: 3.88ms
197:   learn: 0.0157380      total: 256ms   remaining: 2.58ms
198:   learn: 0.0156695      total: 257ms   remaining: 1.29ms
199:   learn: 0.0155247      total: 258ms   remaining: 0us

```

CatBoostClassifier:

```

Best params: {'bagging_temperature': 0, 'depth': 7, 'iterations': 200,
'l2_leaf_reg': 1, 'learning_rate': 0.1, 'min_data_in_leaf': 20, 'rsm': 0.8}

```

=== Repeated Stratified 5×2 CV on TRAIN ===

```

CV accuracy : 0.8267 ± 0.0677
CV precision: 0.8023 ± 0.0864
CV recall   : 0.8025 ± 0.0837
CV f1       : 0.8007 ± 0.0756
CV roc_auc  : 0.8841 ± 0.0492

```

=== Final Performance on HOLD-OUT TEST ===

```

Accuracy : 0.8676
Precision: 0.8387
Recall   : 0.8667
F1 Score : 0.8525
ROC_AUC  : 0.9289

```

Classification Report on Test Set:

	precision	recall	f1-score	support
0	0.89	0.87	0.88	38
1	0.84	0.87	0.85	30
accuracy			0.87	68
macro avg	0.87	0.87	0.87	68
weighted avg	0.87	0.87	0.87	68

Confusion Matrix:

```

[[33  5]
 [ 4 26]]


```

Top feature importances:

	Feature	Importance
3	eTIV	28.136046
2	EDUC	23.609229
4	nWBV	21.381349
1	Age	17.861711
0	M/F	9.011665

```
[ ]: fpr, tpr, _ = roc_curve(y_test, y_prob)
roc_auc = roc_auc_score(y_test, y_prob)

plt.figure()
plt.plot(fpr, tpr, label=f'ROC curve (AUC = {roc_auc:.3f})')
plt.plot([0, 1], [0, 1], linestyle='--', linewidth=1)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('CatBoost Receiver Operating Characteristic')
plt.legend(loc='lower right')
plt.show()
```



```
[ ]: cm = confusion_matrix(y_test, y_pred)
tn, fp, fn, tp = cm.ravel()


labels = np.array([
    [f"TN = {tn}", f"FP = {fp}"],
    [f"FN = {fn}", f"TP = {tp}"]
])
```

```

fig, ax = plt.subplots()
ax.matshow(cm, cmap=plt.cm.Blues, alpha=0.3)
for i in range(2):
    for j in range(2):
        ax.text(j, i, labels[i, j],
                va='center', ha='center', fontsize=12, fontweight='bold')

plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix with Labels")
plt.show()

```



```

[ ]: # Find all FN indices in the full test set
FN_all = (~y_pred) & (y_test == 1)
FN_indices = y_test[FN_all].index
print("False Negative indices:", FN_indices)

```

False Negative indices: Index([300, 299, 51, 94], dtype='int64')


```
[ ]: FN_sample_test_idx = X_test.index.get_indexer_for([300, 299, 51, 94])
FN_sample_test_idx
```

```
[ ]: array([31, 37, 49, 61])
```

```
[ ]: CB_explainer = shap.Explainer(CatBoost_md1)
CB_shap = CB_explainer(X_test)
print(type(CB_explainer))
```


```
<class 'shap.explainers._tree.TreeExplainer'>
```

```
[ ]: print("Values dimensions: %s" % (CB_shap.values.shape,))
print("Data dimensions: %s" % (CB_shap.data.shape,))
```


```
Values dimensions: (68, 5)
```

```
Data dimensions: (68, 5)
```

```
[ ]: sb.reset_orig()
shap.plots.bar(CB_shap)
```



```
[ ]: shap.plots.beeswarm(CB_shap)
```




```
[ ]: print("X_test.iloc[37]: ")
      print(X_test.iloc[37])
      print(y_test.iloc[37], y_pred[37])
      print("-----")
      print("X_test.iloc[31]: ")
      print(X_test.iloc[31])
      print(y_test.iloc[31], y_pred[31])
```


```
X_test.iloc[37]:
M/F      1.000
Age      80.000
EDUC     15.000
eTIV    1732.000
nWBV      0.685
Name: 299, dtype: float64
1 0
-----
X_test.iloc[31]:
M/F      1.000
Age      82.000
EDUC     15.000
eTIV    1729.000
nWBV      0.698
Name: 300, dtype: float64
1 0
```

```
[ ]: # CB-FN-37
      expected_value = CB_explainer.expected_value
```

```
shap.decision_plot(expected_value, CB_shap.values[37], X_test.iloc[37],  
highlight=0)
```



```
[ ]: # CB-FN-31  
expected_value = CB_explainer.expected_value  
shap.decision_plot(expected_value, CB_shap.values[31], X_test.iloc[31],  
highlight=0)
```



```
[ ]: # CB-FN-37
shap.initjs()
expected_value = CB_explainer.expected_value
shap.force_plot(expected_value, CB_shap.values[37], X_test.iloc[37])
```

<IPython.core.display.HTML object>

```
[ ]: <shap.plots._force.AdditiveForceVisualizer at 0x786eec91ba50>
```

```
[ ]: # CB-FN-31
shap.initjs()
expected_value = CB_explainer.expected_value
shap.force_plot(expected_value, CB_shap.values[31], X_test.iloc[31])
```

<IPython.core.display.HTML object>

```
[ ]: <shap.plots._force.AdditiveForceVisualizer at 0x786eeb9ffe90>
```

```
[ ]: lime_CB_explainer = lime.lime_tabular.LimeTabularExplainer(X_test.values,
                                                                feature_names=X_test.columns,
                                                                class_names=['Nondemented', 'Demented'])
```

```
[ ]: # CB-FN-37
lime_CB_explainer.explain_instance(X_test.iloc[37].values,\
                                   CatBoost_md1.predict_proba,\
                                   num_features=6).\
                                   show_in_notebook(predict_proba=True)
```

<IPython.core.display.HTML object>

```
[ ]: # CB-FN-31
lime_CB_explainer.explain_instance(X_test.iloc[31].values,\
                                   CatBoost_md1.predict_proba,\
                                   num_features=6).\
                                   show_in_notebook(predict_proba=True)
```

<IPython.core.display.HTML object>

```
[ ]: # store all fn_results
fn_results = []
feature_counter = Counter()

CB_FN_indices = [31, 37, 49, 61]

for fn_idx in CB_FN_indices:
    CB_instance_values = X_test.iloc[fn_idx].values
```

```

exp = lime_CB_explainer.explain_instance(
    CB_instance_values,
    CatBoost_mdl.predict_proba,
    num_features=6
)

exp_list = exp.as_list()

pushed_non = [f for f, w in exp_list if w < 0]
pushed_dem = [f for f, w in exp_list if w > 0]

fn_results.append({
    'Index': fn_idx,
    'Pushed_Nondemented': pushed_non,
    'Pushed_Demented': pushed_dem
})

feature_counter.update(pushes_non)

fn_df = pd.DataFrame(fn_results)

top_causes = pd.DataFrame(feature_counter.most_common(), columns=['Feature', 'Count'])

print(fn_df.head())

print("\n=== CB False Negative Feature Frequency ===")
print(top_causes)

```

	Index	Pushed_Nondemented \
0	31	[0.00 < M/F <= 1.00, Age > 80.25, eTIV > 1669.00]
1	37	[0.00 < M/F <= 1.00, eTIV > 1669.00, 75.00 < A...
2	49	[EDUC > 16.25, 0.73 < nWBV <= 0.76, 1491.50 < ...
3	61	[0.00 < M/F <= 1.00, Age > 80.25, 0.73 < nWBV ...

	Pushed_Demented
0	[nWBV <= 0.70, 12.00 < EDUC <= 15.00]
1	[nWBV <= 0.70, 12.00 < EDUC <= 15.00]
2	[M/F <= 0.00, Age <= 71.00]
3	[1391.00 < eTIV <= 1491.50, 12.00 < EDUC <= 15...

=== CB False Negative Feature Frequency ===

	Feature	Count
0	0.00 < M/F <= 1.00	3
1	Age > 80.25	2
2	eTIV > 1669.00	2
3	0.73 < nWBV <= 0.76	2

```

4      75.00 < Age <= 80.25      1
5      EDUC > 16.25      1
6  1491.50 < eTIV <= 1669.00      1

```

removing SES didn't notably decrease CatBoost model's performance. Fewer FN cases(6→4), recall is improved

7.3.5 Model training-XGBoost_remove SES and Gender

```

[ ]: # Try remove SES feature
X = df.drop(['Group', 'CDR', 'MMSE', 'ASF', 'M/F', 'SES'], axis = 1)
y = df['Group']

# Split dataset into training and test sets (80% training, 20% testing)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)

```

```

[ ]: X_test.columns

```

```

[ ]: Index(['Age', 'EDUC', 'eTIV', 'nWBV'], dtype='object')

```

```

[ ]: xgb = XGBClassifier(random_state=42,n_jobs=-1)

param_grid = {
    'n_estimators': [100, 200],
    'max_depth': [3, 5, 7],
    'learning_rate': [0.001, 0.01, 0.1],
    'subsample': [0.7, 0.8, 0.9],
    'colsample_bytree': [0.7, 0.8, 1.0],
    'min_child_weight': [1, 3, 5],
}

inner_cv = RepeatedStratifiedKFold(
    n_splits=5, n_repeats=2, random_state=42
)

grid = GridSearchCV(
    estimator=xgb,
    param_grid=param_grid,
    scoring='f1',
    cv=inner_cv,
    n_jobs=-1,
    verbose=0
)

grid.fit(X_train, y_train)
best_model = grid.best_estimator_
print("XGBoost classifier:")

```

```

print("Best params:", grid.best_params_)

scoring = {
    'accuracy': 'accuracy',
    'precision': 'precision',
    'recall': 'recall',
    'f1': 'f1',
    'roc_auc': 'roc_auc'
}
cv_res = cross_validate(
    best_model, X_train, y_train,
    cv=inner_cv, scoring=scoring, n_jobs=-1, return_train_score=False
)

def show(name):
    vals = cv_res[f'test_{name}']
    print(f"CV {name:<9}: {vals.mean():.4f} ± {vals.std():.4f}")

print("\n=== Repeated Stratified 5×2 CV on TRAIN ===")
for m in ['accuracy', 'precision', 'recall', 'f1', 'roc_auc']:
    show(m)

y_pred = best_model.predict(X_test)
y_prob = best_model.predict_proba(X_test)[:, 1]

acc = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred)
rec = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
auc = roc_auc_score(y_test, y_prob)

print("\n=== Final Performance on HOLD-OUT TEST ===")
print(f"Accuracy : {acc:.4f}")
print(f"Precision: {prec:.4f}")
print(f"Recall : {rec:.4f}")
print(f"F1 Score : {f1:.4f}")
print(f"ROC_AUC : {auc:.4f}")

print("\nClassification Report on Test Set:\n", classification_report(y_test,
    ↪ y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))

feat_imp = pd.DataFrame({
    "Feature": X_train.columns,
    "Importance": best_model.feature_importances_
}).sort_values(by="Importance", ascending=False)
print("\nTop feature importances:\n", feat_imp)

```

```
XGBoost_md1 = best_model
```

XGBoost classifier:

Best params: {'colsample_bytree': 1.0, 'learning_rate': 0.1, 'max_depth': 7, 'min_child_weight': 1, 'n_estimators': 200, 'subsample': 0.9}

=== Repeated Stratified 5×2 CV on TRAIN ===

CV accuracy : 0.7781 ± 0.0621

CV precision: 0.7500 ± 0.0845

CV recall : 0.7424 ± 0.0859

CV f1 : 0.7433 ± 0.0713

CV roc_auc : 0.8492 ± 0.0526

=== Final Performance on HOLD-OUT TEST ===

Accuracy : 0.8235

Precision: 0.7647

Recall : 0.8667

F1 Score : 0.8125

ROC_AUC : 0.8868

Classification Report on Test Set:

	precision	recall	f1-score	support
0	0.88	0.79	0.83	38
1	0.76	0.87	0.81	30
accuracy			0.82	68
macro avg	0.82	0.83	0.82	68
weighted avg	0.83	0.82	0.82	68

Confusion Matrix:

```
[[30  8]
 [ 4 26]]
```

Top feature importances:


	Feature	Importance
1	EDUC	0.382794
3	nWBV	0.239561
2	eTIV	0.204472
0	Age	0.173173

```
[ ]: fpr, tpr, _ = roc_curve(y_test, y_prob)
     roc_auc = roc_auc_score(y_test, y_prob)

     plt.figure()
     plt.plot(fpr, tpr, label=f'ROC curve (AUC = {roc_auc:.3f})')
```



```
plt.plot([0, 1], [0, 1], linestyle='--', linewidth=1)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('XGBoost Receiver Operating Characteristic')
plt.legend(loc='lower right')
plt.show()
```



```
[ ]: cm = confusion_matrix(y_test, y_pred)
tn, fp, fn, tp = cm.ravel()

labels = np.array([
    [f"TN = {tn}", f"FP = {fp}"],
    [f"FN = {fn}", f"TP = {tp}"]
])


fig, ax = plt.subplots()
ax.matshow(cm, cmap=plt.cm.Blues, alpha=0.3)
for i in range(2):
    for j in range(2):
        ax.text(j, i, labels[i, j],
```

```

        va='center', ha='center', fontsize=12, fontweight='bold')

plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix with Labels")
plt.show()

```



```

[ ]: # Find all FN indices in the full test set
FN_all = (~y_pred) & (y_test == 1)
FN_indices = y_test[FN_all].index
print("False Negative indices:", FN_indices)

# Find all FP indices in the full test set
FP_all = (y_pred) & (y_test == 0)
FP_indices = y_test[FP_all].index
print("False Positive indices:", FP_indices)

# Find all TP indices in the full test set
TP_all = (y_pred) & (y_test == 1)

```

```

TP_indices = y_test[TP_all].index
print("True Positive indices:", TP_indices)

# Find all FN indices in the full test set
TN_all = (~y_pred) & (y_test == 0)
TN_indices = y_test[TN_all].index
print("True Negative indices:", TN_indices)

```

```

False Negative indices: Index([52, 300, 51, 94], dtype='int64')
False Positive indices: Index([167, 85, 146, 198, 130, 199, 7, 64],
dtype='int64')
True Positive indices: Index([124, 332, 250, 317, 154, 25, 90, 106, 172, 285,
87, 215, 127, 3,
239, 162, 299, 345, 72, 39, 89, 88, 16, 329, 365, 275],
dtype='int64')
True Negative indices: Index([ 84, 122, 311, 48, 336, 213, 9, 210, 113, 363,
66, 5, 153, 291,
370, 158, 102, 243, 132, 306, 337, 196, 351, 179, 309, 197, 362, 209,
333, 96],
dtype='int64')

```

```

[ ]: FN_sample_test_idx = X_test.index.get_indexer_for([300, 299, 51, 94])
FP_sample_test_idx = X_test.index.get_indexer_for([167, 146, 198, 130, 199, 7, 64])
TP_sample_test_idx = X_test.index.get_indexer_for([124, 332, 250, 317, 154, 25, 90, 106, 172, 285, 87, 215, 127, 3, 239, 162, 345, 72, 39, 89, 88, 16, 329, 365, 275])
TN_sample_test_idx = X_test.index.get_indexer_for([84, 122, 311, 48, 336, 213, 9, 210, 113, 363, 66, 5, 153, 291, 370, 158, 102, 243, 132, 306, 337, 196, 351, 179, 309, 197, 362, 209, 333, 96])

```

```

[ ]: print("FN_sample_test_idx: ", FN_sample_test_idx)
print("FP_sample_test_idx: ", FP_sample_test_idx)
print("TP_sample_test_idx: ", TP_sample_test_idx)
print("TN_sample_test_idx: ", TN_sample_test_idx)

```

```

FN_sample_test_idx: [31 37 49 61]
FP_sample_test_idx: [17 21 42 50 53 59 60]
TP_sample_test_idx: [ 1  2  6  7  9 11 13 14 16 20 22 28 29 30 33 35 36 39 40
41 47 55 56 64
65 67]
TN_sample_test_idx: [ 0  3  4  5  8 10 12 15 18 19 23 24 25 26 27 32 34 38 43
44 45 46 48 51
52 54 57 58 62 63 66]

```


```
[ ]: XGB_explainer = shap.Explainer(XGBoost_md1)
XGB_shap = XGB_explainer(X_test)
print(type(XGB_explainer))
```

```
<class 'shap.explainers._tree.TreeExplainer'>
```


```
[ ]: print("Values dimensions: %s" % (XGB_shap.values.shape,))
print("Data dimensions: %s" % (XGB_shap.data.shape,))
```

```
Values dimensions: (68, 4)
Data dimensions: (68, 4)
```

```
[ ]: sb.reset_orig()
shap.plots.bar(XGB_shap)
```



```
[ ]: shap.plots.beeswarm(XGB_shap)
```



From the shap plot above, we can see that M/F(M=0, F=1) female and older age tend to push class to nondemented side.

7.3.6 Model training-LightGBM_remove SES and Gender

```
[ ]: LightGBM = LGBMClassifier(random_state=42, verbosity=-1)

param_grid = {
    'n_estimators': [100, 200],
    'max_depth': [3, 5, 7],
    'learning_rate': [0.001, 0.01, 0.1],
    'num_leaves': [31, 50, 70],
}

inner_cv = RepeatedStratifiedKFold(
    n_splits=5, n_repeats=2, random_state=42
)

grid = GridSearchCV(
    estimator=LightGBM,
    param_grid=param_grid,
    scoring='f1',
    cv=inner_cv,
    n_jobs=-1,
    verbose=0
)

grid.fit(X_train, y_train)
best_model = grid.best_estimator_
print("LightGBM classifier:")
```

```

print("Best params:", grid.best_params_)

scoring = {
    'accuracy': 'accuracy',
    'precision': 'precision',
    'recall': 'recall',
    'f1': 'f1',
    'roc_auc': 'roc_auc'
}
cv_res = cross_validate(
    best_model, X_train, y_train,
    cv=inner_cv, scoring=scoring, n_jobs=-1, return_train_score=False
)

def show(name):
    vals = cv_res[f'test_{name}']
    print(f"CV {name:<9}: {vals.mean():.4f} ± {vals.std():.4f}")

print("\n=== Repeated Stratified 5×2 CV on TRAIN ===")
for m in ['accuracy', 'precision', 'recall', 'f1', 'roc_auc']:
    show(m)

y_pred = best_model.predict(X_test)
y_prob = best_model.predict_proba(X_test)[:, 1]

acc = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred)
rec = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
auc = roc_auc_score(y_test, y_prob)

print("\n=== Final Performance on HOLD-OUT TEST ===")
print(f"Accuracy : {acc:.4f}")
print(f"Precision: {prec:.4f}")
print(f"Recall : {rec:.4f}")
print(f"F1 Score : {f1:.4f}")
print(f"ROC_AUC : {auc:.4f}")

print("\nClassification Report on Test Set:\n", classification_report(y_test,
    ↪ y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))

feat_imp = pd.DataFrame({
    "Feature": X_train.columns,
    "Importance": best_model.feature_importances_
}).sort_values(by="Importance", ascending=False)
print("\nTop feature importances:\n", feat_imp)

```

```
LightGBM_md1 = best_model
```

LightGBM classifier:

Best params: {'learning_rate': 0.1, 'max_depth': 5, 'n_estimators': 200, 'num_leaves': 31}

=== Repeated Stratified 5×2 CV on TRAIN ===

CV accuracy : 0.7632 ± 0.0448

CV precision: 0.7329 ± 0.0758

CV recall : 0.7333 ± 0.0536

CV f1 : 0.7294 ± 0.0406

CV roc_auc : 0.8323 ± 0.0565

=== Final Performance on HOLD-OUT TEST ===

Accuracy : 0.7941

Precision: 0.7500

Recall : 0.8000

F1 Score : 0.7742

ROC_AUC : 0.8860

Classification Report on Test Set:

	precision	recall	f1-score	support
0	0.83	0.79	0.81	38
1	0.75	0.80	0.77	30
accuracy			0.79	68
macro avg	0.79	0.79	0.79	68
weighted avg	0.80	0.79	0.79	68

Confusion Matrix:

```
[[30  8]
 [ 6 24]]
```


Top feature importances:

	Feature	Importance
2	eTIV	674
3	nWBV	423
0	Age	340
1	EDUC	200

```
[ ]: fpr, tpr, _ = roc_curve(y_test, y_prob)
     roc_auc = roc_auc_score(y_test, y_prob)

     plt.figure()
     plt.plot(fpr, tpr, label=f'ROC curve (AUC = {roc_auc:.3f})')
```

```
plt.plot([0, 1], [0, 1], linestyle='--', linewidth=1)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('LightGBM Receiver Operating Characteristic')
plt.legend(loc='lower right')
plt.show()
```



```
[ ]: cm = confusion_matrix(y_test, y_pred)
tn, fp, fn, tp = cm.ravel()

labels = np.array([
    [f"TN = {tn}", f"FP = {fp}"],
    [f"FN = {fn}", f"TP = {tp}"]
])

fig, ax = plt.subplots()
ax.matshow(cm, cmap=plt.cm.Blues, alpha=0.3)
for i in range(2):
    for j in range(2):
        ax.text(j, i, labels[i, j],
```




```

        va='center', ha='center', fontsize=12, fontweight='bold')

plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix with Labels")
plt.show()

```



```

[ ]: # Find all FN indices in the full test set
FN_all = (~y_pred) & (y_test == 1)
FN_indices = y_test[FN_all].index
print("False Negative indices:", FN_indices)

```

False Negative indices: Index([52, 300, 345, 51, 94], dtype='int64')

```

[ ]: FN_sample_test_idx = X_test.index.get_indexer_for([154, 300, 299, 39, 51, 94,
↪329])
FN_sample_test_idx

```

```

[ ]: array([ 9, 31, 37, 41, 49, 61, 64])

```

```
[ ]: LGB_explainer = shap.Explainer(LightGBM_md1)
LGB_shap = LGB_explainer(X_test)
print(type(LGB_explainer))
```


```
<class 'shap.explainers._tree.TreeExplainer'>
```

```
[ ]: print("Values dimensions: %s" % (LGB_shap.values.shape,))
print("Data dimensions: %s" % (LGB_shap.data.shape,))
```


```
Values dimensions: (68, 4)
```

```
Data dimensions: (68, 4)
```

```
[ ]: sb.reset_orig()
shap.plots.bar(LGB_shap)
```



```
[ ]: shap.plots.beeswarm(LGB_shap)
```



7.3.7 Model training-CatBoost_remove SES and Gender

```
[ ]: CatBoost = CatBoostClassifier(random_state=42)

param_grid = {
    'min_data_in_leaf': [20, 40, 60],
    'rsm': [0.7, 0.8, 1.0],
    'iterations': [100, 200],
    'depth': [3, 5, 7],
    'learning_rate': [0.001, 0.01, 0.1],
    'l2_leaf_reg': [1, 3, 5],
    'bagging_temperature': [0, 0.5, 1.0],
}

inner_cv = RepeatedStratifiedKFold(
    n_splits=5, n_repeats=2, random_state=42
)

grid = GridSearchCV(
    estimator=CatBoost,
    param_grid=param_grid,
    scoring='f1',
    cv=inner_cv,
    n_jobs=-1,
    verbose=0
)

grid.fit(X_train, y_train)
best_model = grid.best_estimator_
print("CatBoostClassifier:")
```

```

print("Best params:", grid.best_params_)

scoring = {
    'accuracy': 'accuracy',
    'precision': 'precision',
    'recall': 'recall',
    'f1': 'f1',
    'roc_auc': 'roc_auc'
}
cv_res = cross_validate(
    best_model, X_train, y_train,
    cv=inner_cv, scoring=scoring, n_jobs=-1, return_train_score=False
)

def show(name):
    vals = cv_res[f'test_{name}']
    print(f"CV {name:<9}: {vals.mean():.4f} ± {vals.std():.4f}")

print("\n=== Repeated Stratified 5×2 CV on TRAIN ===")
for m in ['accuracy', 'precision', 'recall', 'f1', 'roc_auc']:
    show(m)

y_pred = best_model.predict(X_test)
y_prob = best_model.predict_proba(X_test)[:, 1]

acc = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred)
rec = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
auc = roc_auc_score(y_test, y_prob)

print("\n=== Final Performance on HOLD-OUT TEST ===")
print(f"Accuracy : {acc:.4f}")
print(f"Precision: {prec:.4f}")
print(f"Recall : {rec:.4f}")
print(f"F1 Score : {f1:.4f}")
print(f"ROC_AUC : {auc:.4f}")

print("\nClassification Report on Test Set:\n", classification_report(y_test,
    ↪ y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))

feat_imp = pd.DataFrame({
    "Feature": X_train.columns,
    "Importance": best_model.feature_importances_
}).sort_values(by="Importance", ascending=False)
print("\nTop feature importances:\n", feat_imp)

```

```
CatBoost_md1 = best_model
```

0:	learn: 0.6433030	total: 48.1ms	remaining: 9.57s
1:	learn: 0.6072231	total: 49.6ms	remaining: 4.91s
2:	learn: 0.5877394	total: 50.6ms	remaining: 3.33s
3:	learn: 0.5621783	total: 51.9ms	remaining: 2.54s
4:	learn: 0.5421796	total: 53ms	remaining: 2.07s
5:	learn: 0.5222894	total: 53.9ms	remaining: 1.74s
6:	learn: 0.5080146	total: 55ms	remaining: 1.52s
7:	learn: 0.4919294	total: 56.1ms	remaining: 1.35s
8:	learn: 0.4779497	total: 57ms	remaining: 1.21s
9:	learn: 0.4587318	total: 58.1ms	remaining: 1.1s
10:	learn: 0.4423055	total: 59.1ms	remaining: 1.02s
11:	learn: 0.4201557	total: 60.9ms	remaining: 954ms
12:	learn: 0.4078103	total: 62.5ms	remaining: 899ms
13:	learn: 0.4023800	total: 63.8ms	remaining: 847ms
14:	learn: 0.3942624	total: 64.9ms	remaining: 800ms
15:	learn: 0.3764204	total: 66ms	remaining: 759ms
16:	learn: 0.3745783	total: 66.5ms	remaining: 716ms
17:	learn: 0.3642996	total: 67.7ms	remaining: 685ms
18:	learn: 0.3580881	total: 68.8ms	remaining: 656ms
19:	learn: 0.3519850	total: 70.1ms	remaining: 631ms
20:	learn: 0.3396768	total: 71.2ms	remaining: 607ms
21:	learn: 0.3359188	total: 72.3ms	remaining: 585ms
22:	learn: 0.3311288	total: 73.5ms	remaining: 565ms
23:	learn: 0.3220628	total: 74.5ms	remaining: 546ms
24:	learn: 0.3180902	total: 75.6ms	remaining: 529ms
25:	learn: 0.3087211	total: 76.7ms	remaining: 513ms
26:	learn: 0.3037849	total: 77.8ms	remaining: 498ms
27:	learn: 0.2990716	total: 78.9ms	remaining: 485ms
28:	learn: 0.2913764	total: 80.2ms	remaining: 473ms
29:	learn: 0.2869979	total: 81.9ms	remaining: 464ms
30:	learn: 0.2856476	total: 83.2ms	remaining: 454ms
31:	learn: 0.2817315	total: 84.3ms	remaining: 442ms
32:	learn: 0.2797068	total: 85.5ms	remaining: 433ms
33:	learn: 0.2781957	total: 86.6ms	remaining: 423ms
34:	learn: 0.2757684	total: 87.6ms	remaining: 413ms
35:	learn: 0.2724397	total: 88.8ms	remaining: 405ms
36:	learn: 0.2662034	total: 89.9ms	remaining: 396ms
37:	learn: 0.2573439	total: 91ms	remaining: 388ms
38:	learn: 0.2537094	total: 92.1ms	remaining: 380ms
39:	learn: 0.2480944	total: 93.3ms	remaining: 373ms
40:	learn: 0.2405935	total: 94.5ms	remaining: 366ms
41:	learn: 0.2383741	total: 95.6ms	remaining: 360ms
42:	learn: 0.2321652	total: 96.7ms	remaining: 353ms
43:	learn: 0.2302174	total: 97.8ms	remaining: 347ms

44:	learn: 0.2276557	total: 98.9ms	remaining: 341ms
45:	learn: 0.2245745	total: 100ms	remaining: 336ms
46:	learn: 0.2199735	total: 102ms	remaining: 331ms
47:	learn: 0.2170957	total: 103ms	remaining: 325ms
48:	learn: 0.2124202	total: 104ms	remaining: 320ms
49:	learn: 0.2119652	total: 105ms	remaining: 315ms
50:	learn: 0.2090747	total: 106ms	remaining: 310ms
51:	learn: 0.2059814	total: 107ms	remaining: 305ms
52:	learn: 0.2032197	total: 108ms	remaining: 300ms
53:	learn: 0.1995946	total: 109ms	remaining: 296ms
54:	learn: 0.1973423	total: 111ms	remaining: 291ms
55:	learn: 0.1948246	total: 112ms	remaining: 287ms
56:	learn: 0.1907309	total: 113ms	remaining: 283ms
57:	learn: 0.1895778	total: 114ms	remaining: 278ms
58:	learn: 0.1884515	total: 115ms	remaining: 274ms
59:	learn: 0.1854442	total: 116ms	remaining: 271ms
60:	learn: 0.1809831	total: 117ms	remaining: 268ms
61:	learn: 0.1782131	total: 119ms	remaining: 264ms
62:	learn: 0.1737153	total: 120ms	remaining: 261ms
63:	learn: 0.1716256	total: 121ms	remaining: 257ms
64:	learn: 0.1670977	total: 122ms	remaining: 254ms
65:	learn: 0.1629564	total: 123ms	remaining: 251ms
66:	learn: 0.1605172	total: 124ms	remaining: 247ms
67:	learn: 0.1575741	total: 125ms	remaining: 244ms
68:	learn: 0.1557338	total: 127ms	remaining: 240ms
69:	learn: 0.1525787	total: 128ms	remaining: 237ms
70:	learn: 0.1496277	total: 129ms	remaining: 234ms
71:	learn: 0.1469448	total: 130ms	remaining: 231ms
72:	learn: 0.1454756	total: 131ms	remaining: 229ms
73:	learn: 0.1438074	total: 133ms	remaining: 226ms
74:	learn: 0.1412499	total: 134ms	remaining: 224ms
75:	learn: 0.1394253	total: 136ms	remaining: 221ms
76:	learn: 0.1378607	total: 137ms	remaining: 218ms
77:	learn: 0.1371424	total: 138ms	remaining: 216ms
78:	learn: 0.1351813	total: 139ms	remaining: 213ms
79:	learn: 0.1327478	total: 141ms	remaining: 212ms
80:	learn: 0.1280091	total: 143ms	remaining: 210ms
81:	learn: 0.1265301	total: 144ms	remaining: 207ms
82:	learn: 0.1230572	total: 145ms	remaining: 204ms
83:	learn: 0.1224970	total: 146ms	remaining: 202ms
84:	learn: 0.1202086	total: 148ms	remaining: 200ms
85:	learn: 0.1184823	total: 149ms	remaining: 197ms
86:	learn: 0.1173750	total: 150ms	remaining: 194ms
87:	learn: 0.1166478	total: 151ms	remaining: 192ms
88:	learn: 0.1140979	total: 152ms	remaining: 190ms
89:	learn: 0.1128446	total: 154ms	remaining: 188ms
90:	learn: 0.1101893	total: 156ms	remaining: 187ms
91:	learn: 0.1088455	total: 158ms	remaining: 185ms

92:	learn: 0.1067866	total: 159ms	remaining: 183ms
93:	learn: 0.1045246	total: 160ms	remaining: 181ms
94:	learn: 0.1039400	total: 161ms	remaining: 178ms
95:	learn: 0.1017326	total: 162ms	remaining: 176ms
96:	learn: 0.0982394	total: 163ms	remaining: 173ms
97:	learn: 0.0977207	total: 164ms	remaining: 171ms
98:	learn: 0.0973069	total: 166ms	remaining: 169ms
99:	learn: 0.0967707	total: 167ms	remaining: 167ms
100:	learn: 0.0944164	total: 168ms	remaining: 165ms
101:	learn: 0.0938359	total: 169ms	remaining: 162ms
102:	learn: 0.0923871	total: 170ms	remaining: 160ms
103:	learn: 0.0916017	total: 171ms	remaining: 158ms
104:	learn: 0.0905539	total: 172ms	remaining: 155ms
105:	learn: 0.0887147	total: 173ms	remaining: 153ms
106:	learn: 0.0875213	total: 174ms	remaining: 151ms
107:	learn: 0.0866559	total: 175ms	remaining: 149ms
108:	learn: 0.0854698	total: 176ms	remaining: 147ms
109:	learn: 0.0842835	total: 177ms	remaining: 145ms
110:	learn: 0.0837254	total: 178ms	remaining: 143ms
111:	learn: 0.0808711	total: 179ms	remaining: 141ms
112:	learn: 0.0803991	total: 180ms	remaining: 139ms
113:	learn: 0.0793167	total: 182ms	remaining: 137ms
114:	learn: 0.0785392	total: 183ms	remaining: 135ms
115:	learn: 0.0766991	total: 184ms	remaining: 134ms
116:	learn: 0.0759175	total: 186ms	remaining: 132ms
117:	learn: 0.0754021	total: 187ms	remaining: 130ms
118:	learn: 0.0742354	total: 189ms	remaining: 128ms
119:	learn: 0.0739136	total: 190ms	remaining: 126ms
120:	learn: 0.0733437	total: 191ms	remaining: 125ms
121:	learn: 0.0729341	total: 191ms	remaining: 122ms
122:	learn: 0.0721195	total: 192ms	remaining: 120ms
123:	learn: 0.0718087	total: 193ms	remaining: 118ms
124:	learn: 0.0703325	total: 194ms	remaining: 117ms
125:	learn: 0.0700263	total: 195ms	remaining: 115ms
126:	learn: 0.0693124	total: 196ms	remaining: 113ms
127:	learn: 0.0687718	total: 198ms	remaining: 111ms
128:	learn: 0.0685893	total: 198ms	remaining: 109ms
129:	learn: 0.0675148	total: 199ms	remaining: 107ms
130:	learn: 0.0669490	total: 200ms	remaining: 106ms
131:	learn: 0.0658678	total: 202ms	remaining: 104ms
132:	learn: 0.0655645	total: 203ms	remaining: 102ms
133:	learn: 0.0646241	total: 204ms	remaining: 100ms
134:	learn: 0.0638795	total: 205ms	remaining: 98.5ms
135:	learn: 0.0631970	total: 206ms	remaining: 96.8ms
136:	learn: 0.0629538	total: 207ms	remaining: 95ms
137:	learn: 0.0626684	total: 208ms	remaining: 93.3ms
138:	learn: 0.0621399	total: 209ms	remaining: 91.5ms
139:	learn: 0.0616209	total: 209ms	remaining: 89.8ms

140:	learn: 0.0604167	total: 210ms	remaining: 88ms
141:	learn: 0.0593850	total: 211ms	remaining: 86.3ms
142:	learn: 0.0588888	total: 212ms	remaining: 84.6ms
143:	learn: 0.0584728	total: 213ms	remaining: 83ms
144:	learn: 0.0575776	total: 214ms	remaining: 81.3ms
145:	learn: 0.0573503	total: 215ms	remaining: 79.6ms
146:	learn: 0.0570246	total: 216ms	remaining: 78ms
147:	learn: 0.0564381	total: 217ms	remaining: 76.3ms
148:	learn: 0.0562654	total: 218ms	remaining: 74.7ms
149:	learn: 0.0556245	total: 219ms	remaining: 73.1ms
150:	learn: 0.0551304	total: 220ms	remaining: 71.5ms
151:	learn: 0.0549703	total: 221ms	remaining: 69.9ms
152:	learn: 0.0543406	total: 222ms	remaining: 68.3ms
153:	learn: 0.0537970	total: 224ms	remaining: 66.8ms
154:	learn: 0.0533642	total: 225ms	remaining: 65.4ms
155:	learn: 0.0528259	total: 227ms	remaining: 63.9ms
156:	learn: 0.0523555	total: 229ms	remaining: 62.7ms
157:	learn: 0.0515599	total: 230ms	remaining: 61.2ms
158:	learn: 0.0514442	total: 232ms	remaining: 59.7ms
159:	learn: 0.0512947	total: 233ms	remaining: 58.2ms
160:	learn: 0.0512666	total: 234ms	remaining: 56.6ms
161:	learn: 0.0510238	total: 235ms	remaining: 55ms
162:	learn: 0.0505527	total: 236ms	remaining: 53.5ms
163:	learn: 0.0498957	total: 237ms	remaining: 51.9ms
164:	learn: 0.0494456	total: 238ms	remaining: 50.5ms
165:	learn: 0.0492686	total: 240ms	remaining: 49.2ms
166:	learn: 0.0489407	total: 241ms	remaining: 47.6ms
167:	learn: 0.0485755	total: 242ms	remaining: 46.1ms
168:	learn: 0.0480240	total: 243ms	remaining: 44.6ms
169:	learn: 0.0471546	total: 245ms	remaining: 43.2ms
170:	learn: 0.0464787	total: 246ms	remaining: 41.7ms
171:	learn: 0.0459925	total: 247ms	remaining: 40.1ms
172:	learn: 0.0455186	total: 248ms	remaining: 38.7ms
173:	learn: 0.0453493	total: 249ms	remaining: 37.2ms
174:	learn: 0.0449091	total: 250ms	remaining: 35.7ms
175:	learn: 0.0443691	total: 251ms	remaining: 34.2ms
176:	learn: 0.0440176	total: 252ms	remaining: 32.7ms
177:	learn: 0.0435075	total: 253ms	remaining: 31.3ms
178:	learn: 0.0432383	total: 254ms	remaining: 29.8ms
179:	learn: 0.0429571	total: 255ms	remaining: 28.3ms
180:	learn: 0.0427847	total: 256ms	remaining: 26.9ms
181:	learn: 0.0422452	total: 257ms	remaining: 25.4ms
182:	learn: 0.0418478	total: 258ms	remaining: 24ms
183:	learn: 0.0414921	total: 259ms	remaining: 22.6ms
184:	learn: 0.0411721	total: 260ms	remaining: 21.1ms
185:	learn: 0.0409214	total: 261ms	remaining: 19.7ms
186:	learn: 0.0405452	total: 262ms	remaining: 18.2ms
187:	learn: 0.0402071	total: 263ms	remaining: 16.8ms

188:	learn: 0.0398397	total: 264ms	remaining: 15.4ms
189:	learn: 0.0393759	total: 265ms	remaining: 14ms
190:	learn: 0.0392496	total: 266ms	remaining: 12.5ms
191:	learn: 0.0388633	total: 267ms	remaining: 11.1ms
192:	learn: 0.0384548	total: 268ms	remaining: 9.73ms
193:	learn: 0.0382617	total: 269ms	remaining: 8.33ms
194:	learn: 0.0381814	total: 270ms	remaining: 6.93ms
195:	learn: 0.0379160	total: 271ms	remaining: 5.54ms
196:	learn: 0.0377706	total: 273ms	remaining: 4.15ms
197:	learn: 0.0374555	total: 274ms	remaining: 2.76ms
198:	learn: 0.0373077	total: 275ms	remaining: 1.38ms
199:	learn: 0.0370488	total: 276ms	remaining: 0ms

CatBoostClassifier:

Best params: {'bagging_temperature': 0, 'depth': 7, 'iterations': 200, 'l2_leaf_reg': 3, 'learning_rate': 0.1, 'min_data_in_leaf': 20, 'rsm': 0.8}

=== Repeated Stratified 5x2 CV on TRAIN ===

CV accuracy : 0.8194 ± 0.0687
 CV precision: 0.7995 ± 0.0876
 CV recall : 0.7855 ± 0.1006
 CV f1 : 0.7893 ± 0.0823
 CV roc_auc : 0.8852 ± 0.0432

=== Final Performance on HOLD-OUT TEST ===

Accuracy : 0.8824
 Precision: 0.8929
 Recall : 0.8333
 F1 Score : 0.8621
 ROC_AUC : 0.9211

Classification Report on Test Set:

	precision	recall	f1-score	support
0	0.88	0.92	0.90	38
1	0.89	0.83	0.86	30
accuracy			0.88	68
macro avg	0.88	0.88	0.88	68
weighted avg	0.88	0.88	0.88	68

Confusion Matrix:

```
[[35  3]
 [ 5 25]]
```


Top feature importances:

	Feature	Importance
2	eTIV	32.923596
3	nWBV	25.126673

```
1    EDUC    24.059391
0    Age     17.890340
```

```
[ ]: fpr, tpr, _ = roc_curve(y_test, y_prob)
roc_auc = roc_auc_score(y_test, y_prob)

plt.figure()
plt.plot(fpr, tpr, label=f'ROC curve (AUC = {roc_auc:.3f})')
plt.plot([0, 1], [0, 1], linestyle='--', linewidth=1)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('CatBoost Receiver Operating Characteristic')
plt.legend(loc='lower right')
plt.show()
```



```
[ ]: cm = confusion_matrix(y_test, y_pred)
tn, fp, fn, tp = cm.ravel()

labels = np.array([
    f"TN = {tn}", f"FP = {fp}"],
```


```

    [f"FN = {fn}", f"TP = {tp}"]
])

fig, ax = plt.subplots()
ax.matshow(cm, cmap=plt.cm.Blues, alpha=0.3)
for i in range(2):
    for j in range(2):
        ax.text(j, i, labels[i, j],
                va='center', ha='center', fontsize=12, fontweight='bold')

plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix with Labels")
plt.show()

```



```

[ ]: # Find all FN indices in the full test set
FN_all = (~y_pred) & (y_test == 1)
FN_indices = y_test[FN_all].index
print("False Negative indices:", FN_indices)

```

```
False Negative indices: Index([52, 300, 345, 51, 94], dtype='int64')
```

```
[ ]: FN_sample_test_idx = X_test.index.get_indexer_for([300, 299, 51, 94])
     FN_sample_test_idx
```

```
[ ]: array([31, 37, 49, 61])
```

```
[ ]: CB_explainer = shap.Explainer(CatBoost_md1)
     CB_shap = CB_explainer(X_test)
     print(type(CB_explainer))
```


```
<class 'shap.explainers._tree.TreeExplainer'>
```

```
[ ]: print("Values dimensions: %s" % (CB_shap.values.shape,))
     print("Data dimensions:   %s" % (CB_shap.data.shape,))
```


```
Values dimensions: (68, 4)
```

```
Data dimensions:   (68, 4)
```

```
[ ]: sb.reset_orig()
     shap.plots.bar(CB_shap)
```



```
[ ]: shap.plots.beeswarm(CB_shap)
```



7.3.8 Model training-CatBoost_remove Age and Gender

```
[ ]: # Initialize a CatBoost classifier
CatBoost = CatBoostClassifier(random_state=42, verbose = False)

# Define the parameter grid for Grid
CatBoost_param_dist = {
    'min_data_in_leaf': [20, 40, 60],
    'rsm': [0.7, 0.8, 1.0],
    'iterations': [100, 200],
    'depth': [3, 5, 7],
    'learning_rate': [0.001, 0.01, 0.1],
    'l2_leaf_reg': [1, 3, 5],
    'bagging_temperature': [0, 0.5, 1.0],
}

grid_search = GridSearchCV(CatBoost, CatBoost_param_dist, cv=5,
    scoring='recall', n_jobs=-1)
grid_search.fit(X_train, y_train)
print("Best hyperparameters found by GridSearchCV:")
print(grid_search.best_params_)

CatBoost_md1 = grid_search.best_estimator_
y_pred = CatBoost_md1.predict(X_test)
```

```
-----
KeyboardInterrupt                                Traceback (most recent call last)
/tmp/ipython-input-3511066618.py in <cell line: 0>()
    14
```

```

15 grid_search = GridSearchCV(CatBoost, CatBoost_param_dist, cv=5,
↳ scoring='recall', n_jobs=-1)
---> 16 grid_search.fit(X_train, y_train)
17 print("Best hyperparameters found by GridSearchCV:")
18 print(grid_search.best_params_)

/usr/local/lib/python3.11/dist-packages/sklearn/base.py in wrapper(estimator,
↳ *args, **kwargs)
1387         )
1388     ):
-> 1389         return fit_method(estimator, *args, **kwargs)
1390
1391     return wrapper

/usr/local/lib/python3.11/dist-packages/sklearn/model_selection/_search.py in
↳ fit(self, X, y, **params)
1022         return results
1023
-> 1024         self._run_search(evaluate_candidates)
1025
1026         # multimetric is determined here because in the case of a
↳ callable

/usr/local/lib/python3.11/dist-packages/sklearn/model_selection/_search.py in
↳ _run_search(self, evaluate_candidates)
1569     def _run_search(self, evaluate_candidates):
1570         """Search all candidates in param_grid"""
-> 1571         evaluate_candidates(ParameterGrid(self.param_grid))
1572
1573

/usr/local/lib/python3.11/dist-packages/sklearn/model_selection/_search.py in
↳ evaluate_candidates(candidate_params, cv, more_results)
968         )
969
--> 970         out = parallel(

971             delayed(_fit_and_score)(
972                 clone(base_estimator),

/usr/local/lib/python3.11/dist-packages/sklearn/utils/parallel.py in
↳ __call__(self, iterable)
75         for delayed_func, args, kwargs in iterable
76     )
---> 77     return super().__call__(iterable_with_config)
78
79

```

```

/usr/local/lib/python3.11/dist-packages/joblib/parallel.py in __call__(self,
↳ iterable)
    2070         next(output)
    2071
-> 2072         return output if self.return_generator else list(output)
    2073
    2074     def __repr__(self):

/usr/local/lib/python3.11/dist-packages/joblib/parallel.py in _get_outputs(self,
↳ iterator, pre_dispatch)
    1680
    1681         with self._backend.retrieval_context():
-> 1682             yield from self._retrieve()
    1683
    1684     except GeneratorExit:

/usr/local/lib/python3.11/dist-packages/joblib/parallel.py in _retrieve(self)
    1798         self._jobs[0].get_status(timeout=self.timeout) ==
↳ TASK_PENDING
    1799         ):
-> 1800             time.sleep(0.01)
    1801             continue
    1802

KeyboardInterrupt:

```

CatBoost Best hyperparameters found by GridSearchCV:

```
{'bagging_temperature': 0, 'depth': 5, 'iterations': 200, 'l2_leaf_reg': 1, 'learning_rate': 0.1, 'min_data_in_leaf': 20, 'rsm': 1.0}
```

```

[ ]: print("CatBoost: ")
print("\nClassification Report on Test Set:")
print(classification_report(y_test, y_pred))
print("-----")

print("Accuracy:", accuracy_score(y_test, y_pred))
print("Precision:", precision_score(y_test, y_pred))
print("Recall:", recall_score(y_test, y_pred))
print("F1 Score:", f1_score(y_test, y_pred))
print("ROC_AUC:", roc_auc_score(y_test, y_pred))
print("-----")

# Get feature importances
feature_importances = pd.DataFrame({
    "Feature": X_test.columns,
    "Importance": CatBoost_md1.feature_importances_

```

```

})

feature_importances = feature_importances.sort_values(by="Importance",
↪ascending=False)
feature_importances

```

CatBoost:

Classification Report on Test Set:

	precision	recall	f1-score	support
0	0.85	0.92	0.89	38
1	0.89	0.80	0.84	30
accuracy			0.87	68
macro avg	0.87	0.86	0.86	68
weighted avg	0.87	0.87	0.87	68

```

-----
Accuracy: 0.8676470588235294
Precision: 0.8888888888888888
Recall: 0.8
F1 Score: 0.8421052631578947
ROC_AUC: 0.8605263157894737
-----

```

```

[ ]:  Feature  Importance
2    eTIV    37.023811
3    nWBV    25.111532
1     SES    19.110291
0    EDUC    18.754366

```

```

[ ]: cm = confusion_matrix(y_test, y_pred)
tn, fp, fn, tp = cm.ravel()


labels = np.array([
    [f"TN = {tn}", f"FP = {fp}"],
    [f"FN = {fn}", f"TP = {tp}"]
])

fig, ax = plt.subplots()
ax.matshow(cm, cmap=plt.cm.Blues, alpha=0.3)
for i in range(2):
    for j in range(2):
        ax.text(j, i, labels[i, j],
                va='center', ha='center', fontsize=12, fontweight='bold')

```



```
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix with Labels")
plt.show()
```



```
[ ]: # Find all FN indices in the full test set
FN_all = (~y_pred) & (y_test == 1)
FN_indices = y_test[FN_all].index
print("False Negative indices:", FN_indices)
```

False Negative indices: Index([332, 52, 300, 299, 51, 94], dtype='int64')

```
[ ]: FN_sample_test_idx = X_test.index.get_indexer_for([332, 52, 300, 299, 51, 94])
FN_sample_test_idx
```

```
[ ]: array([ 2, 30, 31, 37, 49, 61])
```

```
[ ]: CB_explainer = shap.Explainer(CatBoost_md1)
CB_shap = CB_explainer(X_test)
print(type(CB_explainer))
```


```
<class 'shap.explainers._tree.TreeExplainer'>
```

```
[ ]: print("Values dimensions: %s" % (CB_shap.values.shape,))  
print("Data dimensions: %s" % (CB_shap.data.shape,))
```


```
Values dimensions: (68, 4)
```

```
Data dimensions: (68, 4)
```

```
[ ]: sb.reset_orig()  
shap.plots.bar(CB_shap)
```



```
[ ]: shap.plots.beeswarm(CB_shap)
```




```
[ ]: print("X_test.iloc[2]: ")
      print(X_test.iloc[2])
      print(y_test.iloc[2], y_pred[2])
      print("-----")
      print("X_test.iloc[31]: ")
      print(X_test.iloc[31])
      print(y_test.iloc[31], y_pred[31])
```


```
X_test.iloc[2]:
EDUC      12.000
SES        3.000
eTIV     1755.000
nWBV       0.696
Name: 332, dtype: float64
1 0
```

```
-----
X_test.iloc[31]:
EDUC      15.000
SES        2.000
eTIV     1729.000
nWBV       0.698
Name: 300, dtype: float64
1 0
```

```
[ ]: # CB-FN-2
      expected_value = CB_explainer.expected_value
      shap.decision_plot(expected_value, CB_shap.values[2], X_test.iloc[2],
      ↪ highlight=0)
```



```
[ ]: # CB-FN-31
expected_value = CB_explainer.expected_value
shap.decision_plot(expected_value, CB_shap.values[31], X_test.iloc[31],
↳highlight=0)
```



```
[ ]: # CB-FN-2
shap.initjs()
expected_value = CB_explainer.expected_value
shap.force_plot(expected_value, CB_shap.values[2], X_test.iloc[2])
```

<IPython.core.display.HTML object>

```
[ ]: <shap.plots._force.AdditiveForceVisualizer at 0x7dda42593b50>
```

```
[ ]: # CB-FN-31
shap.initjs()
expected_value = CB_explainer.expected_value
shap.force_plot(expected_value, CB_shap.values[31], X_test.iloc[31])
```

<IPython.core.display.HTML object>

```
[ ]: <shap.plots._force.AdditiveForceVisualizer at 0x7dda147b7690>
```

```
[ ]: lime_CB_explainer = lime.lime_tabular.LimeTabularExplainer(X_test.values,
                                                                feature_names=X_test.columns,
                                                                class_names=['Nondemented',
↳'Demented'])
```

```
[ ]: # CB-FN-2
lime_CB_explainer.explain_instance(X_test.iloc[2].values,\
                                   CatBoost_md1.predict_proba,\
                                   num_features=6).\
                                   show_in_notebook(predict_proba=True)
```

<IPython.core.display.HTML object>

```
[ ]: # CB-FN-31
lime_CB_explainer.explain_instance(X_test.iloc[31].values,\
                                   CatBoost_md1.predict_proba,\
                                   num_features=6).\
                                   show_in_notebook(predict_proba=True)
```

<IPython.core.display.HTML object>

```
[ ]: # store all fn_results
fn_results = []
feature_counter = Counter()

CB_FN_indices = [ 2, 30, 31, 37, 49, 61]

for fn_idx in CB_FN_indices:
    CB_instance_values = X_test.iloc[fn_idx].values

    exp = lime_CB_explainer.explain_instance(
        CB_instance_values,
        CatBoost_md1.predict_proba,
        num_features=6
    )

    exp_list = exp.as_list()

    pushed_non = [f for f, w in exp_list if w < 0]
    pushed_dem = [f for f, w in exp_list if w > 0]

    fn_results.append({
        'Index': fn_idx,
        'Pushed_Nondemented': pushed_non,
        'Pushed_Demented': pushed_dem
    })

    feature_counter.update(pushed_non)

fn_df = pd.DataFrame(fn_results)
```

```

top_causes = pd.DataFrame(feature_counter.most_common(), columns=['Feature', 'Count'])

print(fn_df.head())

print("\n=== CB False Negative Feature Frequency ===")
print(top_causes)

```

	Index	Pushed_Nondemented \
0	2	[eTIV > 1669.00]
1	30	[EDUC > 16.25, SES <= 2.00]
2	31	[eTIV > 1669.00, SES <= 2.00]
3	37	[eTIV > 1669.00, SES <= 2.00]
4	49	[EDUC > 16.25, SES <= 2.00, 0.73 < nWBV <= 0.76]

	Pushed_Demented
0	[nWBV <= 0.70, EDUC <= 12.00, 2.00 < SES <= 3.00]
1	[1491.50 < eTIV <= 1669.00, 0.70 < nWBV <= 0.73]
2	[nWBV <= 0.70, 12.00 < EDUC <= 15.00]
3	[nWBV <= 0.70, 12.00 < EDUC <= 15.00]
4	[1491.50 < eTIV <= 1669.00]

=== CB False Negative Feature Frequency ===

	Feature	Count
0	SES <= 2.00	5
1	eTIV > 1669.00	3
2	EDUC > 16.25	2
3	0.73 < nWBV <= 0.76	2

After removing gender and age features, only LightGBM performance drop, Catboost and XGBoost are almost the same.

We have calculated XGBoost model feature importance: Age top 5, gender top2; LightGBM model feature importance: Age top 3, gender top6; CatBoost model feature importance: Age top 4, gender top6. From model feature importance, LightGBM didn't prioritize age and gender over the other two models.

From shap value side: (MeanAbsSHAP(Age) + MeanAbsSHAP(Gender) proportion of total SHAP.)

XGBoost = 1.3/4.68=0.277; LightGBM=1.67/5.7=0.293; CatBoost = 1.29/4.23=0.305. LightGBM's proportion is not the highest.

so from both feature importance and SHAP proportion, LightGBM is not obviously relying on Age and Gender more than XGBoost or CatBoost. That means the performance drop in LightGBM probably isn't just because of "over-reliance" on those two features — it's more likely about how LightGBM's tree structure uses them.